

Wait-free Queues with Polyarithmic Step Complexity

February 1, 2022

1 Previous Work


	Type	Progress Property	Complexity
MichaelS96 [8] DBLP:conf/podc/MichaelS96	Singly Linked List	Lock-free	$\Omega(p)$
MoirNSS05 [10] DBLP:conf/spaa/MoirNSS05	Linked List with Elimination		
HoffmanSS07 [3] DBLP:conf/opodis/HoffmanSS07	Linked List with Baskets		$\Omega(p)$
Ladan-MozesS08 [7] DBLP:journals/dc/Ladan-MozesS08	Doubly linked list	Lock-free	$\Omega(p)$
MilmanKLLP18 [9] DBLP:conf/spaa/MilmanKLLP18	Linked List with batching same type ops	Lock-free	
GidenstamST10 [2] DBLP:conf/opodis/GidenstamST10	Linked list of arrays	Lock-free	$\Omega(p + \text{len}(\text{array}))$
KoganP11 [6] DBLP:conf/ppopp/KoganP11	Linkedlist with head and tail pointers	Wait-free	$O(p \times \text{bakery-max})$

Table 1: Linearizable ME-MD Queues

2 Universal Construction using Tournament Tree with Big CAS Objects

Jayanti [\[4\]](#) [DBLP:conf/podc/Jayanti98a](#) proved an $\Omega(\log p)$ lower bound on the worst-case shared-access time complexity of p -process universal constructions. He also introduced [\[1\]](#) [DBLP:conf/podc/ChandraJT98](#) a construction that achieves $O(\log^2 p)$ shared accesses. Here, we first introduce a universal construction using $O(\log p)$ CAS operations [\[5\]](#) [DBLP:conf/fsttcs/JayantiP05](#). We use the universal construction as a stepping stone towards our queue algorithm, so we will not explain it in too much detail.

We introduce our universal construction in Algorithm [alg1](#) with a tournament tree with p leaves and height $\log(p)$ is shared among p processes. Nodes are CAS objects, and each leaf is assigned to one process exclusively. Each CAS object stores a sequence of operations. When process P_i wishes to apply an operation op on the implemented object, it appends op to its assigned leaf and tries to propagate it up to the root. The sequence of operations stored in Node n 's CAS object shows the order of the operations propagated up to n . The history of operations stored at the root is the linearization ordering. The operation op is linearized when it is appended to the root.



Leaf l_i stores the sequence of the operations invoked by P_i . The algorithm uses a subroutine $\text{REFRESH}(n)$ that concatenates new operations from node n 's children (that have not already been propagated to n) to the sequence of operations stored in n and tries to




Figure 1: Propagate Step in Universal Construction

CAS the new sequence into n . In other words, $\text{REFRESH}(n)$ tries to append n 's children's new operations to n 's sequence. After a process adding a new operations to its leaf, it has to propagate new operations up to the root. $\text{PROPAGATE}(n)$ tries to append n 's new operations to the root n by recursively calling $\text{REFRESH}(n)$. In each step if a $\text{REFRESH}(n)$ fails, it means another CAS operation has succeeded; if so, it tries to $\text{REFRESH}(n)$ again. If the second attempt fails too, another process has already appended the operations the current PROPAGATE is trying to append. Operations that were in l_i before $\text{PROPAGATE}(l_i.\text{parent})$ was invoked are guaranteed to be added to the root by the time the $\text{PROPAGATE}(l_i.\text{parent})$ finishes.

alg1

Algorithm Universal Construction Idea

<pre> 1: response Do(operation op, pid i) 2: l_i.APPEND(op) 3: PROPAGATE(parent of l_i) 4: Run the sequence stored in root 5: return op's response from line 4 6: end Do 7: void PROPAGATE(node n) 8: if $n == \text{root}$ then return 9: else if !REFRESH(n) then 10: REFRESH(n) 11: end if 12: PROPAGATE(parent of n) 13: end PROPAGATE </pre>	<pre> 14: boolean REFRESH(node n) 15: old= READ(n) 16: new= ops that n's children contain but old does not 17: new= old.new 18: return n.CAS(old, new) 19: end REFRESH </pre>
---	---

$O(\log n)$ CAS operations are invoked to do a PROPAGATE, but the CAS words store sequences of unbounded length. The problem is that we are trying to store unbounded sequence of operations in each node n (see Figure 2). However, to compute the result of an operation, we only use the total ordering that is stored at the root. Although we use a similar construction for our queue implementation, we develop an implicit representation of the sequence of operations, so that we can use reasonable size CAS objects and still achieve polyarithmic step complexity.

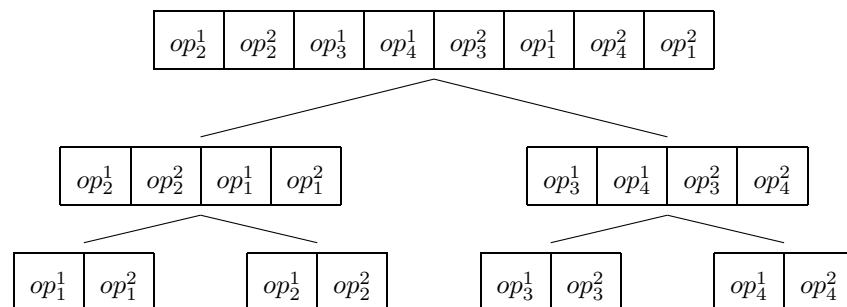


Figure 2: Universal Construction: op_j^i denotes the i th operation from process j . In each node, we store the ordering of all the operations propagated up to it.

3 Block Tree

We apply two ideas from universal construction to create a new linearizable data structure agreeing on a sequence of elements among processes. First, there is a shared tournament tree among processes, in which each process appends its element to its leaf in the tree and then tries to propagate it up to the root by performing `REFRESH()` operations at each node. Second, each operation is linearized when its element is appended to the root.

3.1 Sequence of propagated operations

The basic idea behind the universal construction is to create a linearization of operations invoked by processes. If we design a fast shared data structure in which processes can append elements to a sequence, we can use it to implement practical, fast implementation of an object O by using the sequence data structure to keep track of the sequence of operations on O . In the following sections, we introduce our solution called Block Tree.

3.2 Sequence of Sets of Concurrent Operations

In the universal construction, we order new concurrent operations at each `REFRESH()` and maintain that order in the path up to the root. However, we can instead keep track of sets of concurrent operations and create the total ordering of all operations at the root (see Figure 3).




Figure 3: In each internal node, we store the set of all the operations propagated together, and one can arbitrarily linearize the sets of concurrent operations among themselves. Since we linearize operations when they are added to the root, ordering the blocks in the root is important.

The definition of linearizability allows concurrent operations to be reordered arbitrarily. Thus, a group of concurrent operations can be appended to our root sequence as one block without specifying the order among the operations.

3.3 Using arrays instead of big CAS Objects

We used unbounded CAS objects storing sequences as big words in the universal construction. One can represent sequences as arrays to overcome this implementation problem. Each array element will store one of the blocks of concurrent operations described in section 3.2.

3.4 Augmenting Tree to make Refresh() Step faster

Copying operation sequences from children to their parent in a `REFRESH()` takes time proportioned to the number of operations being copied. This is time-consuming, so we propose a way to augment the tree to calculate lines 15,16 in $O(\log p)$ steps which reads new operation and concatenates them with old operations. Instead of representing the set of operations by explicitly listing them in a node, we represent a set of operations implicitly by recording which of the children's sets were unioned to create the set. Having operation sequences stored at leaves, we can deduce a set of operations in a node using this implicit representation. (see Figure 4.)




fig:block

Figure 4: Showing concurrent operation sets with blocks. Each block consists of a pair(left, right) indicating the number of operations from the left and the right child, respectively. Block (12,6) in the root contains blocks (10,2) from the left child and (6,0) from the right child. Blocks between two lines || are propagated together to the parent. For example, Blocks (2,3) and (4,1) from the leftmost leaf and (0,2) from its sibling are propagated together into the block (10,2) in their parent. The number underneath a group of blocks in a node indicates which block in the node's parent those blocks were propagated to.

Each block b in node n is the aggregation of blocks in the children of n that are newly read by the `PROPAGATE()` step that created block b . For example, the third block in the root (8,36) is created by merging block (5,3) from the left child and (14,6) and (0,16) from the right child. Block (5,3) also points to elements from blocks (0,5) and (1,2).

Definition 1. {Existence of an operation in a block} Operation op exists in block b if it has propagated up to block b .

Definition 2. {Subblock} The blocks that are aggregated into block b in a `PROPAGATE()` step are called subblocks of b . Block b_1 is a subblock of b_2 if and only if b_1 is a block in node v and in b_2 is a block in the parent of v and b_1 's elements exists in b_2 's elements.

We choose to linearize operations in a block from the left child before those from the right child as a convention. Operations within a block of the root can be ordered in any way that is convenient. In effect, this means that if there are concurrent new blocks in a `REFRESH()` step from several processes we linearize them in the order of their process ids. So for example operations aggregated in block (10,2) are in the order (2,3),(4,1),(0,2). All blocks from the left child with come before the right child and the order of blocks of each child is preserved among themselves.

In a `PROPAGATE()` invocation path from a leaf to root, there will be `REFRESH()` steps with merges from 2, 4, 8, ..., p processes. So in a complete propagation, at most $2p$ blocks are merged into one block. (maybe useful for analysis)

3.5 Using pointers and prefix sum to make `GetIndex(i)` faster

`GETINDEX(i)` returns the i th operation stored in the block tree sequence. We do that by finding the block b_i containing i th element in the root, and then recursively finding the subblock of b_i which contains i th element. To make this recursive search faster, instead of iterating over all elements in sequence of blocks we store prefix sum of number of elements in the blocks sequence and pointers to make `BSearch` faster.

Furthermore, in each block, we store the prefix sum of left and right elements. Moreover, for each block, we store two pointers to the last left and right subblock of it (see fig 6 and 5).

fig::pointersprefix


3,8	6,0	14,6	0,16
0,0	3,8	9,8	23,14

0,3	4,2	5,1	3,5
0,0	0,3	4,5	9,6

2,1	5,0	4,2	3,4	6,3
0,0	2,1	7,1	11,3	14,7

fig:prefix

Figure 5: Using Prefix sums in blocks. When we want to find block b elements in its children, we can use binary search. The number below each block shows the count of elements in the previous blocks.



pointer

Figure 6: Block have pointers to the starting block of theirs for each child.

Starting from the root, $\text{GETINDEX}(i)$ BSearches i in the prefix sum array to find block containing i th operation, then continues recursively calling $\text{GETELEMENT}(b, i)$ to find i th element of block b . From lemma [lem:block_size](#) we know a block size is at most p . So BSearch takes at most $(O)(\log p)$, since with knowing pointers of a block and its previous block we can determine the base (domain ?) to search and its size is $O(p)$.

3.6 Block Tree Algorithm

Our Block Tree is a linearizable implementation of a data structure that stores a sequence of elements. It has two methods (see Algorithm [alg2](#)), $\text{APPEND}(e)$ which appends element e to the sequence, and $\text{GET}(i)$ which returns the i th element in the sequence.

Design of a block tree Each process is assigned to a leaf in a shared tournament tree. Thus, for example, the leaf node for process p_i contains an array of elements by p_i in the order they were invoked. Each internal node of the tree contains an array of blocks of elements. Block b in node n is created in a $\text{PROPAGATE}()$ step and is merged block of new blocks at the time of $\text{PROPAGATE}()$ reading n 's children blocks. Each block consists of pointers left and right, to the last block merged into itself from left and right child in that order. Moreover, two numbers, left and right, indicate the count of elements in the blocks from the left and right child consecutively. Furthermore, prefix left, and right can be computed from the prefix sum of left and right values. Elements of block b can be determined recursively ($\text{GETELEMENTS}(b)$). The i th element in the sequence can be determined in $O(\log^2 p)$ steps by recursively finding i th element in block b ($\text{GETELEMENT}(i)$). After element e is propagated (appended to a block into the root), its index can be computed with $\text{GETINDEX}(op)$.

In order to compute elements of a block faster we store prefix-sum blocks (block i has tuple (right-sum=#right ops in previous block, left-sum=#left ops in previous blocks) [See Figure [fig:prefix](#)]). Here is the algorithm to get elements of a block.

Specification A block tree is a shared data structure that stores a sequence of elements. It has two methods **Append(e)** and **Get(i)**. **Append(e)** adds **e** to the end of the sequence and returns the index of **e** in the sequence. **Get(i)** returns *i*th element stored in the sequence.

SubBlock Block **s** is a subblock of **b** if **s** is between blocks **start..end** in **n** from Lines 41,42 of **CreateBlock()**.

Membership Element **e** is a member of block **b** in:

- internal node **n**, if **e** is a member of **s** that **s** is a subblock of **b**.
- leaf node **n**, if **e** belongs to $n.dir.blocks[b'.end_{dir}+1..b.end_{dir}]$ for $dir \in \{left, right\}$ which **b'** is the previous block of **b** in **n**.

Order of elements inside node Element **d** is before element **e** in node **n**, if:

- The block containing **d** is before the block containing **e**.
- **e** and **d** are in the same block and **d** is in the left child and **e** is in the right child.
- **d** is before **e** in the same child's order.

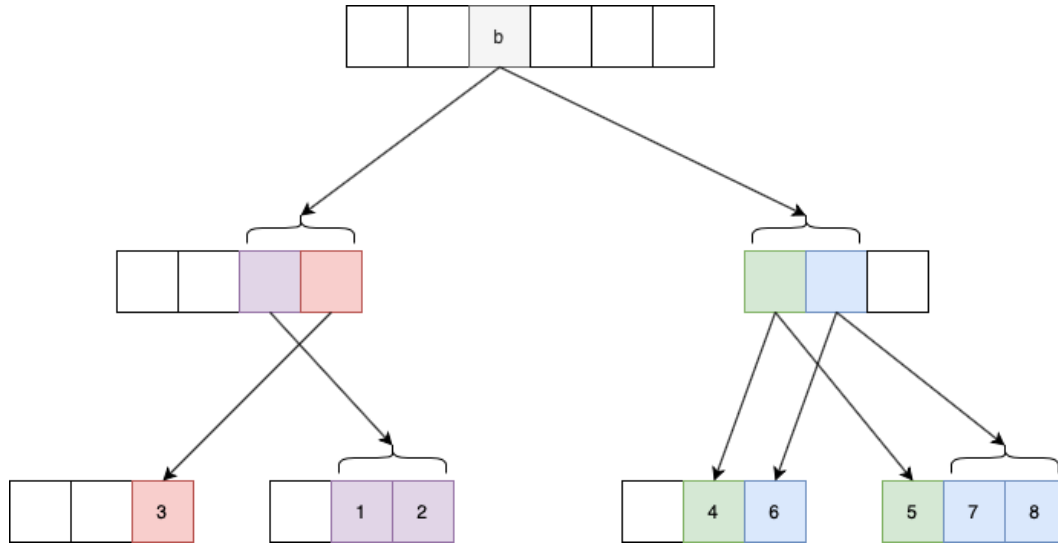


Figure 7: Order of elements in **b**: elements in leaves are ordered with numerical order in the drawing.

CreateBlock() **CreateBlock(n)** returns a block containing new operations of **n**'s children. $b'.end_{left}$ stores the index of the rightmost subblock of left child of **b**'s previous block. Other attributes are assigned values followed by definition.

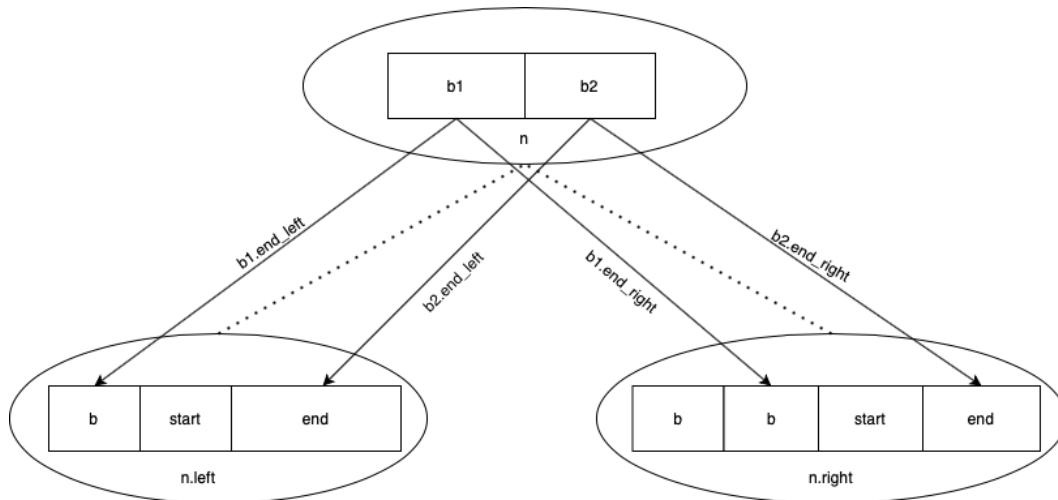


Figure 8: Snapshot of a **CreateBlock()**

Double Refresh Elements in n 's children's blocks before line 13 are guaranteed to be in n 's blocks after Line 15.

Proof. `CreateBlock()` reads blocks in the children that does not exist in the parent and aggregates them into one block. If a `Refresh` procedure returns true it means it has appended the block created by `CreateBlock()` into the parent node's sequence. So suppose two `Refreshes` fail. Since the first `Refresh` was not successful, it means another CAS operation by a `Refresh`, concurrent to the first `Refresh`, was successful before the second `Refresh`. So it means the second failed `Refresh` is concurrent with a successful `Refresh` that assuredly has read block before the mentioned line 13. After all it means if any of the `Refresh` attempts were successful the claim is true, and also if both fail the mentioned claim still holds. \square

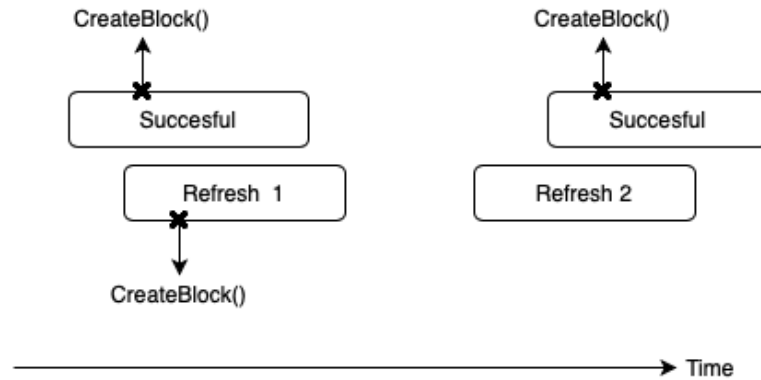


Figure 9: The second failed `Refresh` is assuredly concurrent to a Successful `Refresh()` with `CreateBlock` line after first failed `Refresh`'s `CreateBlock()`.

Disjunction Blocks in node n 's contain disjoint sets of elements.

Proof. Without loss of generality, assume blocks b_1 , b_2 contain common element e from the left child, and b_2 is after b_1 in n 's sequence of blocks. So block start of b_2 's `CreateBlock()` is after block end of b_1 's end. Since b_2 's start is the end of the block before itself, it cannot be before b_1 's end. \square

Total Order Sequence represented by the Block Tree is the sequence of the blocks stored in the root.

Linearization Points `Get(i)` is linearized when it terminates. `Append(e)` is linearized right after when a block containing e is appended to the root, if there are multiple elements appended together, they are linearized by the defined order in the root.

Subblocks Upperbound Block b has at most p subblocks.

Proof. If there are more than p subblocks, then there is more than one block from process p_l . `Append(e)` finishes after propagating and appending e to the root(line 9). So these blocks cannot be appended to root already, so p_l has invoked two concurrent `Append()`s(line 1) without terminating the first one. \square

Computing `Get(n, b, i)` To find the i th element in block b of node n , we search among subblocks of b that is bounded by p . Subblocks of a block are within the start and end block of the `CreateBlock()` procedure of it.

How `Refresh(n)` works.

1. Read n 's counter and head
2. Create block b
3. CAS b into n

4. If previous succeed:

- (a) Update sup of b's ending subblocks
- (b) Increment children's counters




Figure 10: Sup and timer in a node, numbers on blocks are their time values.

Computing superblock

1. Value read for `super[b.time]` in line 71 is not null.

Proof. `Index()` is invoked after finishing `Propagate()` in line 10. For each value `c_dir` read in lines 23, `super` is set before incrementing in lines 26,27. □

2. `super[]` preserves order from child to parent; if in a child block `b` is before `c` then `b.time ≤ c.time` and `super[b.time] ≤ super[c.time]`

Proof. Follows from the order of lines 37, 26, 27. □

3. `super[i+1]-super[i] ≤ p`

Proof. In a Refresh with successful CAS in line 24, `super` and `counter` are set for each child in lines 26,27. Assume the current value of the counter in node `n` is `i+1` and still `super[i+1]` is not set. If an instance of successful `Refresh(n)` finishes `super[i+1]` is set a new value and a block is added after `n.parent[sup[i]]`. There could be at most p successful unfinished concurrent instances of `Refresh()` that have not reached line 27. So the distance between `super[i+1]` and `super[i]` is less than p . □

4. Superblock of `b` is within range $\pm 2p$ of the `super[b.time]`.

Proof. `super[i]` is the index of the superblock of a block containing block `b`. It is trivial to see that `n.super` and `n.b.counter` are increasing. `super(b)` is the real superblock of `b`. `super(t)` is the index of the superblock of the last block with time `t`. If `b.time` is `t` we have:

$$super[t] - p \leq super[t-1] \leq super(t-1) \leq super(b) \leq super(t+1) \leq super(t+1) \leq super[t] + p$$

□

Algorithm Block Tree

Structure► *element* *e*

- *int* *pid*
- *int* *loc*: location in *l_{pid}.ops*

► *node* *n*

- **node* *left*, *right*, *parent*
- *block[]* *blocks*: blocks stored in *n*
- *int* *head*= 1: index of first empty cell of blocks
- *int* *counter*= 0
- *int[]* *super*: *super[i]* stores the index of a superblock in parent that contains some block of this node whose time is *i*

► *leaf node* *l_i* extends *node*

- *operation[]* *ops*: elements that are invoked to *Append()* to the block tree by process *i*

► *block* *b*

- *int* *num_{left}*, *sum_{left}*
#operations from the left subblocks of *b*, prefix sum of *num_{left}*
- *int* *num_{right}*, *sum_{right}*
#operations from the right subblocks of *b*, prefix sum of *num_{right}*
- *int* *sum*
operation in *b*
- *int* *end_{left}*, *end_{right}*
index of *b*'s last subblock
- *int* *time*

```

1: int APPEND(operation op, int pid)
2:   op.loc= lpid.head
3:   block b= NEW(block)
4:   b.time= op.loc
5:   b.sum= 1
6:   lpid.ops[op.loc]= op
7:   lpid.blocks[op.loc]= b
8:   op.head+= 1
9:   PROPAGATE(lpid.parent)
10:  return INDEX(lpid, op.loc, b)
11: end APPEND

```

```

12: void PROPAGATE(node n)

```

```

13:   if not REFRESH(n) then

```

```

14:     REFRESH(n)

```

```

15:   end if

```

```

16:   if n is root then return

```

```

17:   end if

```

```

18:   PROPAGATE(n.parent)

```

```

19: end PROPAGATE

```

```

20: boolean REFRESH(node n)

```

```

21:   i= n.head

```

```

22:   c= n.counter

```

```

23:   new, cleft, cright= CREATEBLOCK(n, i+1, c)

```

```

24:   if CAS(n.blocks[i+1], null, new) then

```

```

25:     for each dir in {left, right} do

```

```

26:         CAS(n.dir.super[cdir], null, i+1)

```

```

27:         CAS(n.dir.counter, cdir, cdir+1)

```

```

28:     end for

```

```

29:     return true

```

```

30:   else

```

```

31:     return false

```

```

32:   end if

```

```

33:   CAS(n.head, i, i+1)

```

```

34: end INDEX

```

```

35: block, int, int CREATEBLOCK(node n, int i, int c)

```

▷ Creates a block to insert into *n.blocks[i]*

```

36:   block b= NEW(block)

```

```

37:   b.time= c

```

```

38:   for each dir in {left, right} do

```

```

39:     j= n.dir.head

```

```

40:     start= n.blocks[i-1].enddir

```

```

41:     end= n.dir.blocks[j]

```

```

42:     cdir= n.dir.blocks[j].time

```

```

43:     b.enddir= j

```

```

44:     b.numdir= end.sum - start.sum

```

```

45:     b.sum#dir= n.blocks[i-1].sum#dir + b.numdir

```

```

46:   end for

```

```

47:   b.sum= b.sumleft + b.sumright

```

```

48:   return b, cleft, cright

```

```

49: end CREATEBLOCK

```

Algorithm Block Tree Continued

47: *element* GET(*int* *i*)

48: *res* = BSEARCH(*root*, *sum*, *i*, 0, *root*.*head*)

49: **return** GET(*root*, *res*, *i* - *root*.*blocks*[*res* - 1].*sum*)

50: **end** GET

↪ Precondition: *n*.*blocks*[*start*..*end*] contains a block with field *f* $\geq i$

51: *int* BSEARCH(*node* *n*, *field* *f*, *int* *i*, *int* *start*, *int* *end*)

▷ Searches the value *i* of the given prefix sum *type* on the node *n* in the domain from *start* to *end* blocks. *f* is one of *sum*, *sum*_{left}, *sum*_{right}. Does binary search on *field* *f* of the blocks stored in *n*.*blocks* and returns the index of the leftmost block in *n*.*blocks*[*start*..*end*] whose *field* *f* is $\geq i$.

52: **return** result block's index

53: **end** BSEARCH

54: *element* GET(*node* *n*, *int* *b*, *int* *i*)

▷ Returns the *i*th operation in *b*th block of node *n*

55: **if** *n* is leaf **then** **return** *n*.*ops*[*i*]

56: **else**

57: **if** $i \leq n$.*blocks*[*b*].*num*_{left} **then**

▷ *i* exists in left child of *n*

58: *subBlock* = BSEARCH(*n*.*left*.*sum*, *i*, *n*.*blocks*[*b* - 1].*end*_{left} + 1, *n*.*blocks*[*b*].*end*_{left})

59: **return** GET(*n*.*left*, *subBlock*, *i* - *n*.*left*.*blocks*[*subBlock* - 1].*sum*)

60: **else**

61: *i* = *i* - *n*.*blocks*[*b*].*num*_{left}

62: *subBlock* = BSEARCH(*n*.*right*.*sum*, *i*, *n*.*blocks*[*b* - 1].*end*_{right}, *n*.*blocks*[*b*].*end*_{right})

63: **return** GET(*n*.*right*, *subBlock*, *i* - *n*.*right*.*blocks*[*subBlock* - 1].*sum*)

64: **end if**

65: **end if**

66: **end** GET

↪ Precondition: *i*th operation in node *n* is in block *b* of node *n*.

67: *index* INDEX(*node* *n*, *int* *i*, *int* *b*)

▷ Returns rank in the root of the *i*th operation in node *n*.

68: **if** *n* is root **then** **return** *i*

69: **else**

70: $dir = (n.parent.left == n) ? left : right$

71: *superBlock* = BSEARCH(*n*.*parent*, *n*.*sum*_{dir}, *i*, $sup[n.blocks[b].time] - p$, $sup[n.blocks[b].time] + p$)

72: **if** *dir* is left **then**

73: *i* += *n*.*parent*.*blocks*[*superBlock* - 1].*sum*_{right}

74: **else**

75: *i* += *n*.*parent*.*blocks*[*superBlock*].*sum*_{left} + *n*.*parent*.*blocks*[*superBlock*].*sum*_{left} + *n*.*blocks*[*nBlock* - 1].*sum*

76: **end if**

77: **return** INDEX(*n*.*parent*, *i*, *superBlock*)

78: **end if**

79: **end** INDEX

4 Implementing Queue using Block Tree

4.1 Idea in a nutshell

With the ideas introduced in block tree we are going to create a shared wait-free queue with $O(\log^2 p + \log n)$ steps. A queue stores a sequence of elements and supports two operations, enqueue and dequeue. **Enqueue**(*e*) appends element *e* to the sequence stored. **Dequeue**() removes and returns the first element among in the sequence. If the queue is empty it returns **null**. Knowing index *i* is the tail of the queue, we can return the dequeue response using **Get**(*i*). So in the rest we modify block tree to compute *i* for each **Dequeue**() to achieve a FIFO queue.

Next, we describe how to use block tree to implement queues. The block tree, maintains the history of all operations, not only the current state of the queue. Now consider the following history of operations. What should each **Dequeue**() return?

ENQ(5)	ENQ(2)	DEQ()	ENQ(3)	DEQ()	DEQ()	DEQ()	ENQ(4)	ENQ(6)	DEQ()
--------	--------	-------	--------	-------	-------	-------	--------	--------	-------

Table 2: An example histoy of operations on the queue

Definition 3. A non-null dequeue is one that returns a non-null value.

In the example above, **Dequeue**() operations return 5, 2, 3, **null**, 4 in order. Before **ENQ(4)** the queue gets empty so the last **DEQ()** returns **null**. If the queue is non-empty and *r* **Dequeue**() operations have returned a non-null response, then *i*th **Dequeue**() returns the input of the *r* + 1th **Enqueue**(). So, in order to answer a Dequeue, it's sufficent to know the size of the queue and the number of previous non-null dequeues.

In the Block Tree, we did not store the sequence of operations explicitly but instead stored blocks of concurrent operations to optimize **Propagate**() steps and increase parallelism. So now the problem is to find the result of each Dequeue. From lemma ^{lem:block_size} we know we can linearize operations in a block in any order; here, we choose to decide to put Enqueue operations in a block before Dequeue operations. In the next example, operations in a cell are concurrent. **DEQ()** operations return **null**, 5, 2, 1, 3, 4, **null** respectively. We will next describe how these values can be computed efficiently.

DEQ()	ENQ(5), ENQ(2), ENQ(1), DEQ()	ENQ(3), DEQ()	ENQ(4), DEQ(), DEQ(), DEQ(), DEQ()
-------	-------------------------------	---------------	------------------------------------

Table 3: An example history of operation blocks on the queue

Now, we claimed that by knowing the current size of the queue and the number of non-null dequeue operations before the current dequeue, we could compute the index of the resulting **Enqueue**(). We apply this approach to blocks; if we store the size of the queue after each block of operations happens and the number of non-null dequeues dequeues till a block, we can compute each dequeue's index of result in $O(1)$ steps.

	DEQ()	ENQ(5), ENQ(2), ENQ(1), DEQ()	ENQ(3), DEQ()	ENQ(4), DEQ(), DEQ(), DEQ(), DEQ()
#enqueues	0	3	1	1
#dequeues	1	1	1	4
#non-null dequeues	0	1	2	5
size	0	2	2	0

Table 4: Augmented history of operation blocks on the queue

Size and the number of non-null dequeues for b th block could be computed this way:

$size[b] = \max(size[b-1] + enqueuees[b] - dequeues[b], 0)$

$non_null_dequeues[b] = non_null_dequeues[b-1] + dequeues[b] - size[b-1] - enqueuees[b]$

Given DEQ is in block b , $response(DEQ)$ would be:

$(size[b-1] - index\ of\ DEQ\ in\ the\ block's\ dequeues \geq 0) ? ENQ[non_null_dequeues[b-1] + index\ of\ DEQ\ in\ the\ block's\ dequeues]$
 $: null;$

5 Main Algorithm

Specification A Queue is a shared data structure that stores a sequence of elements. It has two methods **Enqueue(e)** and **Dequeue()**. **Enqueue(e)** adds e to the end of the sequence. **Dequeue()** returns the first element stored in the sequence and removes it from the sequence.




Figure 11: Fields stored in the Queue nodes.

fig::queue

5.1 Pseudocode description

Tournament Tree In order to reach an agreement on the order of operations among p processes, we use a Tournament Tree. Leaf l_i is assigned to a process i . Each process adds op to its leaf. In each internal node an ordering of operations in its subtree is stored. All processes agree on the total ordering of all operations stored in the root. This ordering will be the linearization of the operations.

Implicit Storing Blocks For efficiency, instead of storing explicit sequence of operations in nodes of the Tournament Tree, we use Blocks. A Block is a constant size object that implicitly represents a sequence of operations. In each node there is an array of Blocks.

Definition 4 (Block). A block is an object that stores some statistics described in Algorithm Queue.

Definition 5 (Subblock). Block b is a subblock of $n.blocks[i]$ if it is in $n.left.blocks[n.blocks[i-1].end_{left}+1..n.blocks[i].end_{left}]$ or $n.right.blocks[n.blocks[i-1].end_{right}+1..n.blocks[i].end_{right}]$.

Block b contains subblocks in the left and right children. WLOG left subblocks of b are some consecutive blocks in the left child starting from where previous block of b has ended to the the end of b . See Figure [8](#). [fig::createBlock](#)

Definition 6 (Membership of an operation in a block). Operation e is a member of block b in:

- leaf node n , if e belongs to $n.ops[b's\ index]$.

► Operation

- *Object* element: argument of an enqueue, if null this is a dequeue
- *int* loc: location in the leaf's ops

► Node

- **Node* left, right, parent
- *Block[]* blocks: index 0 contains an empty block with all fields equal to 0
- *int* head= 1: index of the first empty cell of blocks
- *int* counter= 0
- *int[]* super: super[i] stores the index of a superblock in parent that contains some block of this node whose time is field i

► Leaf extends Node

▷ l_i is the Leaf for process i and i is in range 0 to p-1

- *Operation[]* ops: invoked operations

► Root extends Node

- override *Root Block[]* blocks

► Block

- *int* num_{enq-left}, sum_{enq-left}
#enqueues from subblocks in left child, prefix sum of num_{enq-left}
- *int* num_{deq-left}, sum_{deq-left}
#dequeues from subblocks in left child, prefix sum of num_{deq-left}
- *int* num_{enq-right}, sum_{enq-right}
#enqueues from subblocks in right child, prefix sum of num_{enq-right}
- *int* num_{deq-right}, sum_{deq-right}
#dequeues from subblocks in right child, prefix sum of num_{deq-right}
- *int* num_{enq}, num_{deq}
enqueue, dequeue operations in the block
- *int* sum_{enq}, sum_{deq}
sum of # enqueue, dequeue operations in blocks up to this one
- *int* num, sum
total # operations in block, prefix sum of num
- *int* end_{left}, end_{right}
index of this block's last subblock in the left and right child
- *int* group

► Root Block extends Block

- *int* size
size of queue after this block's operations finish

- *int* sum_{non-null deq}

count of non-null dequeues up to this block

```

1: void ENQUEUE(Object e)
2:   op= NEW(operation)
3:   op.element= e
4:   APPEND(op)
5: end ENQUEUE

6: Object DEQUEUE()
7:   op= NEW(operation)
8:   APPEND(op)                                ▷ ith DEQ in bth block
9:   <i, b>= INDEX(lpid, op.loc, 1)
10:  res= COMPUTEHEAD(i, b)                    ▷ Index of the enqueue whose argument
                                             should be returned
11:  if res is null then
12:    return null
13:  else
14:    return GET(res)
15:  end if
16: end DEQUEUE

17: int COMPUTEHEAD(int i, int b)              ▷ Computes head of the queue when
                                             ith dequeue in bth block occurs. The dequeue should return the argument
                                             of the head enqueue.
18:  if root.blocks[b-1].size + root.blocks[b].numenq - i < 0 then
19:    return -1
20:  else return root.blocks[b-1].sumnon-null deq + i
21:  end if
22: end COMPUTEHEAD

23: void APPEND(operation op)
24:  pid= id of process performing APPEND
25:  op.loc= lpid.head
26:  op.head+= 1
27:  block b= NEW(block)
28:  b.group= op.loc
29:  if op.element==null then b.sumdeq=1
30:  else b.sumenq=1
31:  end if
32:  lpid.ops[op.loc]= op
33:  lpid.blocks[op.loc]= b
34:  PROPAGATE(lpid.parent)
35: end APPEND

```

addOP

Algorithm Queue Continued

```

34: void PROPAGATE(node n)
35:   if not REFRESH(n) then
36:     REFRESH(n)
37:   end if
38:   if n.parent is null then
39:     PROPAGATE(n.parent)
40:   end if
41: end PROPAGATE

42: boolean REFRESH(node n)
43:   h= n.head
44:   t= n.counter
45:   <new, tleft, tright>= CREATEBLOCK(n, h, t)
46:   if new.num==0 then return true
47:   else if CAS(n.blocks[h], null, new) then
48:     for each dir in {left, right} do
49:       CAS(n.dir.super[tdir], null, h+1)
50:       CAS(n.dir.counter, tdir, cdir+1)
51:     end for
52:     CAS(n.head, h, h+1)
53:     return true
54:   else
55:     CAS(n.head, h, h+1)
56:     return false
57:   end if
58: end REFRESH

59: element GET(int i)                                ▷ Returns ith Enqueue.
60:   res= BSEARCH(root, sumenq, i, 0, root.head)
61:   return GET(root, res, i-root.blocks[res-1].sumenq)
62: end GET

    ~▷ Precondition: n.blocks[start..end] contains a block with field f ≥ i

63: int BSEARCH(node n, field f, int i, int start, int end)
    ▷ Does binary search for the value i of the given
    prefix sum feild. f is one of sum, sumleft, sumright. Returns the index of
    the leftmost block in n.blocks[start..end] whose field f is ≥ i.

64: end BSEARCH

```

```

65: <Block, int, int> CREATEBLOCK(node n, int i, int t)
    ▷ Creates a block to insert into n.blocks[i] with time field=t. Returns
    the created block as well as values read from each child counter feild.

66:   block b= NEW(block)
67:   b.group= t
68:   for each dir in {left, right} do
69:     lastLine69:   lastIndex= n.dir.head
70:     prevLine70:   prevIndex= n.blocks[i-1].enddir
71:     lastBlock= n.dir.blocks[lastIndex]
72:     prevBlock= n.dir.blocks[prevIndex]
73:     tdir= n.dir.counter
74:     b.enddir= lastIndex
75:     b.numenq-dir= lastBlock.sumenq - prevBlock.sumenq
76:     b.numdeq-dir= lastBlock.sumdeq - prevBlock.sumdeq
77:     b.sumenq-dir= n.blocks[i-1].sumenq-dir + b.numenq-dir
78:     b.sumdeq-dir= n.blocks[i-1].sumdeq-dir + b.numdeq-dir
79:   end for
80:   b.numenq= b.numenq-left + b.numenq-right
81:   b.numdeq= b.numdeq-left + b.numdeq-right
82:   b.num= b.numenq + b.numdeq
83:   b.sum= n.blocks[i-1].sum + b.num
84:   if n.parent is null then
85:     b.size= max(root.blocks[i-1].size + b.numenq - b.numdeq, 0)
86:     b.sumnon-null deq= root.blocks[i-1].sumnon-null deq + max(
      b.numdeq - root.blocks[i-1].size - b.numenq, 0)
87:   end if
88:   return b, tleft, tright

89: end CREATEBLOCK

```

Algorithm Queue Continued

↪ Precondition: $n.blocks[b]$ contains $\geq i$ enqueues.

```
84: element GET(node n, int b, int i)                                ▷ Returns the ith Enqueue in bth block of node n
85:   if n is leaf then return n.ops[b]
86:   else
87:     if  $i \leq n.blocks[b].num_{enq-left}$  then                                ▷ i exists in left child of n
88:       subBlock= BSEARCH(n.left, sumenq, i, n.blocks[b-1].endleft+1, n.blocks[b].endleft)
89:       return GET(n.left, subBlock, i-n.left.blocks[subBlock-1].sumenq)
90:     else
91:       i= i-n.blocks[b].numenq-left
92:       subBlock=BSEARCH(n.right, sumenq, i, n.blocks[b-1].endright+1, n.blocks[b].endright)
93:       return GET(n.right, subBlock, i-n.right.blocks[subBlock-1].sumenq)
94:     end if
95:   end if
96: end GET
```

↪ Precondition: *b*th block of node *n* has propagated up to the root and *i*th dequeue resides in node *n* is in block *b* of node *n*.

```
97: <int, int> INDEX(node n, int b, int i)                            ▷ Returns the order in the root among dequeues, of ith dequeue in bth block of node n
98:   if n is root then return root.blocks[b-1]+i, b
99:   else
100:    dir= (n.parent.left==n)? left: right
101:    superBlock= BSEARCH(n.parent, n.sumdeq-dir, i, super[n.blocks[b].group]-p, super[n.blocks[b].group]+p)
102:    if dir is left then
103:      i+= n.parent.blocks[superBlock-1].sumdeq-right
104:    else
105:      i+= n.parent.blocks[superBlock-1].sumdeq + n.blocks[superBlock].sumdeq-left
106:    end if
107:    return INDEX(n.parent, superBlock, i)
108:  end if
109: end INDEX
```

- internal node n , if e is a member of s that s is a subblock of b .

We store ordering among **operations** in the tournament tree constructed by **nodes**. In each **node** we store pointers to its relatives, an array of **blocks** and an index to the first empty **block**. Furthermore in **leaf** nodes there is an array of **operations** where each **operation** is stored in one cell with the same index in **blocks**. There is a **counter** in each **node** incrementing after a successful **Refresh()** step. It means after that some bunch of **blocks** in a node have propagated into the parent then the **counter** increases. Each new **block** added to a node sets its **time** regarding **counter**. This helps us to know which blocks have aggregated together to a block, not precisely though. We also store the index of the aggregated **block** of a **block** with **time** i in **super**[i].

In each **block** we store 4 essential stats that implicitly summarize which operations are in the block **num_{enq-left}**, **num_{deq-left}**, **num_{enq-right}**, **num_{deq-right}**. In order to make **BSearch()**es faster we store prefix sums as well and there are some more general stats that help to make pseudocode more readable but not necessary.

To compute the head of the **queue** before a **dequeue** two more fields are stored in the root **size** and **sum_{non-null deq}**. **size** in a **block** shows the number of elements after the **block** has finished and **sum_{non-null deq}** is the total number of non-null dequeues till the **block**.

Enqueue(e) just appends an operation with element e to the root. **Dequeue()** appends an operation to the root and computes its ordering and the **enqueue** operation containing the head before it calling **ComputeHead()** and then **gets** and returns the operation's element.

Append(op) adds op to the invoking process's leaf's **ops** and **blocks**, propagates it up to the root and if the op is a dequeue returns its order in residing block in the root and the block's index. As we said later **Propagate()** assuredly aggregates new blocks to a block in the parent by calling **Refresh()** two times. **Refresh(n)** creates a block, tries to CAS it into the pn 's **blocks** and if it was successful updates **super** and **counter** in both of n 's children.

We only want to know the **element** of **enqueue** operations and compute ordering for **dequeue** operations. That's the reason here **Get()** searches between enqueues only and **Index()** returns ordering of a dequeue among dequeues. **Get(n, b, i)** decides the requested element is in which child of n and continues to search recursively. **index(n, i, b)** calculates the ordering of the given operation in n 's parent each step and finally returns the result among total ordering.

5.2 Complexity Analysis

Enqueue() operations do a constant number of steps and an **Append()** which calls **Propagate()**. Let p be the number of processes and m be the count of operations invoked. In a **Propagate()** step there are a constant number of steps taken at each level and the height of the tree is $\log(p)$. So **Enqueue()** takes $O(\log p)$ steps.

Get() searches among all blocks at first and then iterates over a path from the root to a leaf of the tree and in each step searches in a domain of size p , so it takes $\log^2 p + \log m$. **Index()** is a path from a leaf to the root which each step calls a $\log p$ step search so it takes $\log^2 p$. **Dequeue()** calls **Append()**, it also calls **Index()** and **Get()** that take $\log^2 p$, $\log^2 p + \log n$.

5.3 Linearizability Proof

Definition 7. If $n.blocks[i] == b$ we call i the *index* of block b in node n . Block b is before block b' in node n if and only if b 's index is smaller than b' 's. Block b is propagated to node n or set S if b is in $n.blocks$ or S or is a subblock of a block in $n.blocks$ or S .

Definition 8. Block b in node n is in *Established*(n, t) if $n.head$ is greater than b 's index at time t .

Lemma 9 (headProgress). $n.head$ is non-decreasing over time.

Proof. Simply because $n.head$ is only incremented. □

Lemma 10 (headPosition). The value read in Line 52($h = n.head$) might be 1 bit behind the first empty block.

Proof. Because at the end of every **Refresh()** with block size greater than 0 (Lines 53,56) **n.head** is incremented. Maybe some process goes to sleep before incrementing the head, but after sleeping if **h** does not increase then CAS in Line 52 is going to be failed and nothing is going to be appended to **n.blocks**. \square

establishOrder **Lemma 11** (establishedOrder). *If time $t < \text{time } t'$, then $\text{Established}(n, t) \subseteq \text{Established}(n, t')$.*

Proof. Because blocks are only appended(not modified) with CAS to **n.blocks[n.head]** and **n.head** is non-decreasing. \square

createBlock **Lemma 12** (createBlock). *Suppose $\text{CreateBlock}(n, h, x)$ is invoked at time t . The blocks propagated to $\text{Established}(n.\text{left}, t)$ and $\text{Established}(n.\text{right}, t)$ that are not propagated to $\text{Established}(n, t)$, are subblock of the block returned by $\text{CreateBlock}(n, h, x)$.*

Proof. We prove the claim for the left child. Blocks in **n.left.blocks[n.blocks[i-1].end_{left}+1..n.blocks[i].end_{left}]** are all the new established operations at time t by definition of Subblock. Line 70 is after t and since the head is only increasing (Lemma [lem::headProgress](#)) the lemma holds. See Figure [fig::createBlock](#). The right child is the same. \square

trueRefresh **Lemma 13** (trueRefresh). *Suppose $\text{Refresh}(n)$'s $\text{CAS}(n.\text{blocks}[h], \text{null}, \text{new})$ returns **true**. Let t be the time $\text{Refresh}(n)$ is invoked, blocks propagated to $\text{Established}(n.\text{left}, t)$ and $\text{Established}(n.\text{right}, t)$ are propagated to in $\text{Established}(n, t)$ after $\text{CAS}(n.\text{blocks}[h], \text{null}, \text{new})$.*

Proof. By Lemma [lem::createBlock](#) **new** contains **n**'s childrens' established blocks before Line 43 which is appended to **n.blocks** by CAS in Line 48. \square

falseRefresh **Lemma 14** (falseRefresh). *If instance r of $\text{Refresh}(n)$ returns **false**, then there is another successful instance r' of $\text{Refresh}(n)$ that has performed a successful $\text{CAS}(n.\text{blocks}[h], \text{null}, \text{new})$ (Line 49) after Line 43($h = n.\text{head}$) of r .*

Proof. If there is no other concurrent successful $\text{Refresh}(n)$ then $\text{Refresh}(n)$ would succeed in Line 48. So there is another $\text{Refresh}(n)$, that has to CASed successfully its block in **n.blocks[h]** after Line 43 of $\text{Refresh}(n)$. Otherwise the other $\text{Refresh}(n)$ should have read $h' > h$ instead of h for **n.head** (Line 52). \square

leRefresh **Lemma 15** (Double Refresh). *Consider two consecutive instances r_1, r_2 of $\text{Refresh}(n)$ by the same process (Lines 35,36). Let t be the time before r_1 invoked. After r_2 's CAS all the blocks propagated to $\text{Established}(n.\text{left}, t)$ and $\text{Established}(n.\text{right}, t)$ are in $\text{Established}(n, t)$.*

Proof.

If Line 35 (first Refresh which we call R_1) returns **true**, the claim is held by Lemma [lem::trueRefresh](#) 13. If not, then there is another successful instance of $\text{Refresh}()$ R'_1 by Lemma [lem::falseRefresh](#) 14. R'_1 may or may not have propagated some subblocks of **new** to **n**. It is obvious that the **new** constructed by the second Refresh in Line 36 contains the blocks in **new** by R_1 which R'_1 did not contain, since **n.head** is only increasing (Maybe Lemma [lem::oldnewOrder](#) ??). If R_2 succeeds by Lemma [lem::trueRefresh](#) 13 the claim holds. If not, it is deduced that **n.blocks[h]** was not **null** before R_2 's CAS. Furthermore, **n.blocks[h]** was **null** before reading **h** by R_1 . So there is a successful $\text{Refresh}()$ after the read of **h** in R_1 and before the CAS of R_2 . This $\text{Refresh}()$ contains all the new established operations before Line 35 (Maybe Lemma [lem::oldnewOrder](#) ??) and by Lemma [lem::trueRefresh](#) 13 our claim holds. See Figure [fig::doubleRefresh](#). \square

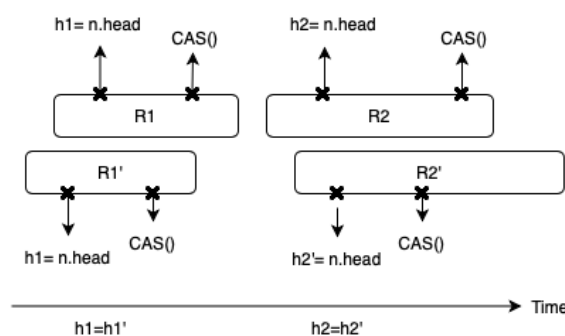


Figure 12: R'_2 's CAS is executed after $h1 = n.\text{head}$.

Block **new** is created of new established subblocks of children of **n** (Lemma ^{lem::createBlock} 12, Line 46). If CAS in Line 48 succeeds then by Lemma ^{lem::trueRefresh} 13 new established blocks will be in **n**.

lyRefresh

Lemma 16 (Double Refresh). *All operations in **n**'s children's blocks before line 35 are guaranteed to be in **n**'s blocks after Line 37.*

Proof. Suppose block **b** with index **i** is in the the left child of **n** before the line 35. By Lemma ^{head} 26 it follows that **n.left.head** is greater than **i**. **Refresh()** calls **CreateBlock()** and creates a block from blocks between **n.blocks[n.head].end_{left}** and **n.left.head** in the left child, which contains **b** as well. First it tries to append it in **n.blocks.head** and if it was succsuful it continues recursivley. If not it tries again, and if the second call of **Refresh()** in Line 36 fails. It means there is another **Refresh** which has reah its **i** after the Line 35, so it contains **b** as well. \square

CreateBlock() reads blocks in the children that do not exist in the parent and aggregates them into one block. If a **Refresh()** procedure returns true it means it has appended the block created by **CreateBlock()** into the parent node's sequence. So suppose two **Refreshes** fail. Since the first **Refresh()** was not successful, it means another CAS operation by a **Refresh**, concurrent to the first **Refresh()**, was successful before the second **Refresh()**. So it means the second failed **Refresh** is concurrent with a successful **Refresh()** that assuredly has read block before the mentioned line 35. After all it means if any of the **Refresh()** attempts were successful the claim is true, and also if both fail the mentioned claim still holds.

append

Lemma 17 (Append). *When **Append(op)** is finished, **op** appears exactly once in a block of **root.blocks**.*

Proof. **Append(op)** adds **op** to **l_p.blocks** (Line ^{addOP} 32) and **Propagate()** recursively propagates **op** up to the root. By lemma ^{doublyRefresh} 16 we know that operation **op** propagates from child to parent at each level. \square

blockSize

Lemma 18 (Block Size Upper Bound). *Each block in a node contains at most one operation from each process.*

Proof. Note that **prevIndex** and **lastIndex** defined in lines ^{lastIndexLine} 69, 70 are the indices defined in the Definition ^{def::subblock} 5. After a block has propagated to the parent the blocks between **prevIndex** and **lastIndex** make up to the parent (Lemma ^{doublyRefresh} 16). The number of new operations which have not propagated yet to the parent cannot be more than **p**. If so by the law of pigeonholes there is a process which has appended two concurrent operations. \square

ocksBound

Lemma 19 (Subblocks Upperbound). *Each block in a node has at most **p** subblocks.*

Proof. From Line ^{addOP} 32 it is induced directly that each block contains at least one operation. Now it follows directly by Lemma ^{blockSize} 18. \square

ordering

Definition 20 (Ordering of operations inside a node). \blacktriangleright Note that from Lemma ^{blockSize} 18 we know there is at most one operation from each process in a given block.

- $E(n, i)$ is the sequence of enqueue operations that are member of **n.blocks[i]** ordered by process id.
- $D(n, i)$ is the sequence of dequeue operations that are member of **n.blocks[i]** ordered by process id.
- $D(n) = D(n, 1).D(n, 2).D(n, 3)...$
- $L = E(root, 1).D(root, 1).E(root, 2).D(root, 2).E(root, 3).D(root, 3)...$

Theorem 21. *The queue implementation is linearizable.*

Proof. We show that the ordering **L** stored in the root, satisfies the properties of a linearizable ordering.

1. If op_1 ends before op_2 begins in **E**, then op_1 comes before op_2 in **T**.

\blacktriangleright This is followed by Lemma ^{append} 17. The time op_1 ends it is in root, before op_2 , by Definition ^{ordering} 20 op_1 is before op_2 .

2. Responses to operations in E are same as they would be if done sequentially in order of L .

► Enqueue operations do not have any response so it does not matter how they are ordered. It remains to prove Dequeue d returns the correct response according to the linearization order. By Lemma ^{computeHead}25 it is deduced that the head of the queue at time of the linearization of d is computed properly. If the Queue is not empty by Lemma ^{get}22 we know that the returning response is the computed index element.

□

^{get} **Lemma 22** (Get). *Get(n, b, i) returns i th Enqueue in $E(n, b)$.*

Proof. It is obvious that `Get(leaf $l, b, 1$)` returns the operations stored in b th block of leaf l . To find the i th enqueue in block b of an internal node n in line 87 it is decided that it resides in the left child or the right child. This decision is made by Definition of $E(n, b)$. After that Lines 88, 92 search the proper subblocks of b . From Definition ^{def::subblock}5 we know the subblocks of the b th block are within the `prevBlock` and `lastBlock` block of the `CreateBlock()`.

□

Lemma 23 (Index). *Index(n, b, i) returns the rank in the $D(\text{root})$ of i th Dequeue in $D(n, b)$.*

Proof. `Index(n, b, i)` computes superblock of i th Dequeue in b th block of n in $n.\text{parent}$ and then computes the order in $D(n.\text{parent}, \text{superblock})$. Then calls `index()` on $n.\text{parent}$ recursively. It is easy to see why the second is correct. Correctness of computing superblock comes from Lemma ^{superBlock}24.

□

^{superBlock} **Lemma 24** (Computing SuperBlock). *If Index(n, b, i) performs line 101, then `superblock` contains i th Dequeue in b th block of node n .*

Proof. 1. Value read for `super[b.time]` in line 101 is not null.

► Values `c_dir` read in lines 23, `super` are set before incrementing in lines 26, 27.

2. `super[]` preserves order from child to parent; if in a child block b is before c then $b.\text{time} \leq c.\text{time}$ and `super[b.time] ≤ super[c.time]`

► Follows from the order of lines 60, 48, 49.

3. `super[i+1]-super[i] ≤ p`

► In a Refresh with successful CAS in line 46, `super` and `counter` are set for each child in lines 48, 49. Assume the current value of the counter in node n is $i+1$ and still `super[i+1]` is not set. If an instance of successful `Refresh(n)` finishes `super[i+1]` is set a new value and a block is added after $n.\text{parent}[\text{sup}[i]]$. There could be at most p successful unfinished concurrent instances of `Refresh()` that have not reached line 49. So the distance between `super[i+1]` and `super[i]` is less than p .

4. Superblock of b is within range $\pm 2p$ of the `super[b.time]`.

► `super[i]` is the index of the superblock of a block containing block b , followed by Lemma ^{superCounter}27. It is trivial to see that $n.\text{super}$ and $n.b.\text{counter}$ are increasing. `super(b)` is the real superblock of b . `super(t)` is the index of the superblock of the last block with time t . If $b.\text{time}$ is t we have:

$$\text{super}[t] - p \leq \text{super}[t-1] \leq \text{super}(t-1) \leq \text{super}(b) \leq \text{super}(t+1) \leq \text{super}(t+1) \leq \text{super}[t] + p$$

□

^{computeHead} **Lemma 25** (Computing Queue's Head). *Let Q be state of the queue if the operations before i th Dequeue in $L(\text{root})$ are applied on the Queue sequentially and X be the head of Q . If Q is empty `ComputeHead(i, b)` returns -1, otherwise returns index in $E(\text{root}, b)$ of X .*

^{head} **Lemma 26** (Validity of head). *No two blocks are written in the same index in $n.\text{blocks}$.*

Proof. **head** is incremented in lines 51, 54 after trying to append a block to the index of the last **head** read. If it was successful, we have to do this, but if it was unsuccessful, it means it has appended to the index before, so we have to update the **head**. If a process dies before line 51, another process will increment **head** in line 54. \square

erCounter

Lemma 27 (Validity of **super** and **counter**). *If $\text{super}[i] \neq \text{null}$, then $\text{super}[i]$ in node n is the superblock of a block with $\text{time}=i$.*

Proof. After a successful CAS in line 46 **super** and **counter** are modified in both children. **super**[*i*] is supposed to be the superblock of a block with **time**=*i* and **counter** is the timer in each node. **super**[*i*] and **counter**=*i* are expected to update after a bunch of blocks with **time**=*i* have been aggregated together into a block in the parent. If the process dies before line 48 these values remain unchanged and incoming blocks will get the same **time**. We claim that our algorithm still works since at most p processes die and it will not change our complexity. If a process dies right after line 48, then **counter** will remain the same and **super**[*i*] is correct. Furthermore we are sure that when **super**[*i*] is read it will not be **null**. \square

search

Lemma 28 (Search Ranges). *Preconditions of all invocation of **BSearch** are satisfied.*

Proof. Line 83: **Get**(*i*) is called if the result of a dequeue is not null. The search is among all blocks in the root.

Line 88: This search tries to find the *i*th enqueue, knowing that it is in the left child. Search is done over the left subblocks. The start and end of the range are followed by definition. Line 92 is the same.

Line 101: Here, the goal is to find the superblock. We know the distance between answer and the **super**[*i*] is at most p , since at most p processes could die. \square

Overall The main difference between the main algorithm and the block tree is that we separate enqueues and dequeues, compute the number of non-null dequeues and the Queue's size in each block in the root. We emphasize that these changes work correctly; every other claim in Block trees applies here.

6 Make first level search fast

References

andraJT98

- [1] T. D. Chandra, P. Jayanti, and K. Tan. A polylog time wait-free construction for closed objects. In B. A. Coan and Y. Afek, editors, *Proceedings of the Seventeenth Annual ACM Symposium on Principles of Distributed Computing, PODC '98, Puerto Vallarta, Mexico, June 28 - July 2, 1998*, pages 287–296. ACM, 1998.

nstamST10

- [2] A. Gidenstam, H. Sundell, and P. Tsigas. Cache-aware lock-free queues for multiple producers/consumers and weak memory consistency. In C. Lu, T. Masuzawa, and M. Mosbah, editors, *Principles of Distributed Systems - 14th International Conference, OPODIS 2010, Tozeur, Tunisia, December 14-17, 2010. Proceedings*, volume 6490 of *Lecture Notes in Computer Science*, pages 302–317. Springer, 2010.

ffmanSS07

- [3] M. Hoffman, O. Shalev, and N. Shavit. The baskets queue. In E. Tovar, P. Tsigas, and H. Fouchal, editors, *Principles of Distributed Systems, 11th International Conference, OPODIS 2007, Guadeloupe, French West Indies, December 17-20, 2007. Proceedings*, volume 4878 of *Lecture Notes in Computer Science*, pages 401–414. Springer, 2007.

ayanti98a

- [4] P. Jayanti. A time complexity lower bound for randomized implementations of some shared objects. In B. A. Coan and Y. Afek, editors, *Proceedings of the Seventeenth Annual ACM Symposium on Principles of Distributed Computing, PODC '98, Puerto Vallarta, Mexico, June 28 - July 2, 1998*, pages 201–210. ACM, 1998.

► *PRBTree*[*RBTNode*] *rootBlocks*

A persistent redblack tree supporting *append(n),get(i),split(j)*.
append(n) returns true in case successful.

► *RBTNode*

- *RootBlock* *block*
- *int* *age*
number of finished operations in the block
- *int* *step*
If this node is the root of the *RBTree*, then *step* shows the number of appended nodes to the *RBTree*

► *leaf* extends *Node*

- *int*[] *response*
leaf.response[i] stores response of *leaf.ops[i]*
- *int* *maxOld*
Index of the youngest old block in the root that this process has seen yet.

```

1: boolean REFRESH(node n)                                ▷ if n is root
2:   new=CreateBlock()                                    ▷ TODO
3:   if new.num==0 then return true
4:   else if RBTAPPEND(new) then
5:     ..
6:   end if
7: end REFRESH

7: <int, int> INDEX(node n, int b, int i)                  ▷ Returns the order in the
root among dequeues, of ith dequeue in bth block of node n
8:   if n is root then return search in the current RBTree
9:   else..
10:  end if
11: end INDEX

12: Object DEQUEUE()
13:   ..
14:   Ni= RBTnode n that n.block contains the invocation of Dequeue
15:   Nr= RBTnode n that n.block contains the response of Dequeue
16:   Ni.age= Ni.age+1                                     ▷ this is a shared counter
17:   Nr.age= Nr.age+1
18:   if Ni or Nr become old then update maxOld
19:   end if
20: end DEQUEUE

21: void RBTAPPEND(block b)                                ▷ adds block b to the rootBlocks
22:   root= rootBlocks.root
23:   if root.step%p2==0 then
24:     Help()
25:     CollectGarbage()
26:   end if
27:   new= RBTnode(b,0,null)
28:   return rootBlocks.append(new)
29: end RBTAPPEND

30: void HELP                                                ▷ Helps pending operations
31:   last= l.head-1
32:   for leaf l in leaves do
33:     if l.blocks[last] is not null then
34:       if op is DEQ then
35:         run Dequeue() for l.ops[last] after Propagate()
36:         write the response to l.responses[last]
37:       end if
38:     end if
39:   end for
40: end HELP

41: void COLLECTGRABAGE                                    ▷ Collects the old root blocks.
42:   l=FindYoungestOld(Root.Blocks.root)
43:   t1,t2= RBT.split(l)
44:   RBTRoot.CAS(t2.root)
45: end COLLECTGRABAGE

46: Block FINDYOUNGESTOLD(b)
47:   return          read all maxOld values among leaves and decide the
largest one        ▷ There is no need to do a sasDK;Lnapgshot.
48: end FINDYOUNGESTOLD

49: response FALLBACK(op i)
50:   if operation i in leaf l cannot find its desired RootBlock then
51:     return l.response[i]
52:   end if
53: end FALLBACK

```

ayantiP05

- [5] P. Jayanti and S. Petrovic. Logarithmic-time single deleter, multiple inserter wait-free queues and stacks. In R. Ramanujam and S. Sen, editors, *FSTTCS 2005: Foundations of Software Technology and Theoretical Computer Science, 25th International Conference, Hyderabad, India, December 15-18, 2005, Proceedings*, volume 3821 of *Lecture Notes in Computer Science*, pages 408–419. Springer, 2005.

/KoganP11

- [6] A. Kogan and E. Petrank. Wait-free queues with multiple enqueueers and dequeuers. In C. Cascaval and P. Yew, editors, *Proceedings of the 16th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming, PPOPP 2011, San Antonio, TX, USA, February 12-16, 2011*, pages 223–234. ACM, 2011.

-MozesS08

- [7] E. Ladan-Mozes and N. Shavit. An optimistic approach to lock-free FIFO queues. *Distributed Comput.*, 20(5):323–341, 2008.

ichaelS96

- [8] M. M. Michael and M. L. Scott. Simple, fast, and practical non-blocking and blocking concurrent queue algorithms. In J. E. Burns and Y. Moses, editors, *Proceedings of the Fifteenth Annual ACM Symposium on Principles of Distributed Computing, Philadelphia, Pennsylvania, USA, May 23-26, 1996*, pages 267–275. ACM, 1996.

manKLLP18

- [9] G. Milman, A. Kogan, Y. Lev, V. Luchangco, and E. Petrank. BQ: A lock-free queue with batching. In C. Scheideler and J. T. Fineman, editors, *Proceedings of the 30th on Symposium on Parallelism in Algorithms and Architectures, SPAA 2018, Vienna, Austria, July 16-18, 2018*, pages 99–109. ACM, 2018.

MoirNSS05

- [10] M. Moir, D. Nussbaum, O. Shalev, and N. Shavit. Using elimination to implement scalable and lock-free FIFO queues. In P. B. Gibbons and P. G. Spirakis, editors, *SPAA 2005: Proceedings of the 17th Annual ACM Symposium on Parallelism in Algorithms and Architectures, July 18-20, 2005, Las Vegas, Nevada, USA*, pages 253–262. ACM, 2005.