

1 Pseudocode

Algorithm Fields description

◇ Shared

- *Tree* *tree* : A binary tree of *Nodes*. *root* is a pointer to the root node.

◇ Local

- **Node* *leaf*: a pointer to the process's leaf in the tree.

◇ Structures

► *Node*

- **Node* *left*, *right*, *parent* : initialized when creating the tree.
- *BlockList* *blocks*
- *int* *size*= 1: #blocks in *blocks*.
- *int* *numpropagated*= 0 : # groups of blocks that have been propagated from the node to its parent. Since it is incremented after propagating, it may be behind by 1.
- *int[]* *super*: *super[i]* stores an approximate index of the superblock of the blocks in *blocks* whose *group* field have value *i*.

► *Root* extends *Node*

- *PBRT* *blocks*
BlockList is implemented with a persistent red-black tree.

► *NonRootNode* extends *Node*

- *Block[]* *blocks*
BlockList is implemented with an array.

► *Leaf* extends *NonRootNode*

- *int* *lastdone*
Stores the index of the block in the root such that the process that owns this leaf has most recently finished the. A block is finished if all of its operations are finished. *enqueue(e)* is finished if *e* is returned by some *dequeue()* and *dequeue()* is finished when it computes its response. *put the definitions before the pseudocode*

► *Block*

▷ For a *block* in a *blocklist* we define *the prefix for the block* to be the blocks in the *BlockList* up to and including the *block*. *put the definitions before the pseudocode*

- *int* *group* : the value read from *numpropagated* when appending this block to the node.

► *LeafBlock* extends *Block*

- *Object* *element* : Each block in a leaf represents a single operation. For enqueue operations *element* is the input of the enqueue and for dequeue operations it is null.
- *Object* *response* : stores the response of the operation in the *LeafBlock*.
- *int* *sum_{enq}*, *sum_{deq}* : # enqueue, dequeue operations in the prefix for the block

► *InternalBlock* extends *Block*

- *int* *end_{left}*, *end_{right}* : index of the last subblock of the block in the left and right child
- *int* *sum_{enq-left}* : # enqueue operations in the prefix for *left.blocks[end_{left}]*
- *int* *sum_{deq-left}* : # dequeue operations in the prefix for *left.blocks[end_{left}]*
- *int* *sum_{enq-right}* : # enqueue operations in the prefix for *right.blocks[end_{right}]*
- *int* *sum_{deq-right}* : # dequeue operations in the prefix for *right.blocks[end_{right}]*

► *RootBlock* extends *InternalBlock*

- *int* *length* : length of the queue after performing all operations in the prefix for this block
 - *counter* *numfinished* : number of finished operations in the block
 - *int* *order* : the index of the block in the *BlockList* containing the block.
-

Variable naming:

- *b_{op}*: index of the block containing operation *op*
- *r_{op}*: rank of operation *op* i.e. the ordering among the operations of its type according to linearization ordering

Abbreviations:

- *blocks[b].sum_x*=*blocks[b].sum_{x-left}*+*blocks[b].sum_{x-right}* (for *b*≥0 and *x* ∈ {*enq*, *deq*})
- *blocks[b].sum*=*blocks[b].sum_{enq}*+*blocks[b].sum_{deq}* (for *b*≥0)
- *blocks[b].num_x*=*blocks[b].sum_x*-*blocks[b-1].sum_x*
(for *b*>0 and *x* ∈ {∅, *enq*, *deq*, *enq-left*, *enq-right*, *deq-left*, *deq-right*}, *blocks[0].num_x*=0)

Algorithm Queue

```
201: void ENQUEUE(Object e) ▷ Creates a block with element e and appends it to the tree.
202:   block newBlock= NEW(LeafBlock)
203:   newBlock.element= e
204:   newBlock.sumenq= leaf.blocks[leaf.size].sumenq+1
205:   newBlock.sumdeq= leaf.blocks[leaf.size].sumdeq
206:   leaf.APPEND(newBlock)
207: end ENQUEUE

208: Object DEQUEUE()
209:   block newBlock= NEW(LeafBlock) ▷ Creates a block
    with null value element, appends it to the tree, computes its order among
    operations, then computes and returns its response.
210:   newBlock.element= null
211:   newBlock.sumenq= leaf.blocks[leaf.size].sumenq
212:   newBlock.sumdeq= leaf.blocks[leaf.size].sumdeq+1
213:   leaf.APPEND(newBlock)
214:   return leaf.HELPDEQUEUE()

215: end DEQUEUE

216: int, int FINDRESPONSE(int i, int b) ▷ Computes the rank and
    index of the block in the root of the enqueue that is the response of the ith
    dequeue in the root's bth block. Returns <-1,--> if the queue is empty.
217:   if root.blocks[b-1].length + root.blocks[b].numenq - i < 0 then
218:     return <-1,-->
219:   else
    ▷ We call the dequeues that
    return a value non-null dequeues. rth non-null dequeue returns the element
    of th rth enqueue. We can compute # non-null dequeues in the prefix for
    a block this way: #non-null dequeues= length - #enqueuees. Note that the
    ith dequeue in the given block is not a non-null dequeue.
220:     renq= root.blocks[b-1].sumenq- root.blocks[b-1].length + i
221:     return <root.blocks.get(enq, renq).order, renq>
222:   end if
223: end FINDRESPONSE
```

Algorithm Node

```

301: void PROPAGATE()
302:   if not REFRESH() then
303:     REFRESH()
304:   end if
305:   if this is not root then
306:     parent.PROPAGATE()
307:   end if
308: end PROPAGATE

309: boolean REFRESH()
310:   s= size
311:   <new, npleft, npright>= CREATEBLOCK(h)
312:   if new.num==0 then return true
313:   else if blocks.tryAppend(new, s) then
314:     for each dir in {left, right} do
315:       CAS(dir.super[npdir], null, h+1)
316:       CAS(dir.numpropagated, npdir, npdir+1)
317:     end for
318:     CAS(size, h, h+1)
319:     return true
320:   else
321:     CAS(size, h, h+1)
322:     return false
323:   end if
324: end REFRESH

325: int BSEARCH(field f, int i, int start, int end)
326: end BSEARCH

327: <Block, int, int> CREATEBLOCK(int i)
328:   block newBlock= NEW(block)
329:   newBlock.group= numpropagated
330:   newBlock.order= i
331:   for each dir in {left, right} do
332:     indexlast= dir.size
333:     indexprev= blocks[i-1].enddir
334:     newBlock.enddir= indexlast
335:     blocklast= dir.blocks[indexlast]
336:     blockprev= dir.blocks[indexprev]
337:     > newBlock includes dir.blocks[indexprev+1..indexlast].
338:     thisdir= dir.numpropagated
339:     newBlock.sumenq-dir= blocks[i-1].sumenq-dir + blocklast.sumenq
340:     - blockprev.sumenq
341:     newBlock.sumdeq-dir= blocks[i-1].sumdeq-dir + blocklast.sumdeq
342:     - blockprev.sumdeq
343:   end for
344:   if this is root then
345:     newBlock.length= max(root.blocks[i-1].length + b.numenq -
346:     b.numdeq, 0)
347:   end if
348:   return <b, npleft, npright>
349: end CREATEBLOCK

```

↪ Precondition: blocks[start..end] contains a block with field $f \geq i$

325: int BSEARCH(field f, int i, int start, int end)

↪ Does binary search for the value i of the given prefix sum field. Returns the index of the leftmost block in blocks[start..end] whose field f is $\geq i$.

326: end BSEARCH

Algorithm Node

↪ Precondition: $\text{blocks}[b].\text{num}_{\text{enq}} \geq i$

```
401: element GETENQ(int b, int i)
402:   if this is leaf then
403:     return blocks[b].element
404:   else if  $i \leq \text{blocks}[b].\text{num}_{\text{enq-left}}$  then ▷ i exists in the left child of this node
405:     subBlock= left.BSEARCH( $\text{sum}_{\text{enq}}$ , i,  $\text{blocks}[b-1].\text{end}_{\text{left}}+1$ ,  $\text{blocks}[b].\text{end}_{\text{left}}$ ) ▷ Search range of left child's subblocks of blocks[b].
406:     return left.GET( $i-\text{left.blocks}[\text{subBlock}-1].\text{sum}_{\text{enq}}$ , subBlock)
407:   else
408:     i= i- $\text{blocks}[b].\text{num}_{\text{enq-left}}$ 
409:     subBlock= right.BSEARCH( $\text{sum}_{\text{enq}}$ , i,  $\text{blocks}[b-1].\text{end}_{\text{right}}+1$ ,  $\text{blocks}[b].\text{end}_{\text{right}}$ ) ▷ Search range of right child's subblocks of blocks[b].
410:     return right.GET( $i-\text{right.blocks}[\text{subBlock}-1].\text{sum}_{\text{enq}}$ , subBlock)
411:   end if
412: end GETENQ
```

↪ Precondition: b th block of the node has propagated up to the root and $\text{blocks}[b].\text{num}_{\text{enq}} \geq i$.

```
413: <int, int> INDEXDEQ(int b, int i) ▷ Returns the rank of i-th dequeue in the bth block of the node, among the dequeues in the root.
414:   if this is root then
415:     return <b, i>
416:   else
417:     dir= (parent.left==n)? left: right ▷ check if a left or a right child
418:     superBlock= parent.BSEARCH( $\text{sum}_{\text{deq-dir}}$ , i,  $\text{super}[\text{blocks}[b].\text{group}]-p$ ,  $\text{super}[\text{blocks}[b].\text{group}]+p$ )
▷ superblock's group has at most p difference with the value stored in super[].
419:     if dir is right then
420:       i+=  $\text{blocks}[\text{superBlock}].\text{sum}_{\text{deq-left}}$  ▷ consider the dequeues from the right child
421:     end if
422:     return this.parent.INDEXDEQ(superBlock, i)
423:   end if
424: end INDEX
```

Algorithm Leaf

```
501: void APPEND(block blk) ▷ Append is only called by the owner of the leaf.
502:     size+=1
503:     blk.group= size
504:     blocks[size]= blk
505:     parent.PROPAGATE()
506: end APPEND

507: Object HELPDEQUEUE()
508:     <bdeq, rdeq>= INDEXDEQ(leaf.size, 1) ▷ r is the rank among the dequeues of the dequeue of the bdeqth block in the root containing.
509:     <renq, renq>= FINDRESPONSE(rdeq, bdeq) ▷ renq is the rank of the enqueue whose element is the response to the dequeue in the block containing it and
bdeq is the index of that block of it in the blocklist. If the response is null then rdeq is -1.
510:     if renq==-1 then
511:         output= null
512:         root.blocks[bdeq].numfinished.inc() ▷ shared counter
513:         if root.blocks[bdeq].numfinished==root.blocks[bdeq].num then
514:             lastdone= bdeq
515:         end if
516:     else
517:         output= GETENQ(benq, renq) ▷ getting the reponse's element.
518:         root.blocks[benq].numfinished.inc()
519:         root.blocks[benq].numfinished.inc()
520:         if root.blocks[bdeq].numfinished==root.blocks[bdeq].num then
521:             lastdone= bdeq
522:         else if root.blocks[benq].numfinished==root.blocks[benq].num then
523:             lastdone= benq
524:         end if
525:     end if
526:     return output
527: end DEQUEUE
```

Algorithm BlockList

▷ : Supports two operations `blocks.tryAppend(Block b)`, `blocks[i]`. Initially empty, when `blocks.tryAppend(b,`

`n)` returns true `b` is appended to `blocks[n]` and `blocks[i]` returns i th block in the blocks. If some instance of `blocks.tryAppend(b, n)` returns false there is a concurrent instance of `blocks.tryAppend(b', n)` which has returned true. `blocks[0]` contains an empty block with all fields equal to 0 and `endleft`, `endright` pointers to the first block of the corresponding children.

◇ *PBRT implementation*

A persistent red-black tree supporting `append(b, key)`, `get(key=i)`, `split(j)`. `append(b, key)` returns true in case successful. Since order, `sumenq` are both strictly increasing we can use one of them for another.

`root`: pointer to the root of the PBRT

601: `boolean TRYAPPEND(block blk, int n)`

▷ adds block `b` to the `root.blocks[n]`

602: **if** `root.size % p2 == 0` **then**

▷ Help every often p^2 operations appended to the root.

603: `Help()`

604: `CollectGarbage()`

605: **end if**

606: `blk.numfinished = 0`

607: `*oldRoot = &root.blocks.root`

608: `*newRoot = root.blocks.Append(blk).root`

609: **return** `CAS(root, oldRoot, newRoot)`

610: **end TRYAPPEND**

◇ *Array implementation*

`blocks[]`: array of blocks

611: `boolean TRYAPPEND(block blk, int n)`

612: **return** `CAS(blocks[n], null, blk)`

613: **end TRYAPPEND**

801: `void HELP`

▷ Helps pending operations

802: **for** leaf `l` **in** `leaves` **do**

814: `Block FINDMOSTRECENTDONE(b)`

803: `last = l.size - 1` ▷ `l.blocks[last]` can not be null because size increases after appending, see lines 603-602.

815: **for** leaf `l` **in** `leaves` **do**

816: `max = Max(l.maxOld, max)`

804: **if** `l.blocks[last].element == null` **then** ▷ operation is dequeue

817: **end for**

805: `l.blocks[last].response = l.HELPDEQUEUE()`

818: **return** `max`

▷ This snapshot suffices.

806: **end if**

819: **end FINDMOSTRECENTDONE**

807: **end for**

808: **end HELP**

820: `response FALLBACK(op i)` ▷ how to use as exception handling? by adding try catch in all the methods reading the root?

809: `void COLLECTGARBAGE` ▷ Collects the root blocks that are done.

821: **if** `root.blocks.get(numenq), i` is null **then** ▷ this enqueue was already finished

810: `s = FindMostRecentDone(Root.Blocks.root)` ▷ Lemma: If block `b` is done after helping then all blocks before `b` are done as well.

822: **return** `this.leaf.response(block.order)`

811: `t1, t2 = RBT.split(order, s)`

823: **end if**

812: `RBTRoot.CAS(t2.root)`

824: **end FALLBACK**

813: **end COLLECTGARBAGE**

2 Proof of Linearizability

Definition 1. If $n.blocks[i] == b$ we call i the *index* of block b in node n . Block b is before block b' in node n if and only if b 's index is smaller than b' 's. Block b is propagated to node n or set S if b is in $n.blocks$ or S or is a subblock of a block in $n.blocks$ or S .

Definition 2. Block b in node n is in *Established*(n, t) if $n.head$ is greater than b 's index at time t .

Progress

Lemma 3 (headProgress). $n.head$ is non-decreasing over time.

Position

Lemma 4 (headPosition). The value read in Line 52($h=n.head$) might be 1 bit behind the first empty block.

Order

Lemma 5 (establishedOrder). If time $t < t'$, then $Established(n, t) \subseteq Established(n, t')$.

CreateBlock

Lemma 6 (createBlock). Suppose $CreateBlock(n, h, x)$ is invoked at time t . The blocks propagated to $Established(n.left, t)$ and $Established(n.right, t)$ that are not propagated to $Established(n, t)$, are subblock of the block returned by $CreateBlock(n, h, x)$.

Refresh

Lemma 7 (trueRefresh). Suppose $Refresh(n)$'s $CAS(n.blocks[h], null, new)$ returns *true*. Let t be the time $Refresh(n)$ is invoked, blocks propagated to $Established(n.left, t)$ and $Established(n.right, t)$ are propagated to in $Established(n, t)$ after $CAS(n.blocks[h], null, new)$.

FalseRefresh

Lemma 8 (falseRefresh). If instance r of $Refresh(n)$ returns *false*, then there is another successful instance r' of $Refresh(n)$ that has performed a successful $CAS(n.blocks[h], null, new)$ (Line 49) after Line 43($h = n.head$) of r .

DoubleRefresh

Lemma 9 (Double Refresh). Consider two consecutive instances r_1, r_2 of $Refresh(n)$ by the same process (Lines 35,36). Let be the time before r_1 invoked. After r_2 's CAS all the blocks propagated to $Established(n.left, t)$ and $Established(n.right, t)$ are in $Established(n, t)$.

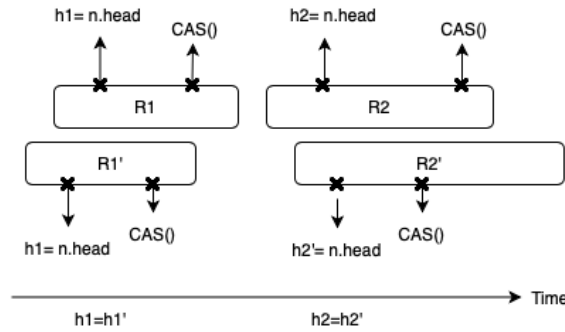


Figure 1: R_2 's CAS is executed after $h1=n.head$.

Block *new* is created of new established subblocks of children of n (Lemma 6, Line 46). If CAS in Line 48 succeeds then by Lemma 7 new established blocks will be in n .

DoubleRefresh

Lemma 10 (Double Refresh). All operations in n 's children's blocks before line 35 are guaranteed to be in n 's blocks after Line 37.

$CreateBlock()$ reads blocks in the children that do not exist in the parent and aggregates them into one block. If a $Refresh()$ procedure returns *true* it means it has appended the block created by $CreateBlock()$ into the parent node's sequence. So suppose two $Refresh$ es fail. Since the first $Refresh()$ was not successful, it means another CAS operation by a $Refresh$, concurrent to the first $Refresh()$, was successful before the second $Refresh()$. So it means the second failed $Refresh$ is concurrent with a successful $Refresh()$ that assuredly has read block before the mentioned line 35. After all it means if any of the $Refresh()$ attempts were successful the claim is true, and also if both fail the mentioned claim still holds.

Append

Lemma 11 (Append). When $Append(op)$ is finished, op appears exactly once in a block of $root.blocks$.

BlockSize

Lemma 12 (Block Size Upper Bound). Each block in a node contains at most one operation from each process.

SubblocksBound

Lemma 13 (Subblocks Upperbound). Each block in a node has at most p subblocks.

ordering	<p>Definition 14 (Ordering of operations inside a node). ► Note that from Lemma 12 we know there is at most one operation from each process in a given block.</p> <ul style="list-style-type: none"> • $E(n, i)$ is the sequence of enqueue operations that are member of $\mathbf{n.blocks[i]}$ ordered by process id. • $D(n, i)$ is the sequence of dequeue operations that are member of $\mathbf{n.blocks[i]}$ ordered by process id. • $D(n) = D(n, 1).D(n, 2).D(n, 3)...$ • $L = E(\mathit{root}, 1).D(\mathit{root}, 1).E(\mathit{root}, 2).D(\mathit{root}, 2).E(\mathit{root}, 3).D(\mathit{root}, 3)...$ <p>Theorem 15. <i>The queue implementation is linearizable.</i></p>
get	<p>Lemma 16 (Get). <i>$\mathit{Get}(\mathbf{n}, \mathbf{b}, \mathbf{i})$ returns ith Enqueue in $E(n, b)$.</i></p>
	<p>Lemma 17 (Index). <i>$\mathit{Index}(n, b, i)$ returns the rank in the $D(\mathit{root})$ of ith Dequeue in $D(n, b)$.</i></p>
superBlock	<p>Lemma 18 (Computing SuperBlock). <i>If $\mathit{Index}(\mathbf{n}, \mathbf{b}, \mathbf{i})$ performs line 101, then $\mathbf{superblock}$ contains ith Dequeue in \mathbf{b}th block of node \mathbf{n}.</i></p>
computeHead	<p>Lemma 19 (Computing Queue's Head). <i>Let Q be state of the queue if the operations before ith Dequeue in $L(\mathit{root})$ are applied on the Queue sequentially and X be the head of Q. If Q is empty $\mathit{ComputeHead}(\mathbf{i}, \mathbf{b})$ returns -1, otherwise returns index in $E(\mathit{root}, b)$ of X.</i></p>
head	<p>Lemma 20 (Validity of head). <i>No two blocks are written in the same index in $\mathbf{n.blocks}$.</i></p>
superCounter	<p>Lemma 21 (Validity of super and counter). <i>If $\mathbf{super[i]} \neq \mathbf{null}$, then $\mathbf{super[i]}$ in node \mathbf{n} is the superblock of a block with $\mathbf{time=i}$.</i></p>
search	<p>Lemma 22 (Search Ranges). <i>Preconditions of all invocation of $\mathbf{BSearch}$ are satisfied.</i></p>