

---

**Algorithm** Fields description

---

◇ *Shared*

- *Tree tree* : A binary tree of Nodes. *root* is a pointer to the root node.

◇ *Local*

- *\*Node leaf* : a pointer to the process's leaf in the tree.

◇ *Structures*► *Node*

- *\*Node left, right, parent* : initialized when creating the tree.
- *BlockList blocks*
- *int size= 1* : #blocks in blocks.
- *int num<sub>propagated</sub>= 0* : # groups of blocks that have been propagated from the node to its parent. Since it is incremented after propagating, it may be behind by 1.
- *int[] super*: *super[i]* stores an approximate index of the superblock of the blocks in *blocks* whose *group* field have value *i*.

► *Root extends Node*

- *PBRT blocks*  
BlockList is implemented with a persistent red-black tree.

► *NonRootNode extends Node*

- *Block[] blocks*  
BlockList is implemented with an array.

► *Leaf extends NonRootNode*

- *int last<sub>done</sub>*  
Stores the index of the block in the root such that the process that owns this leaf has most recently finished the. A block is finished if all of its operations are finished. *enqueue(e)* is finished if *e* is returned by some *dequeue()* and *dequeue()* is finished when it computes its response. *put the definitions before the pseudocode*

► *Block*

▷ For a *block* in a *blocklist* we define *the prefix for the block* to be the blocks in the *BlockList* up to and including the *block*.  
*put the definitions before the pseudocode*

- *int group* : the value read from *num<sub>propagated</sub>* when appending this block to the node.

► *LeafBlock extends Block*

- *Object element* : Each block in a leaf represents a single operation. For enqueue operations *element* is the input of the enqueue and for dequeue operations it is null.
- *Object response* : stores the response of the operation in the *LeafBlock*.
- *int sum<sub>enq</sub>, sum<sub>deq</sub>* : # enqueue, dequeue operations in the prefix for the block

► *InternalBlock extends Block*

- *int end<sub>left</sub>, end<sub>right</sub>* : index of the last subblock of the block in the left and right child
- *int sum<sub>enq-left</sub>* : # enqueue operations in the prefix for *left.blocks[end<sub>left</sub>]*
- *int sum<sub>deq-left</sub>* : # dequeue operations in the prefix for *left.blocks[end<sub>left</sub>]*
- *int sum<sub>enq-right</sub>* : # enqueue operations in the prefix for *right.blocks[end<sub>right</sub>]*
- *int sum<sub>deq-right</sub>* : # dequeue operations in the prefix for *right.blocks[end<sub>right</sub>]*

► *RootBlock extends InternalBlock*

- *int length* : length of the queue after performing all operations in the prefix for this block
- *counter num<sub>finished</sub>* : number of finished operations in the block
- *int order* : the index of the block in the *BlockList* containing the block.

---

*Variable naming:*

- *b<sub>op</sub>*: index of the block containing operation *op*
- *r<sub>op</sub>*: rank of operation *op* i.e. the ordering among the operations of its type according to linearization ordering

*Abbreviations:*

- *blocks[b].sum<sub>x</sub>=blocks[b].sum<sub>x-left</sub>+blocks[b].sum<sub>x-right</sub>* (for  $b \geq 0$  and  $x \in \{\text{enq}, \text{deq}\}$ )
- *blocks[b].sum=blocks[b].sum<sub>enq</sub>+blocks[b].sum<sub>deq</sub>* (for  $b \geq 0$ )
- *blocks[b].num<sub>x</sub>=blocks[b].sum<sub>x</sub>-blocks[b-1].sum<sub>x</sub>*  
(for  $b > 0$  and  $x \in \{\emptyset, \text{enq}, \text{deq}, \text{enq-left}, \text{enq-right}, \text{deq-left}, \text{deq-right}\}$ , *blocks[0].num<sub>x</sub>=0*)

---

**Algorithm Queue**

---

```
201: void ENQUEUE(Object e)  ▷ Creates a block with element e and appends it to the tree.
202:   block newBlock= NEW(LeafBlock)
203:   newBlock.element= e
204:   newBlock.sumenq= leaf.blocks[leaf.size].sumenq+1
205:   newBlock.sumdeq= leaf.blocks[leaf.size].sumdeq
206:   leaf.APPEND(newBlock)
207: end ENQUEUE

208: Object DEQUEUE()
209:   block newBlock= NEW(LeafBlock)  ▷ Creates a block
    with null value element, appends it to the tree, computes its order among
    operations, then computes and returns its response.
210:   newBlock.element= null
211:   newBlock.sumenq= leaf.blocks[leaf.size].sumenq
212:   newBlock.sumdeq= leaf.blocks[leaf.size].sumdeq+1
213:   leaf.APPEND(newBlock)
214:   return leaf.HELPDEQUEUE()

215: end DEQUEUE

216: int, int FINDRESPONSE(int i, int b)  ▷ Computes the rank and
    index of the block in the root of the enqueue that is the response of the ith
    dequeue in the root's bth block. Returns <-1,--> if the queue is empty.
217:   if root.blocks[b-1].length + root.blocks[b].numenq - i < 0 then
218:     return <-1,-->
219:   else
    ▷ We call the dequeues that
    return a value non-null dequeues. rth non-null dequeue returns the element
    of th rth enqueue. We can compute # non-null dequeues in the prefix for
    a block this way: #non-null dequeues= length - #enqueuees. Note that the
    ith dequeue in the given block is not a non-null dequeue.
220:     renq= root.blocks[b-1].sumenq- root.blocks[b-1].length + i
221:     return <root.blocks.get(enq, renq).order, renq>
222:   end if
223: end FINDRESPONSE
```

---

---

**Algorithm Node**


---

```

301: void PROPAGATE()
302:   if not REFRESH() then
303:     REFRESH()
304:   end if
305:   if this is not root then
306:     parent.PROPAGATE()
307:   end if
308: end PROPAGATE

309: boolean REFRESH()
310:   s= size
311:   <new, npleft, npright>= CREATEBLOCK(h)
312:   if new.num==0 then return true
313:   else if blocks.tryAppend(new, s) then
314:     for each dir in {left, right} do
315:       CAS(dir.super[npdir], null, h+1)
316:       CAS(dir.numpropagated, npdir, npdir+1)
317:     end for
318:     CAS(size, h, h+1)
319:     return true
320:   else
321:     CAS(size, h, h+1)
322:     return false
323:   end if
324: end REFRESH

327: <Block, int, int> CREATEBLOCK(int i)
328:   block newBlock= NEW(block)
329:   newBlock.group= numpropagated
330:   newBlock.order= i
331:   for each dir in {left, right} do
332:     indexlast= dir.size
333:     indexprev= blocks[i-1].enddir
334:     newBlock.enddir= indexlast
335:     blocklast= dir.blocks[indexlast]
336:     blockprev= dir.blocks[indexprev]
337:     thisdir= dir.numpropagated
338:     newBlock.sumenq-dir= blocks[i-1].sumenq-dir + blocklast.sumenq
339:     - blockprev.sumenq
340:     newBlock.sumdeq-dir= blocks[i-1].sumdeq-dir + blocklast.sumdeq
341:     - blockprev.sumdeq
342:   end for
343:   newBlock.length= max(root.blocks[i-1].length + b.numenq -
344:   b.numdeq, 0)
345:   return <b, npleft, npright>
346: end CREATEBLOCK

```

↪ Precondition: blocks[start..end] contains a block with field  $f \geq i$

```

325: int BSEARCH(field f, int i, int start, int end)
326:   i of the given prefix sum field. Returns the index of the leftmost block in
   blocks[start..end] whose field f is  $\geq i$ .
326: end BSEARCH

```

---

---

**Algorithm** Node

---

↪ Precondition:  $\text{blocks}[b].\text{num}_{\text{enq}} \geq i$

```
401: element GETENQ(int b, int i)
402:   if this is leaf then
403:     return blocks[b].element
404:   else if  $i \leq \text{blocks}[b].\text{num}_{\text{enq-left}}$  then                                ▷ i exists in the left child of this node
405:     subBlock= left.BSEARCH( $\text{sum}_{\text{enq}}$ , i,  $\text{blocks}[b-1].\text{end}_{\text{left}}+1$ ,  $\text{blocks}[b].\text{end}_{\text{left}}$ )    ▷ Search range of left child's subblocks of  $\text{blocks}[b]$ .
406:     return left.GET( $i-\text{left.blocks}[\text{subBlock}-1].\text{sum}_{\text{enq}}$ , subBlock)
407:   else
408:     i=  $i-\text{blocks}[b].\text{num}_{\text{enq-left}}$ 
409:     subBlock= right.BSEARCH( $\text{sum}_{\text{enq}}$ , i,  $\text{blocks}[b-1].\text{end}_{\text{right}}+1$ ,  $\text{blocks}[b].\text{end}_{\text{right}}$ )    ▷ Search range of right child's subblocks of  $\text{blocks}[b]$ .
410:     return right.GET( $i-\text{right.blocks}[\text{subBlock}-1].\text{sum}_{\text{enq}}$ , subBlock)
411:   end if
412: end GETENQ
```

↪ Precondition:  $b$ th block of the node has propagated up to the root and  $\text{blocks}[b].\text{num}_{\text{enq}} \geq i$ .

```
413: <int, int> INDEXDEQ(int b, int i)                                ▷ Returns the rank of  $i$ th dequeue in the  $b$ th block of the node, among the dequeues in the root.
414:   if this is root then
415:     return <b, i>
416:   else
417:     dir= (parent.left==n)? left: right                                ▷ check if a left or a right child
418:     superBlock= parent.BSEARCH( $\text{sum}_{\text{deq-dir}}$ , i,  $\text{super}[\text{blocks}[b].\text{group}]-p$ ,  $\text{super}[\text{blocks}[b].\text{group}]+p$ )
                                                                ▷ superblock's group has at most  $p$  difference with the value stored in  $\text{super}[]$ .
419:     if dir is right then
420:       i+=  $\text{blocks}[\text{superBlock}].\text{sum}_{\text{deq-left}}$                                 ▷ consider the dequeues from the right child
421:     end if
422:     return this.parent.INDEXDEQ(superBlock, i)
423:   end if
424: end INDEX
```

---

---

**Algorithm** Leaf

---

```
501: void APPEND(block blk) ▷ Append is only called by the owner of the leaf.
502:     size+=1
503:     blk.group= size
504:     blocks[size]= blk
505:     parent.PROPAGATE()
506: end APPEND

507: Object HELPDEQUEUE()
508:     <bdeq, rdeq>= INDEXDEQ(leaf.size, 1) ▷ r is the rank among the dequeues of the dequeue of the bdeqth block in the root containing.
509:     <renq, renq>= FINDRESPONSE(rdeq, bdeq) ▷ renq is the rank of the enqueue whose element is the response to the dequeue in the block containing it and
bdeq is the index of that block of it in the blocklist. If the response is null then rdeq is -1.
510:     if renq== -1 then
511:         output= null
512:         root.blocks[bdeq].numfinished.inc() ▷ shared counter
513:         if root.blocks[bdeq].numfinished==root.blocks[bdeq].num then
514:             lastdone= bdeq
515:         end if
516:     else
517:         output= GETENQ(benq, renq) ▷ getting the reponse's element.
518:         root.blocks[benq].numfinished.inc()
519:         root.blocks[benq].numfinished.inc()
520:         if root.blocks[bdeq].numfinished==root.blocks[bdeq].num then
521:             lastdone= bdeq
522:         else if root.blocks[benq].numfinished==root.blocks[benq].num then
523:             lastdone= benq
524:         end if
525:     end if
526:     return output
527: end DEQUEUE
```

---

---

**Algorithm** BlockList

---

▷ : Supports two operations `blocks.tryAppend(Block b)`, `blocks[i]`. Initially empty, when `blocks.tryAppend(b,`

`n)` returns true `b` is appended to `blocks[n]` and `blocks[i]` returns *i*th block in the blocks. If some instance of `blocks.tryAppend(b, n)` returns false there is a concurrent instance of `blocks.tryAppend(b', n)` which has returned true. `blocks[0]` contains an empty block with all fields equal to 0 and `endleft`, `endright` pointers to the first block of the corresponding children.

◇ *PBRT implementation*

A persistent red-black tree supporting `append(b, key)`, `get(key=i)`, `split(j)`. `append(b, key)` returns true in case successful. Since order, `sumenq` are both strictly increasing we can use one of them for another.

`root`: pointer to the root of the PBRT

601: `boolean TRYAPPEND(block blk, int n)`

▷ adds block `b` to the `root.blocks[n]`

602:   **if** `root.size % p2 == 0` **then**

▷ Help every often  $p^2$  operations appended to the root.

603:     `Help()`

604:     `CollectGarbage()`

605:   **end if**

606:   `blk.numfinished = 0`

607:   `*oldRoot = &root.blocks.root`

608:   `*newRoot = root.blocks.Append(blk).root`

609:   **return** `CAS(root, oldRoot, newRoot)`

610: **end TRYAPPEND**

◇ *Array implementation*

`blocks[]`: array of blocks

611: `boolean TRYAPPEND(block blk, int n)`

612:   **return** `CAS(blocks[n], null, blk)`

613: **end TRYAPPEND**

---

801: `void HELP`

▷ Helps pending operations

802:   **for** leaf `l` **in** `leaves` **do**

814: `Block FINDMOSTRECENTDONE(b)`

803:     `last = l.size - 1` ▷ `l.blocks[last]` can not be null because size increases after appending, see lines 603-602.

815:   **for** leaf `l` **in** `leaves` **do**

816:     `max = Max(l.maxOld, max)`

804:     **if** `l.blocks[last].element == null` **then** ▷ operation is dequeue

817:   **end for**

805:     `l.blocks[last].response = l.HELPDEQUEUE()`

818:   **return** `max`

▷ This snapshot suffices.

806:   **end if**

819: **end FINDMOSTRECENTDONE**

807:   **end for**

808: **end HELP**

820: `response FALLBACK(op i)` ▷ how to use as exception handling? by adding try catch in all the methods reading the root?

809: `void COLLECTGARBAGE` ▷ Collects the root blocks that are done.

821:   **if** `root.blocks.get(numenq), i` is null **then** ▷ this enqueue was already finished

810:   `s = FindMostRecentDone(Root.Blocks.root)` ▷ Lemma: If block `b` is done after helping then all blocks before `b` are done as well.

822:   **return** `this.leaf.response(block.order)`

811:   `t1, t2 = RBT.split(order, s)`

823:   **end if**

812:   `RBTRoot.CAS(t2.root)`

824: **end FALLBACK**

813: **end COLLECTGARBAGE**

---