

1 Pseudocode

Algorithm Tree Fields Description

◇ *Shared*

- A binary tree of Nodes with one leaf for each process. root is the root node.

◇ *Local*

- *Node* leaf: process's leaf in the tree.

◇ *Structures*

► *Node*

- **Node* left, right, parent : initialized when creating the tree.
- *BlockList*
- *int* head= 1: #blocks in blocks. blocks[0] is a block with all integer fields equal to zero.
- *int* num_{propagated}= 0 : # groups of blocks that have been propagated from the node to its parent. Since it is incremented after propagating, it may be behind by 1.

► *Block*

- *int* group : the value read from num_{propagated} when appending this block to the node.

► *LeafBlock* extends *Block*

- *Object* element : Each block in a leaf represents a single operation. If the operation is enqueue(x) then element=x, otherwise element=null.
- *int* sum_{enq}, sum_{deq} : # enqueue, dequeue operations in the prefix for the block

► *InternalBlock* extends *Block*

- *int* end_{left}, end_{right} : indices of the last subblock of the block in the left and right child
- *int* sum_{enq-left} : # enqueue operations in the prefix for left.blocks[end_{left}]
- *int* sum_{deq-left} : # dequeue operations in the prefix for left.blocks[end_{left}]
- *int* sum_{enq-right} : # enqueue operations in the prefix for right.blocks[end_{right}]
- *int* sum_{deq-right} : # dequeue operations in the prefix for right.blocks[end_{right}]

► *RootBlock* extends *InternalBlock*

- *int* size : size of the queue after performing all operations in the prefix for this block
-

Abbreviations:

- blocks[b].sum_x=blocks[b].sum_{x-left}+blocks[b].sum_{x-right} (for b≥0 and x ∈ {enq, deq})
- blocks[b].sum=blocks[b].sum_{enq}+blocks[b].sum_{deq} (for b≥0)
- blocks[b].num_x=blocks[b].sum_x-blocks[b-1].sum_x
(for b>0 and x ∈ {∅, enq, deq, enq-left, enq-right, deq-left, deq-right})

Algorithm Queue

```
201: void ENQUEUE(Object e) ▷ Creates a block with element e and adds it to the tree.
202:   block newBlock= NEW(LeafBlock)
203:   newBlock.element= e
204:   newBlock.sumenq= leaf.blocks[leaf.head].sumenq+1
205:   newBlock.sumdeq= leaf.blocks[leaf.head].sumdeq
206:   leaf.APPEND(newBlock)
207: end ENQUEUE

208: Object DEQUEUE() ▷ Creates a block with null value element, appends it to the tree, computes its order among operations, and returns its response.
209:   block newBlock= NEW(LeafBlock)
210:   newBlock.element= null
211:   newBlock.sumenq= leaf.blocks[leaf.head].sumenq
212:   newBlock.sumdeq= leaf.blocks[leaf.head].sumdeq+1
213:   leaf.APPEND(newBlock)
214:   <bdeq, ideq>= INDEXDEQ(leaf.head, 1)
215:   <benq, ienq>= FINDRESPONSE(bdeq, ideq)
216:   if ienq== -1 then
217:     output= null
218:   else
219:     output= GETENQ(benq, ienq)
220:   end if
221:   return output
222: end DEQUEUE

223: <int, int> FINDRESPONSE(int bd, int id) ▷ Returns <be, ie>, if  $E_{root, b_e, i_e}$  is the response to the  $D_{root, b_d, i_d}$ . Returns <-1, --> if the queue is empty.
224:   if root.blocks[bd-1].size + root.blocks[bd].numenq - id < 0
225:     then
226:       return <-1, -->
227:     else
228:       renq= id - root.blocks[bd-1].size + root.blocks[bd-1].sumenq
229:       return root.DSEARCH(renq, bd)
230:     end if
231:   end FINDRESPONSE
```

deqRest

Algorithm Node

```
301: void PROPAGATE()
firstRefresh: 302:   if not REFRESH() then
secondRefresh: 303:     REFRESH()
304:   end if
305:   if this is not root then
306:     parent.PROPAGATE()
307:   end if
308: end PROPAGATE

309: boolean REFRESH()
readHead: 310:   h= head
311:   <new, npleft, npright>= CREATEBLOCK(s)  ▷ npleft, npright are the
values read from the children's numpropagated field.
312:   if new.num==0 then return true  ▷ The block contains nothing.
313:   else if blocks.tryAppend(new, h) then
314:     for each dir in {left, right} do
315:       CAS(dir.super[npdir], null, h)  ▷ Write would work too.
316:       CAS(dir.numpropagated, npdir, npdir+1)
317:     end for
318:     CAS(head, h, h+1)
319:     return true
320:   else
321:     CAS(head, h, h+1)  ▷ Even if another process wins, help
to increase the head. The winner might have fallen sleep before increasing
322:     return false
323:   end if
324: end REFRESH

325: int BSEARCH(field f, int i, int start, int end)
▷ Does binary search for the value
i of the given prefix sum field. Returns the index of the leftmost block in
blocks[start..end] whose field f is ≥ i.
326: end BSEARCH
```

```
327: <Block, int, int> CREATEBLOCK(int i)  ▷ Creates a block
to be inserted as ith block in blocks. Returns the created block as well as
values read from each child's numpropagated field. These values are used for
incrementing the children's numpropagated field if the block was appended to
blocks successfully.
328:   block newBlock= NEW(block)
329:   newBlock.group= numpropagated
330:   for each dir in {left, right} do
331:     indexlast= dir.head
332:     indexprev= blocks[i-1].enddir
333:     newBlock.enddir= indexlast
334:     blocklast= dir.blocks[indexlast]
335:     blockprev= dir.blocks[indexprev]
336:     ▷ newBlock includes dir.blocks[indexprev+1..indexlast].
337:     npdir= dir.numpropagated
338:     newBlock.sumenq-dir= blocks[i-1].sumenq-dir + blocklast.sumenq
- blockprev.sumenq
339:     newBlock.sumdeq-dir= blocks[i-1].sumdeq-dir + blocklast.sumdeq
- blockprev.sumdeq
340:   end for
341:   if this is root then
342:     newBlock.size = max(root.blocks[i-1].size + newBlock.numenq
- newBlock.numdeq, 0)
343:   end if
344:   return <b, npleft, npright>
345: end CREATEBLOCK
```

Algorithm Root

```
▷ Precondition: root.blocks[end].sumenq ≥ e ≥ 1
801: <int, int> DSEARCH(int e, int end)  ▷ Returns <b,i> if  $E_{root,e} = E_{root,b,i}$ .
802:   start= end-1
803:   while root.blocks[start].sumenq ≥ e do
804:     start= max(start-(end-start), 0)
805:   end while
806:   b= root.BSearch(sumenq, e, start, end)
807:   i= e- root.blocks[b-1].sumenq
808:   return <b,i>
809: end DSEARCH
```

Algorithm Node

\rightsquigarrow Precondition: $\text{blocks}[b].\text{num}_{\text{enq}} \geq i$

```
401: element GETENQ(int b, int i) ▷ Returns the element of  $E_{this,b,i}$ .
402:   if this is leaf then
403:     return blocks[b].element
404:   else if  $i \leq \text{blocks}[b].\text{num}_{\text{enq-left}}$  then ▷  $E_{this,b,i}$  is in the left child of this node.
405:     subBlock= left.BSEARCH(sumenq, i+left.blocks[blocks[b].endleft-1].sumenq, blocks[b-1].endleft+1, blocks[b].endleft)
406:     return left.GETENQ(subBlock, i)
407:   else
408:     i= i-blocks[b].numenq-left
409:     subBlock= right.BSEARCH(sumenq, i+right.blocks[blocks[b].endright-1].sumenq, blocks[b-1].endright+1, blocks[b].endright)
410:     return right.GETENQ(subBlock, i)
411:   end if
412: end GETENQ
```

\rightsquigarrow Precondition: bth block of the node has propagated up to the root and $\text{blocks}[b].\text{num}_{\text{enq}} \geq i$.

```
413: <int, int> INDEXDEQ(int b, int i) ▷ Returns <x, y> if  $D_{this,b,i} = D_{root,x,y}$ .
414:   if this is root then
415:     return <b, i>
416:   else
417:     dir= (parent.left==n)? left: right ▷ check if this node is a left or a right child
418:     superBlock= parent.BSEARCH(sumdeq-dir, i+blocks[b-1].sumdeq, super[blocks[b].group]-p, super[blocks[b].group]+p)
▷ superblock's group has at most  $p$  difference with the value stored in super[].
419:     if dir is right then
420:       i+= blocks[superBlock].numdeq-left ▷ consider the dequeues from the right child
421:     end if
422:     return this.parent.INDEXDEQ(superBlock, i)
423:   end if
424: end INDEXDEQ
```

Algorithm Leaf

```
601: void APPEND(block blk) ▷ Append is only called by the owner of the leaf.
602:   head+=1
603:   blk.group= head
604:   blocks[head]= blk
605:   parent.PROPAGATE()
606: end APPEND
```

Algorithm BlockList

▷ : Supports two operations `blocks.tryAppend(Block b)`, `blocks[i]`. Initially empty, when `blocks.tryAppend(b,`

`n)` returns true b is appended to `blocks[n]` and `blocks[i]` returns *i*th block in the blocks. If some instance of `blocks.tryAppend(b, n)` returns false there is a concurrent instance of `blocks.tryAppend(b', n)` which has returned true.blocks[0] contains an empty block with all fields equal to 0 and `endleft`, `endright` pointers to the first block of the corresponding children.

block[] blocks: array of blocks

int[] super: super[i] stores an approximate index of the superblock of the blocks in blocks whose group field have value i.

```
701: boolean TRYAPPEND(block blk, int n)
702:   return CAS(blocks[n], null, blk)
703: end TRYAPPEND
```

2 Proof of Linearizability

TEST Fix the logical order of definitions (cyclic references).

TEST Is it better to show $\text{ops}(\text{EST}_n, t)$ with EST_n, t ?

Question A good notation for *the index of the b*?

Question How to remove the notion of time? To say $\text{pre}(n, i)$ contains $n.\text{blocks}[0..i]$ instead of $\text{EST}(n, t)$ which $\text{head}=i$ at time t . Is it good? Furthermore, can we remove the notion of established blocks?

Definition 1 (Block). A block is an object storing some statistics, as described in Algorithm Queue. It implicitly represents a set of operations. If $n.\text{blocks}[i] == b$ we call i the *index* of block b . Block b is before block b' in node n if and only if the index of the b is smaller than the index of the b' 's. For a block in a `BlockList` we define *the prefix for the block* to be the blocks in the `BlockList` up to and including the block.

Definition 2 (Subblock). Block b is a *direct subblock* of $n.\text{blocks}[i]$ if it is $\in n.\text{left.blocks}[n.\text{blocks}[i-1].\text{end}_{\text{left}}+1..n.\text{blocks}[i].\text{end}_{\text{left}}] \cup n.\text{right.blocks}[n.\text{blocks}[i-1].\text{end}_{\text{right}}+1..n.\text{blocks}[i].\text{end}_{\text{right}}]$ (See line ^{endDefLine} 533 for the defined range). Block b is a subblock of a $n.\text{blocks}[i]$ if it is a direct subblock of it or subblock of a direct subblock of it.

Definition 3 (Superblock). Block b is direct superblock of block c if c is a direct subblock of b . Block b is superblock of block c if c is a subblock of b .

Definition 4 (Operations of a block). A block lb in a leaf represents one operation which if it is `enqueue(x)` then $lb.\text{element}=x$, otherwise $\text{element}=\text{null}$. The set of operations of block b are the operations in the subblocks of b . We show the set of operations of block b by $\text{ops}(b)$.

For simplicity we say block b is propagated to node n or to a set of blocks S if b is in $n.\text{blocks}$ or S or is a subblock of a block in $n.\text{blocks}$ or S . We also say b contains op if $op \in \text{ops}(b)$.

Definition 5. A block b in $n.\text{blocks}$ is *established* at time t if the last value written into $n.\text{head}$ before t is greater than the index of b in $n.\text{blocks}$ at time t . $\text{EST}_{n, t}$ is the set of established blocks at time t of node n .

Observation 6. Once a block b is written in $n.\text{blocks}[i]$ then $n.\text{blocks}[i]$ never changes.

Lemma 7 (headProgress). $n.\text{head}$ is non-decreasing over time and $n.\text{blocks}[i].\text{end}_{\text{left}} \geq n.\text{blocks}[i-1].\text{end}_{\text{left}}, n.\text{blocks}[i].\text{end}_{\text{right}} \geq n.\text{blocks}[i-1].\text{end}_{\text{right}}$.

Proof. The first claim follows trivially from the pseudocode since $n.\text{head}$ is only incremented. Also when $n.\text{blocks}[i]$ is created its $\text{end}_{\text{left}}, \text{end}_{\text{right}}$ are greater than or equal to the values in $n.\text{blocks}[i-1]$. Since $\text{blocks}[i-1].\text{end}_{\text{dir}} < \text{dir.head} = \text{blocks}[i].\text{end}_{\text{dir}}$ (Lines ^{lastLine, prevLine} 77). □

Lemma 8. Every block has most one direct superblock.

Proof. To show this we are going to refer to the way $n.\text{blocks}[]$ is partitioned while propagating blocks up to $n.\text{parent}$. $n.\text{CreateBlock}(i)$ merges the blocks in $n.\text{left.blocks}[n.\text{blocks}[i-1].\text{end}_{\text{left}}..n.\text{blocks}[i].\text{end}_{\text{left}}]$ and $n.\text{right.blocks}[n.\text{blocks}[i-1].\text{end}_{\text{right}}..n.\text{blocks}[i].\text{end}_{\text{right}}]$ (Lines ^{lastLine, pr} 77). Since $\text{end}_{\text{left}}, \text{end}_{\text{right}}$ are non-decreasing, so the range of the subblocks of $n.\text{blocks}[i]$ which is $(n.\text{blocks}[i-1].\text{end}_{\text{dir}}+1..n.\text{blocks}[i].\text{end}_{\text{dir}}]$ does not overlap with the range of the subblocks of $n.\text{blocks}[i-1]$. □

Corollary 9 (No Duplicates). If op is in $n.\text{blocks}[i]$ then there is no $j \neq i$ such that $op \in \text{ops}(n.\text{blocks}[j])$.

Invariant 10 (headPosition). If the value of $n.\text{head}$ is h then, $n.\text{blocks}[i] = \text{null}$ for $i > h$ and $n.\text{blocks}[i] \neq \text{null}$ for $i < h$.

Proof. The invariant is true initially since 1 is assigned to `n.head` and `n.blocks[x]` is null for every `x`. The truth of the invariant may be affected by writing into `n.blocks` or incrementing `n.head`.

Some value is written into `n.blocks[head]` only in Line 313. It is obvious that writing into `n.blocks[head]` preserves the invariant. The value of `n.head` is modified only in lines ^{incrementHead1}318, ^{incrementHead2}321. Depending on whether the `TryAppend()` in Line ^{cas}313 succeeded or not we show that the claim holds after the increment lines of `n.head` in either case. If `head` is incremented to `h` it is sufficient to show `n.blocks[h] ≠ null` to prove the invariant still holds. In the first case the process applied a successful `TryAppend(new, h)` in line ^{okcas}314, which means `n.blocks[h]` is not null anymore. Note that whether ^{incrementHead1}318 returns true or false after Line `n.head` we know has been incremented from Line ^{readHead}310. The failure case is also the same since it means some value is written into `n.blocks[head]` by some process. \square

Explain More

Lemma 11 (establishedOrder). *If time $t < \text{time } t'$, then $\text{ops}(\text{EST}_n, t) \subseteq \text{ops}(\text{EST}_n, t')$.*

Proof. Blocks are only appended (not modified) with CAS to `n.blocks[n.head]` and `n.head` is non-decreasing, so the set of operations in established blocks of a node can only grow. \square

`CreateBlock()` aggregates the blocks in the children that are not already established in the parent into one block. If a `Refresh()` procedure returns true it means it has appended the block created by `CreateBlock()` into the parent node's sequence. So suppose two `Refreshes` fail. Since the first `Refresh()` was not successful, it means another CAS operation by a `Refresh`, concurrent to the first `Refresh()`, was successful before the second `Refresh()`. So it means the second failed `Refresh` is concurrent with another successful `Refresh()` that assuredly has read block before the mentioned line 35. After all it means if any of the `Refresh()` attempts were successful the claim is true, and also if both fail the mentioned claim still holds.

Lemma 12 (head Increment).

If an $n.Refresh$ instance reaches Line 313 and reads $head=h$ (Line 310) after it terminates $head$ is greater than h .

Proof. If Line 318 or 318 succeeded the claim holds, otherwise another process has incremented the head. \square

Lemma 13 (trueRefresh).

Let t_i be the time an instance of $n.Refresh()$ is invoked and t_t be the time it terminates. Suppose the $TryAppend(new, s)$ of the $n.Refresh()$ returns **true**, then $ops(EST_{n.left, t_i}) \cup ops(EST_{n.right, t_i}) \subseteq ops(EST_n, t_t)$.

Proof. From Lemma 11 we know that $ops(EST_n, t_i) \subseteq ops(EST_n, t_t)$. So it remains to show the operations of $ops(EST_{n.left, t_i}) \cup ops(EST_{n.right, t_i}) - ops(EST_n, t_i)$, which we call *new operations*, are all in $ops(EST_n, t_t)$. If $TryAppend$ returns **true** a block **new** is written into $n.blocks[h]$ (Line 313). We show $ops(EST_{n.left, t_i}) \subseteq ops(EST_n, t_t)$. The proof for the right child's claim is the same. Let $n.left.head$ at t_i be hli . Let $n.Refresh()$ read $head$ equal to h (Line 310). By the lines 332, 331 the new block in $n.blocks[h]$ contains $n.left.blocks[n.blocks[h-1].end_{left}+1..left.head]$. Since $left.head$ is read after t_i then $ops(EST_{n.left, t_i}) \subseteq ops(n.left.blocks[0..left.head])$. By Lemma 11 $ops(n.left.blocks[0..n.blocks[h-1].end_{left}]) \subseteq ops(EST_n, t_i) \subseteq ops(EST_n, t_t)$. Since after line 321 we are sure that the $head$ is incremented (Lemma 12) and $n.head=h+1$ at t_t so the new block is established at t_t and the new block contains the new operations which is what we wanted to show. \square

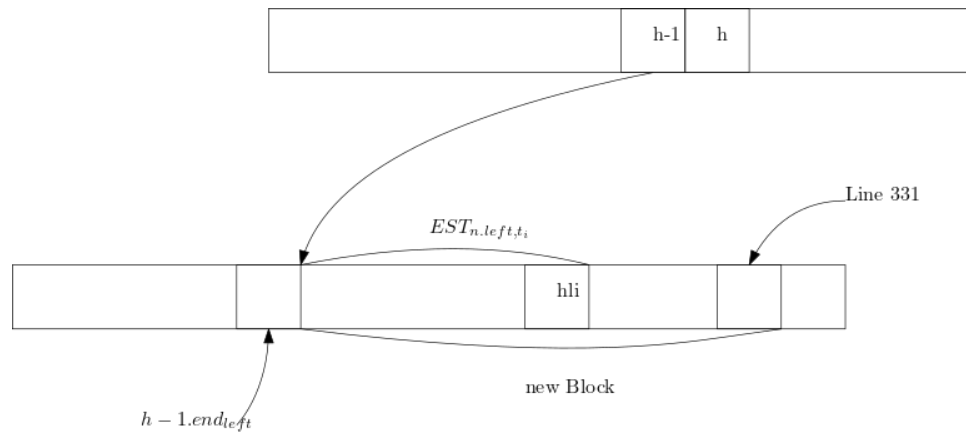


Figure 1: New established operations of the left child are in the new block.

Lemma 14 (Precise True Refresh).

Let t_i be the time an instance of $n.Refresh()$ read the head (Line 310) and t_t be the time its $TryAppend(new, s)$ terminates with and returns **true** (Line 313). We have $ops(EST_{n.left, t_i}) \cup ops(EST_{n.right, t_i}) \subseteq ops(n.blocks)$.

Lemma 15 (Double Refresh). *Consider two consecutive failed instances R_1, R_2 of $\mathbf{n.Refresh}()$ by some process. Let t_1 be the time R_1 is invoked and t_2 be the time R_2 terminated. We have $\text{ops}(\text{EST}_{\mathbf{n.left}}, t_1) \cup \text{ops}(\text{EST}_{\mathbf{n.right}}, t_1) \subseteq \text{ops}(\text{EST}_{\mathbf{n}}, t_2)$.*

Proof.

If Line 313 of R_1 or R_2 returns **true**, then the claim is held by Lemma 13. Let R_1 read i and R_2 read $i+1$ from Line 310. If R_2 reads some value greater than $i+1$ in Line 310 it means a successful instance of $\mathbf{Refresh}()$ started after Line 310 of R_1 and finished its Line 318 or 321 before 310 of R_2 , from Lemma 13 by the end of this instance $\text{ops}(\text{EST}_{\mathbf{n.left}}, t_1) \cup \text{ops}(\text{EST}_{\mathbf{n.right}}, t_1)$ has been propagated.

Since R_2 's $\mathbf{TryAppend}()$ returns **false** then there is another successful instance R'_2 of $\mathbf{n.Refresh}()$ that has done $\mathbf{TryAppend}()$ successfully into $\mathbf{n.blocks}[i+1]$ before R_2 tries to append. In Figure 1 we see why the block R'_2 is appending contains established block in the \mathbf{n} 's children at t_1 , since it create a block reading the head after t_1 . By Lemma 14 after R'_2 's CAS we have $\text{ops}(\text{EST}_{\mathbf{n.left}}, t_1) \cup \text{ops}(\text{EST}_{\mathbf{n.right}}, t_1) \subseteq \text{ops}(\mathbf{n.blocks})$. Also by Lemma 12 of R'_2 head is more than $i+1$ after R'_2 's 321 line, so the block appended by R'_2 to \mathbf{n} is established by then. To summarized t_1 is before R'_2 's read head and R'_2 's CAS is before R_2 's termination. So $\text{ops}(\text{EST}_{\mathbf{n.left}}, t_1) \cup \text{ops}(\text{EST}_{\mathbf{n.right}}, t_1) \subseteq \text{ops}(\text{EST}_{\mathbf{n}}, t_2)$. \square

last sentence need more detail and should be earlier. define i and tell why R_2 prime exists

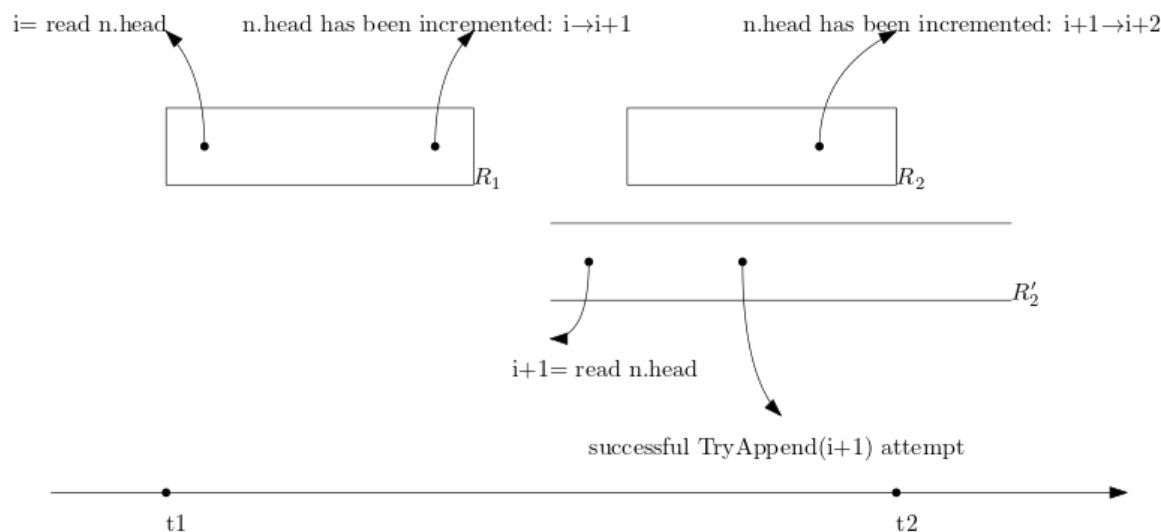


Figure 2: $t_1 < r_1$ reading head $<$ incrementing $\mathbf{n.head}$ from i to $i+1 < R'_2$ reading head $< \mathbf{TryAppend}(i+1) <$ incrementing $\mathbf{n.head}$ from $i+1$ to $i+2 < t_2$

this chain with more depth should be in the proof

Corollary 16 (Propagate Step). *All operations in `n`'s children's established blocks before line `302` are guaranteed to be in `n`'s established blocks after line `303`.*

Proof. Lines `302` and `303` satisfy the preconditions of Lemma `15`. □

Corollary 17. *After `Append(blk)` finishes $\text{ops}(\text{blk}) \subseteq \text{ops}(\text{root.blocks}[\text{x}])$ for some `x` and only one `x`.*

Proof. Follows from Lemma `15`, `9`. □

blockSize **Lemma 18** (Block Size Upper Bound). *Each block contains at most one operation from each process.*

Proof. By proof of contradiction, assume there are more than one operation from process p in block b in node n . A process cannot invoke more than one operations concurrently. From p 's operations in b , let op_1 be the first operation invoked and op_2 be the second one. Note that it is terminated before op_2 started. So before appending op_2 to the tree op_1 exists in every node from the path of p 's leaf to the root. So there is some block b' before b in n containing op_1 . op_1 existing in b and b' contradicts with append. \square

blocksBound **Lemma 19** (Subblocks Upperbound). *Each block has at most p direct subblocks.*

Proof. It follows directly from Lemma blockSize and the observation that each block contains at least one operation, induced from Line addOP. \square

Definition 20 (Ordering of operations inside the nodes). ► Note that processes are numbered from 1 to p , left to right in the leaves of the tree and from Lemma 118 we know there is at most one operation from each process in a given block.

- We call operations strictly before op in the sequence of operations S , prefix of the op .
- $E(n, b)$ is the sequence of enqueue operations $\in \text{ops}(\mathbf{n.blocks}[b])$ ordered by their process id.
- $E_{n,b,i}$ is the i th enqueue in $E(n, b)$.
- $D(n, b)$ is the sequence of dequeue operations $\in \text{ops}(\mathbf{n.blocks}[b])$ ordered by their process id.
- $D_{n,b,i}$ is the i th dequeue in $D(n, b)$.
- Order of the enqueue operations in n : $E(n) = E(n, 1).E(n, 2).E(n, 3)\dots$
- $E_{n,i}$ is the i th enqueue in $E(n)$.
- Order of the dequeue operations in n : $D(n) = D(n, 1).D(n, 2).D(n, 3)\dots$
- $D_{n,i}$ is the i th dequeue in $D(n)$.
- Linearization: $L = E(\text{root}, 1).D(\text{root}, 1).E(\text{root}, 2).D(\text{root}, 2).E(\text{root}, 3).D(\text{root}, 3)\dots$

Note that in the non-root nodes we only order enqueues and dequeues among the operations of their own type. Since `GetENQ()` only searches among enqueues and `IndexDEQ()` works on dequeues.

Preconditions of all invocation of **BSearch** are satisfied.

Lemma 21 (Get correctness). *If $n.\text{blocks}[b].\text{num}_{\text{enq}} \geq i$ then $n.\text{GetENQ}(b, i)$ returns $E_{n,b,i}$.*

Proof. We are going to prove this lemma by induction on the height of the tree. The base case for the leaves of the tree is pretty straight forward. Since leaf blocks contain exactly one operation then only $\text{GetENQ}(b, 1)$ can be called on leaves. $\text{leaf}.\text{GetENQ}(b, 1)$ returns the operation stored in the b th block of leaf l . For non leaf nodes in Line 404 it is decided that the i th enqueue in block b of internal node n resides in the left child or the right child of n . From Definition ordering 20 we know operations in a block are ordered by their process id. Furthermore $b.\text{sum}_{\text{enq-left}}$ stores the number of `enqueue()` operations from the b 's subblocks of the left child of n . So if i is greater than $b.\text{sum}_{\text{enq-left}}$ it means i th operation is propagated from the right child, otherwise we should search for the i th enqueue in the left child subblocks. By definition def::opsf::subblock 4 and 2 we need to search in subblocks of b which their range is $n.\text{left}.\text{blocks}[n.\text{blocks}[i-1].\text{end}_{\text{left}}+1..n.\text{blocks}[i].\text{end}_{\text{left}}] \cup n.\text{right}.\text{blocks}[n.\text{blocks}[i-1].\text{end}_{\text{right}}+1..n.\text{blocks}[i].\text{end}_{\text{right}}]$. If the enqueue we're looking for was in the right child as there are $b.\text{sum}_{\text{enq-left}}$ enqueues before it we need to search for $i-b.\text{sum}_{\text{enq-left}}$ (Line rightChildGet 409). By definition of $E(n, b)$ operations from the left child come before the operations of the right child. Having sum_{enq} , the prefix sum of the number of enqueues we can compute the direct subblock containing the enqueue we are finding for with binary search. Then $n.\text{child}.\text{GetENQ}(\text{block containing, order in the block})$ is invoked which returns the correct operation by the hypothesis of the induction. \square

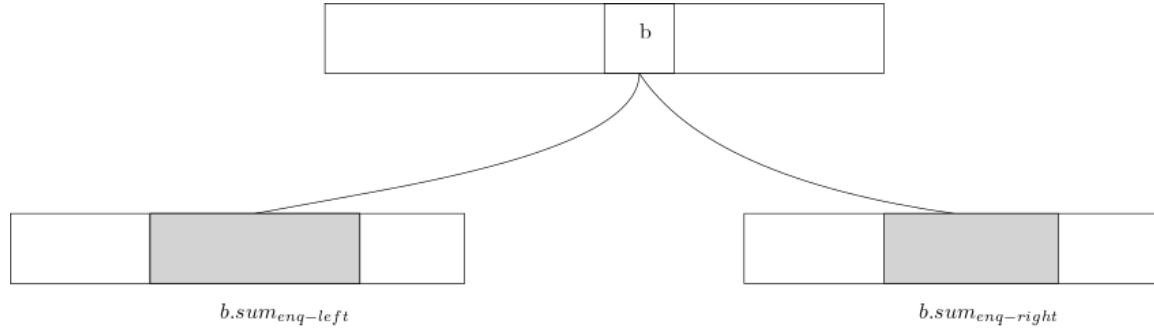


Figure 3: The number of enqueues from the left and the right child

I'm not sure it is going to be long and boring to talk about the parameters, since the reader can find out them.

dsearch

Lemma 22 (DSearch correctness). *If $\text{root.blocks}[\text{end}].\text{num}_{\text{enq}} \geq i$ and $E_{\text{root},i}$ is the response to some $\text{Dequeue}()$ in $\text{root.blocks}[\text{end}]$ then $\text{DSearch}(i, \text{end})$ returns b such that $\text{root.blocks}[b]$ contains $E_{\text{root},b,i}$ in $\Theta(\log(\text{root.blocks}[b].\text{size} + \text{root.blocks}[\text{end}].\text{size}))$ steps.*

Proof. First we show $\text{end} - b \leq \text{root.blocks}[b].\text{size} + \text{root.blocks}[\text{end}].\text{size}$. We know each block size is greater than 0. So every block in $\text{root.blocks}[b..\text{end}]$ contains at least one $\text{Enqueue}()$ or one $\text{Dequeue}()$. There cannot be more than $\text{root.blocks}[b].\text{size}$ $\text{Dequeue}()$ s in $\text{root.blocks}[b+1..\text{end}-1]$, since the queue would become empty after b th block end before end and $E(n, i)$ could not be the response to some DEQ in end . And since the length of the queue would become $\text{root.blocks}[\text{end}].\text{size}$ in the end so there cannot be more than $\text{root.blocks}[\text{end}].\text{size}$ Dequeus in $\text{root.blocks}[b..\text{end}]$. Cause if there was more then the end's length would become more than $\text{root.blocks}[\text{end}].\text{size}$.

Now that we know there are at most $\text{root.blocks}[b].\text{size} + \text{root.blocks}[\text{end}].\text{size}$ distance between end and b then with doubling search in $\log \text{root.blocks}[b].\text{size} + \text{root.blocks}[\text{end}].\text{size}$ steps we reach a block c that the $c.\text{sum}_{\text{enq}}$ is less than i and the distance between c and end is not more than $2 \times \text{root.blocks}[b].\text{size} + \text{root.blocks}[\text{end}].\text{size}$. So the binary search takes $\Theta(\log \text{root.blocks}[b].\text{size} + \text{root.blocks}[\text{end}].\text{size}))$ steps. \square

Lemma 23 (Index correctness). `n.IndexDEQ(b,i)` returns the rank in $D(\text{root})$ of $D_{n,b,i}$.

Proof. We will prove this by induction on the distance of `n` from the `root`. We can see the base case `root.IndexDEQ(b,i)` is trivial (Line indexBaseCase 415). In the non-root nodes `n.IndexDEQ(b,i)` computes the superblock of the i th Dequeue in the b th block of `n` in `n.parent` by Lemma superBlockcomputeSuper 24 (Line 418). After that the order in $D(n.\text{parent}, \text{superblock})$ is computed and `index()` is called on `n.parent` recursively. Then if the operation was propagated from the right child the number of dequeues from the left child are added to it (Line considerRight 420), because the left child operations come before the right child operations (Definition ordering 20). \square

Do I need to talk about the computation of the order in the parent which is based on the definition of ordering of dequeues in a block?

Make sure to show preconditions of all invocation of `BSearch` are satisfied.

Lemma 24 (Computing SuperBlock). After computing line computeSuper 418 of `n.IndexDEQ(b,i)`, `n.parent.blocks[superblock]` contains $D(n,b,i)$.

Proof. Lemmas 28,29,30,31. \square

Lemma 25. Value read for `super[b.group]` in line 418 is not null.

Proof. Values `np_dir` read in lines setNP 337, `super` are set before incrementing in lines setSuperNP 315,316. So before incrementing `num_propagated`, `super[num_propagated]` is set so it cannot be null while reading. \square

Lemma 26. `super[]` preserves order from child to parent; i.e. if in node `n` block `b` is before `c` then `b.group` \leq `c.group`

Proof. Line setGroup 329. Since `num_propagated` is increasing. \square

Lemma 27. Let `b`, `c` be in node `n`, if `b.group` \leq `c.group` then `super[b.group]` \leq `super[c.group]`

Proof. Line setSuper 315. \square

Lemma 28. The number of the blocks with `group=i` in a node is $\leq p$.

Lemma 29. `super[i+1]-super[i]` $\leq p$

Proof. In a Refresh with successful CAS in line 46, `super` and `counter` are set for each child in lines 48,49. Assume the current value of the counter in node `n` is `i+1` and still `super[i+1]` is not set. If an instance of successful `Refresh(n)` finishes `super[i+1]` is set a new value and a block is added after `n.parent[sup[i]]`. There could be at most p successful unfinished concurrent instances of `Refresh()` that have not reached line 49. So the distance between `super[i+1]` and `super[i]` is less than p . \square

Lemma 30 (super property). If `super[i] \neq null` in node `n`, then `super[i]` is the index of the superblock of a block with `time=i` in `n.parent.blocks`.

Lemma 31. Superblock of `b` is within range $\pm 2p$ of the `super[b.time]`.

Proof. `super[i]` is the index of the superblock of a block containing block `b`, followed by Lemma superCounter 30. `super(b)` is the real superblock of `b`. `super(t)` is the index of the superblock of the last block with time `t`. If `b.time` is `t` we have:

$$\text{super}[t] - p \leq \text{super}[t-1] \leq \text{super}(t-1) \leq \text{super}(b) \leq \text{super}(t+1) \leq \text{super}(t+1) \leq \text{super}[t] + p$$

\square

We call the dequeues that return some value *non-null dequeues*. r th non-null dequeue returns the element of the r th enqueue. We can compute # non-null dequeues in the prefix for a block this way: #non-null dequeues = size - #enqueues. Note that the i th dequeue in the given block is not a non-null dequeue.

computeHead

Lemma 32 (Computing Queue's Head block). *Let S be the state of an empty queue if the operations in prefix in L of i th dequeue in $D(\text{root}, b)$ are applied on it. `FindResponse()` returns (b, i) which $E(\text{root}, b, i)$ is the head of the queue in S . If the queue is empty in S then it returns $\langle -1, -- \rangle$.*

Proof. The size of the queue if the operations in the prefix for the b th block in the root are applied with the order of L is stored in the `root.blocks[b].size`. It is computed while creating the block in Line ^{computeLength}342. If the size of a queue is greater than 0 then a `Dequeue()` would decrease the size of the queue, otherwise the size of the queue remains 0. Having size of the queue after the previous block and number of enqueues and dequeues in the block, Line ^{computeLength}342 computes whether the queue becomes empty or the size of it.

$$r_{\text{enq}} = (i_d + \text{root.blocks}[b_d-1].\text{sum}_{\text{deq}}) - (\text{root.blocks}[b_d-1].\text{size} - \text{root.blocks}[b_d-1].\text{sum}_{\text{enq}} + \text{root.blocks}[b_d-1].\text{sum}_{\text{deq}})$$

□

HOW? How to prove mathematically that `ax(root.blocks[i-1].size + b.numenq - b.numdeq, 0)` is the size of the queue after the block. I can only explain it here.

Theorem 33 (Main). *The queue implementation is linearizable.*

Proof. We choose L in Definition ^{ordering}20 to be linearization ordering of operations and prove if we linearize operations as L the queue works consistently. □

Lemma 34. *Operations in a block have a time point in common (There is a time t all the operations are running).*

Lemma 35 (satisfiability). *L can be a linearization ordering.*

Proof. Once some operations are aggregated in one block they will be propagated together up to the root and we can linearize them in any order among themselves (previous lemma). Furthermore in L we arbitrary choose the order to be by process id, since it makes computations in the blocks faster. □

Lemma 36 (correctness). *If operations are applied as L on a sequential queue, the sequence of the responses would be the same as our algorithm.*

Proof. Old parts to review We show that the ordering L stored in the root, satisfies the properties of a linearizable ordering.

1. If op_1 ends before op_2 begins in E , then op_1 comes before op_2 in T .

► This is followed by Lemma ^{append}9. The time op_1 ends it is in root, before op_2 , by Definition ^{ordering}20 op_1 is before op_2 .

2. Responses to operations in E are same as they would be if done sequentially in order of L .

► Enqueue operations do not have any response so it does no matter how they are ordered. It remains to prove Dequeue d returns the correct response according to the linearization order. By Lemma ^{computeHead}32 it is deduced that the head of the queue at time of the linearization of d is computed properly. If the Queue is not empty by Lemma ^{get}21 we know that the returning response is the computed index element.

□

Lemma 37 (Amortized time analysis). *Enqueue() and Dequeue take $O(\log 62p + q)$ steps (amortized analysis), where p is the number of processes and q is the size of the queue at the time of invocation.*