

◇ *Local*

- **Node leaf*: pointer the the process's leaf in the tree

◇ *Shared*

- *Tree* : A binary tree of Nodes is shared among the processes. It can be implemented with a 1 index based array of size p . Such that the root is index 1, the left child and the right child of a node with index i are indices $2i$, $2i+1$ in the array.

◇ *Structures*► *Node*

- **Node left, right, parent*
- *Block[] blocks*: index 0 contains an empty block with all fields equal to 0 and *en* pointers to the first block of the corresponding children. *blocks[i]* returns the i th block stored. In the root node it is implemented with a persistent red-black tree and it is a big array in the other nodes.
- *int head= 1*: index of the first empty cell of blocks
- *int counter= 0*
- *int[] super*: *super[i]* stores the index of a superblock in parent that contains some block of this node whose time is field i

► *leaf extends Node*

- *int[] response*
leaf.response[i] stores response of leaf.ops[i]
- *int last_{done}*

Each process stores the index of the most recent block that the process has finished its last operation. An enqueue operation is finished if it has appended its element to the root and a dequeue operation is finished when it computes its response.

► *Block*

- *int num_{enq-left}, sum_{enq-left}* : #enqueues from subblocks in left child, prefix sum of *num_{enq-left}*
- *int num_{deq-left}, sum_{deq-left}* : #dequeues from subblocks in left child, prefix sum of *num_{deq-left}*
- *int num_{enq-right}, sum_{enq-right}* : #enqueues from subblocks in right child, prefix sum of *num_{enq-right}*
- *int num_{deq-right}, sum_{deq-right}* : #dequeues from subblocks in right child, prefix sum of *num_{deq-right}*
- *int num_{enq}, num_{deq}* : # enqueue, dequeue operations in the block
- *int sum_{enq}, sum_{deq}* : sum of # enqueue, dequeue operations in blocks up to this one
- *int num, sum* : total # operations in block, prefix sum of *num*
- *int end_{left}, end_{right}* : index of the last subblock in the left and right child
- *int group* : id of the group of blocks including this propagated together, more precisely the value read from the node n 's counter when propagating this block to the node n .

► *Leaf Block extends Block*

- *Object element* Each block in a leaf represents an operation. The element shows the operation's argument if it is an enqueue, and if it is a dequeue this value is null.

► *Root Block extends Block*

- *int size* : size of queue after this block's operations finish
- *int sum_{non-null deq}* : count of non-null dequeues up to this block
- *int num_{finished}* : number of finished operations in the block
- *int order* : the index of the block in the node containing it. Useful in the root since in the PBRT we do not keep indices in another way.

```

34: void ENQUEUE(Object e) ▷ Creates a block with element e and appends it
    to the tree.
35:   block b= NEW(leaf block)
36:   b.element= e
37:   b.numenq=1
38:   b.sumenq= this.leaf.blocks[this.leaf.head].sumenq+1
39:   APPEND(b)
40: end ENQUEUE

41: Object DEQUEUE()
42:   block b= NEW(leaf block) ▷ Creates a new block, appends it to the tree, computes its order among operations, then
    computes its response index if it exists and returns the response's element.
43:   b.element= null
44:   b.numdeq=1
45:   b.sumdeq= this.leaf.blocks[this.leaf.head].sumdeq+1 ▷ this is the
    current running process
46:   APPEND(b)
47:   <i, bi>= INDEX(this.leaf, this.leaf.head, 1) ▷ i is the
    order in the root among all dequeues, of the dequeue in the last block in the
    process's leaf. bi is the index of the block in the root containing it. Since
    only one invocation of dequeue is running by this process at one time we
    are allowed to used this.leaf.head safely.
48:   indexresponse= COMPUTEDEQRES(i, b) ▷ indexresponse is the index of the
    enqueue which is the response to the dequeue or -1 if the response is null.
49:   if indexresponse!=-1 then
50:     output= null
51:     bdeq=root.blocks[bi]
52:     bdeq.numfinished.inc() ▷ shared counter
53:     if bdeq.numfinished==bdeq.num then ▷ all the operations in the block
        containing the dequeue are finished.
54:       this.leaf.lastdone= bi
55:     end if
56:   else
57:     output= GET(indexresponse)
58:     br= root.blocks.get(enq, indexresponse).order ▷ index of the
        block in the root contains response enqueue.
59:     benq=root.blocks[br]
60:     benq.numfinished.inc()
61:     bdeq.numfinished.inc()
62:     if benq.numfinished==benq.num then ▷ become done
63:       this.leaf.lastdone= br
64:     else if bdeq.numfinished==bdeq.num then ▷ bdeq comes after benq.
65:       this.leaf.lastdone= bi ▷ this.leaf.lastdone is an increasing
        value.
66:     end if
67:   end if
68:   return output
69: end DEQUEUE

70: void APPEND(block b)
71:   b.group= this.leaf.head ▷ Only this block is propagated from the
    leaf by itself.
72:   this.leaf.blocks[this.leaf.head]= b
73:   this.leaf.head+=1 ▷ Lines 71 to 73 are done by one process at time.
74:   PROPAGATE(this.leaf.parent)
75: end APPEND

76: void PROPAGATE(node n)
77:   if not REFRESH(n) then
78:     REFRESH(n) ▷ Lemma Double Refresh
79:   end if
80:   if n is not root then ▷ To check anode is the root we can check
        its index if the tree is implemented by an array or check if n.parent is not
        null.
81:     PROPAGATE(n.parent)
82:   end if
83: end PROPAGATE

84: element GET(int i) ▷ Returns ith Enqueue.
85:   res= root.blocks.get(enq, i).order
86:   return GET(root, res, i-root.blocks[res-1].sumenq)
87: end GET

88: int COMPUTEDEQRES(int i, int b) ▷ Computes the response of
    the ith dequeue in the root's bth block. Returns the index of the the head
    of the queue or -1 if queue is empty.
89:   if root.blocks[b-1].size + root.blocks[b].numenq - i < 0 then
90:     return -1
91:   else return root.blocks[b-1].sumnon-null deq + i
92:   end if
93: end COMPUTEDEQRES

```

<pre> 34: boolean REFRESH(node n) 35: h= n.head 36: c= n.counter 37: <new, cleft, cright>= CREATEBLOCK(n, h) ▷ cleft, cright are the values read from n's children's counters. 38: new.group= c 39: if new.num==0 then return true ▷ The block contains nothing. 40: else if (n is root and root.blocks.append(new)) or 41: (n is not root and CAS(n.blocks[h], null, new)) then ▷ how to put space in the first of the new line? okcas42: for each dir in {left, right} do 43: CAS(n.dir.super[cdir], null, h+1) ▷ Superblock's Lemma 44: CAS(n.dir.counter, cdir, cdir+1) 45: end for 46: CAS(n.head, h, h+1) 47: return true 48: else 49: CAS(n.head, h, h+1) ▷ Even if another process wins, help to increase the head. It might fall sleep before increasing. 50: return false 51: end if 52: end REFRESH ~ Precondition: n.blocks[start..end] contains a block with field f ≥ i 53: int BSEARCH(node n, field f, int i, int start, int end) ▷ Does binary search for the value i of the given prefix sum field. Returns the index of the leftmost block in n.blocks[start..end] whose field f is ≥ i. 54: end BSEARCH </pre>	<pre> 55: <Block, int, int> CREATEBLOCK(node n, int i) ▷ Creates a block to insert into n.blocks[i]. Returns the created block as well as values read from each child counter field. The values are used for incrementing children's counters if the block was appended to n.blocks successfully. Does it need help? I think no but in that case we do not have to pass these values to the calling line. 56: block b= NEW(block) 57: if n is root then 58: b= NEW(root block) 59: end if 60: b.order= i 61: for each dir in {left, right} do lastLine62: indexlast= n.dir.head prevLine63: indexprev= n.blocks[i-1].enddir 64: blocklast= n.dir.blocks[indexlast] 65: blockprev= n.dir.blocks[indexprev] ▷ n.dir.blocks[indexprev..indexlast] are merged to one block. 66: cdir= n.dir.counter 67: b.enddir= indexlast 68: b.numenq-dir= blocklast.sumenq - blockprev.sumenq 69: b.numdeq-dir= blocklast.sumdeq - blockprev.sumdeq 70: b.sumenq-dir= n.blocks[i-1].sumenq-dir + b.numenq-dir 71: b.sumdeq-dir= n.blocks[i-1].sumdeq-dir + b.numdeq-dir 72: end for 73: b.numenq= b.numenq-left + b.numenq-right 74: b.numdeq= b.numdeq-left + b.numdeq-right 75: b.num= b.numenq + b.numdeq 76: b.sum= n.blocks[i-1].sum + b.num 77: if n is root then 78: b.size= max(root.blocks[i-1].size + b.numenq - b.numdeq, 0) 79: b.sumnon-null deq= root.blocks[i-1].sumnon-null deq + max(b.numdeq - root.blocks[i-1].size - b.numenq, 0) 80: end if 81: return b, cleft, cright 82: end CREATEBLOCK </pre>
----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

↪ Precondition: $n.blocks[b]$ contains $\geq i$ enqueues.

```

84: element GET(node  $n$ , int  $b$ , int  $i$ )                                ▷ Returns the  $i$ th Enqueue in  $b$ th block of node  $n$ 
85:   if  $n$  is leaf then return  $n.blocks[b].element$ 
86:   else
87:     if  $i \leq n.blocks[b].num_{enq-left}$  then                                ▷  $i$  exists in the left child of  $n$ 
88:       subBlock= BSEARCH( $n.left$ ,  $sum_{enq}$ ,  $i$ ,  $n.blocks[b-1].end_{left}+1$ ,  $n.blocks[b].end_{left}$ )
89:       return GET( $n.left$ , subBlock,  $i-n.left.blocks[subBlock-1].sum_{enq}$ )
90:     else
91:        $i = i - n.blocks[b].num_{enq-left}$ 
92:       subBlock=BSEARCH( $n.right$ ,  $sum_{enq}$ ,  $i$ ,  $n.blocks[b-1].end_{right}+1$ ,  $n.blocks[b].end_{right}$ )
93:       return GET( $n.right$ , subBlock,  $i-n.right.blocks[subBlock-1].sum_{enq}$ )
94:     end if
95:   end if
96: end GET

```

↪ Precondition: b th block of node n has propagated up to the root and i th dequeue resides in node n is in block b of node n .

```

97: <int, int> INDEX(node  $n$ , int  $b$ , int  $i$ )                                ▷ Returns the order in the root of  $i$ th dequeue in the  $b$ th block of node  $n$  among dequeues.
98:   if  $n$  is root then return  $i$ ,  $b$ 
99:   else
100:    dir= ( $n.parent.left == n$ )? left: right                                ▷ check  $n$  is a left or a right child
101:    superBlock= BSEARCH( $n.parent$ ,  $n.sum_{deq-dir}$ ,  $i$ ,  $super[n.blocks[b].group]-p$ ,  $super[n.blocks[b].group]+p$ )    ▷ superblock's group has at
    most  $p$  difference with the value stored in  $super[]$ .
102:    if dir is left then
103:       $i += n.parent.blocks[superBlock-1].sum_{deq-right}$ 
104:    else
105:       $i += n.parent.blocks[superBlock-1].sum_{deq} + n.blocks[superBlock].sum_{deq-left}$     ▷ consider dequeues from  $n$ 's right child
106:    end if
107:    return INDEX( $n.parent$ , superBlock,  $i$ )
108:  end if
109: end INDEX

```

```

► PRBTree[rootBlock]
A persistant red-black tree supporting append(b, key),get(key=i),split(j).
append(b, key) returns true in case successful. Since order, sumenqare
both strictly increasing we can use one of them for another.
1: void RBTAPPEND(block b)          ▷ adds block b to the root.blocks
2:   step= root.head
3:   if step%p2==0 then  ▷ Help every often p2 operations appended to the
   root. Used in lemma's using the size of the PBRT.
4:     Help()
5:     CollectGarbage()
6:   end if
7:   b.numfinished= 0
8:   return root.blocks.append(b, b.order)
9: end RBTAPPEND

10: void HELP          ▷ Helps pending operations
11:   for leaf l in leaves do  ▷ if the tree is implemented with an array we
   can iterate over the second half of the array.
12:     last= l.head-1 ▷ l.blocks[last] can not be null because of lines
   11-13.
13:     if l.blocks[last].element==null then  ▷ operation is dequeue
14:       goto 48 with these values <>      ▷ run Dequeue() for
   l.ops[last] after Propagate(). TODO
15:     l.responses[last]= response

16:   end if
17: end for
18: end HELP

19: void COLLECTGARBAGE          ▷ Collects the root blocks that are done.
20:   s=FindMostRecentDone(Root.Blocks.root)  ▷ Lemma: If block b is
   done after helping then all blocks before b are done as well.
21:   t1,t2= RBT.split(order, s)
22:   RBTRoot.CAS(t2.root)
23: end COLLECTGARBAGE

24: Block FINDMOSTRECENTDONE(b)
25:   for leaf l in leaves do
26:     max= Max(l.maxOld, max)
27:   end for
28:   return max          ▷ This snapshot suffies.
29: end FINDYOUNGESTOLD

30: response FALLBACK(op i)          ▷ really necessary?
31:   if root.blocks.get(numenq), i is null then  ▷ this enqueue was already
   finished
32:     return this.leaf.response(block.order)
33:   end if
34: end FALLBACK

```
