

◇ *Shared*

- *Node** *tree[]* : A binary tree of Nodes such that *Tree[0]* is the root and the left child and the right child and the parent of *Tree[i]* are *Tree[2i+1]*, *Tree[2i+2]* and *Tree[i/2]*.

treeRules

◇ *Local*

- *int* *leaf*: *tree[leaf]* is the process's leaf in the tree.

◇ *Structures*

► *Node*

- **Node* *left()*, *right()*, *parent()* : initialized when creating the tree as mentioned in the *tree[]* description ^{treeRules} 1100.
- *BlockList* *blocks*: Supports two operations *blocks.tryAppend(Block b)*, *blocks[i]*. Initially empty, when *blocks.tryAppend(b)* returns true *b* is appended to the end of the list and *blocks[i]* returns *i*th block in the blocks. If some instance of *blocks.tryAppend(b)* returns false there is a concurrent instance of *blocks.tryAppend(b')* which has returned true. *blocks[0]* contains an empty block with all fields equal to 0 and end pointers to the first block of the corresponding children.
- *int* *numpropagated*= 0 : # groups of blocks have been propagated from the node. It may be behind its real value.
- *int[]* *super*: *super[i]* stores the index of the superblock of some block in blocks that its group field is *i*.

► *Root* extends *Node*

- *PBRT* *blocks*
Implemented with a persistent red-black tree.

► *NonRootNode* extends *Node*

- *Block[]* *blocks*
Implemented with an array and using CAS for appending to the head.
- *int* *head*= 1: #blocks in the blocks.

► *Leaf* extends *NonRootNode*

- *int* *lastdone*
Each process stores the index of the most recent block in the root that the process has finished the last operation of the block. *enqueue(e)* is finished if *e* is returned by some *dequeue()* and *dequeue()* is finished when it computes its response.

► *Block* *b* ▷ If *b* is *blocks[i]* (*i*!=0) then *b[-1]* is *blocks[i-1]*.

- *int* *group* : the value read from the node *n.numpropagated* when appending this block to *n*.

► *LeafBlock* extends *Block*

- *Object* *element* : Each block in a leaf represents an operation. The element shows the operation's argument if it is an enqueue, and if it is a dequeue this value is null.
- *int* *sumenq*, *sumdeq* : number of enqueue, dequeue operations in the leaf's containing this block till this block
- *Object* *response*
response stores the response of the operation in the LeafBlock.

► *InternalBlock* extends *Block*

- *int* *sumenq-left* : #enqueues from the subblocks in the left child + *b[-1].sumenq-left*
- *int* *sumdeq-left* : #dequeues from the subblocks in the left child + *b[-1].sumdeq-left*
- *int* *sumenq-right* : #enqueues from the subblocks in the right child + *b[-1].sumenq-right*
- *int* *sumdeq-right* : #dequeues from the subblocks in the right child + *b[-1].sumdeq-right*
- *int* *endleft*, *endright* : index of the last subblock of the block in the left and right child

► *RootBlock* extends *InternalBlock*

- *int* *size* : size of the queue after this block's operations finish
- *counter* *numfinished* : number of finished operations in the block
- *int* *order* : the index of the block in the node containing the block. Useful in the root since in the PBRT we do not keep indices as key.

Conventions

- *b_x*: the block containing *x*
- *r_x*: rank of *x* in the current scope

Unwritten rules

- *blocks[b].sum_x*=*blocks[b].sum_{x-left}*+*blocks[b].sum_{x-right}* (for *b*>0 and *x* ∈ {enq, deq})
- *blocks[b].sum*=*blocks[b].sum_{enq}*+*blocks[b].sum_{deq}* (for *b*>0)
- *blocks[b].num_x*=*blocks[b].sum_x*-*blocks[b-1].sum_x*
(for *b*>0 and *x* ∈ {∅, enq, deq, enq-left, enq-right, deq-left, deq-right})

Algorithm Queue

```

201: void ENQUEUE(Object e)  ▷ Creates a block with element e and appends
    it to the tree.
202:   block newBlock= NEW(LeafBlock)
203:   newBlock.element= e
204:   newBlock.sumenq= tree[leaf].blocks[tree[leaf].head].sumenq+1
205:   newBlock.sumdeq= tree[leaf].blocks[tree[leaf].head].sumdeq
206:   tree[leaf].APPEND(newBlock)
207: end ENQUEUE

208: Object DEQUEUE()
209:   block newBlock= NEW(LeafBlock)  ▷ Creates a block with null value
    element, appends it to the tree, computes its order among operations, then
    computes its response; if it exists returns the response's element.
210:   newBlock.element= null
211:   newBlock.sumenq= tree[leaf].blocks[tree[leaf].head].sumenq
212:   newBlock.sumdeq= tree[leaf].blocks[tree[leaf].head].sumdeq+1
213:   tree[leaf].APPEND(newBlock)
214:   <rdeq, brdeq>= tree[leaf].INDEXDEQ(tree[leaf].head, 1)
    ▷ r is the rank among all dequeues of the dequeue and brdeq is the index of
    the block in the root containing the dequeue.
215:   <renq, brenq>= FINDRESPONSE(rdeq, brdeq)
    ▷ renq is the rank of the enqueue which is the response to the dequeue or
    -1 if the response is null.
216:   if renq== -1 then
217:     output= null
218:     tree[0].blocks[brdeq].numfinished.inc()  ▷ shared counter
219:     if tree[0].blocks[brdeq].numfinished==tree[0].blocks[brdeq].num
        then ▷ all the operations in the block containing the dequeue are finished.
220:       tree[leaf].lastdone= brdeq
221:     end if
222:   else
223:     output= GETENQ(renq, brenq)
224:     root.blocks[brenq].numfinished.inc()
225:     root.blocks[brenq].numfinished.inc()
226:     if root.blocks[br].numfinished==root.blocks[br].num then
227:       tree[leaf].lastdone= brenq
228:     else if root.blocks[brdeq].numfinished==root.blocks[brdeq].num
        then ▷ root.blocks[br] comes after root.blocks[bi].
229:       tree[leaf].lastdone= brdeq
230:     end if
231:   end if
232:   return output
233: end DEQUEUE

234: int, int FINDRESPONSE(int i, int b)  ▷ Computes the rank and
    index of the block in the root of the enqueue that is the response of the ith
    dequeue in the root's bth block. Returns j-1,0 if the queue is empty.
235:   if root.blocks[b-1].size + root.blocks[b].numenq - i < 0 then
236:     return -1, 0
237:   else
238:     response= root.blocks[b-1].sumdeq- root.blocks[b-1].size +
        i
239:     return <res, tree[0].blocks.get(enq, response).order>
240:   end if
241: end FINDRESPONSE

```

Algorithm Node

<pre> 301: void PROPAGATE() 302: if not this.REFRESH() then 303: this.REFRESH() ▷ Lemma Double Refresh 304: end if 305: if this is not root then ▷ To check a node is the root we can check its index is 0. 306: this.parent.PROPAGATE() 307: end if 308: end PROPAGATE 309: boolean REFRESH() 310: h= head 311: <new, npleft, npright>= this.CREATEBLOCK(h) ▷ npleft, npright are the values read from the children's numpropagateds. 312: if new.num==0 then return true ▷ The block contains nothing. 313: else if root.blocks.tryAppend(new) then 314: for each dir in {left, right} do 315: CAS(dir.super[npdir], null, h+1) 316: CAS(dir.numpropagated, npdir, npdir+1) 317: end for 318: CAS(head, h, h+1) 319: return true 320: else 321: CAS(head, h, h+1) ▷ Even if another process wins, help to increase the head. It might fell sleep before increasing. 322: return false 323: end if 324: end REFRESH </pre>	<pre> 327: <Block, int, int> CREATEBLOCK(int i) ▷ Creates a block to be inserted into this.blocks[i]. Returns the created block as well as values read from each childnumpropagated field. The values are used for incrementing children's numpropagateds if the block was appended to this.blocks successfully. 328: block newBlock= NEW(block) 329: newBlock.group= numpropagated 330: newBlock.order= i 331: for each dir in {left, right} do 332: indexlast= dir.head 333: indexprev= blocks[i-1].enddir 334: blocklast= dir.blocks[indexlast] 335: blockprev= dir.blocks[indexprev] 336: ▷ newBlock includes dir.blocks[indexprev+1..indexlast]. 337: ndir= dir.numpropagated 338: newBlock.enddir= indexlast 339: newBlock.sumenq-dir= blocks[i-1].sumenq-dir + blocklast.sumenq - blockprev.sumenq 340: newBlock.sumdeq-dir= blocks[i-1].sumdeq-dir + blocklast.sumdeq - blockprev.sumdeq 341: end for 342: if this is root then 343: newBlock.size= max(root.blocks[i-1].size + b.numenq - b.numdeq, 0) 344: end if 345: return b, npleft, npright 346: end CREATEBLOCK </pre>
--	--

↪ Precondition: blocks[start..end] contains a block with field $f \geq i$

```

325: int BSEARCH(field f, int i, int start, int end)
                                ▷ Does binary search for the value
                                i of the given prefix sum feild. Returns the index of the leftmost block in
                                blocks[start..end] whose field f is  $\geq i$ .
326: end BSEARCH

```

Algorithm Node

↪ Precondition: $n.blocks[b]$ contains i th enqueue in the node.

```
401: element GETENQ(int b, int i)
402:   if  $i \leq blocks[b].num_{enq-left}$  then                                ▷  $i$  exists in the left child of this node
403:     subBlock= leftBSEARCH(sumenq, i, blocks[b-1].endleft+1, blocks[b].endleft)
404:     return left.GET(i-left.blocks[subBlock-1].sumenq, subBlock)
405:   else
406:     i= i-blocks[b].numenq-left
407:     subBlock= rightBSEARCH(sumenq, i, blocks[b-1].endright+1, blocks[b].endright)
408:     return right.GET(i-right.blocks[subBlock-1].sumenq, subBlock)
409:   end if
410: end GETENQ
```

↪ Precondition: b th block of the node has propagated up to the root and i th dequeue in the node is in $blocks[b]$.

```
411: <int, int> INDEXDEQ(int b, int i)                                ▷ Returns the order in the root of  $i$ th dequeue in the  $b$ th block of node  $n$  among dequeues.
412:   dir= (parent.left==n)? left: right                                ▷ check if a left or a right child
413:   superBlock= BSEARCH(parent, sumdeq-dir, i, super[blocks[b].group]-p, super[blocks[b].group]+p)
                                                                ▷ superblock's group has at most  $p$  difference with the value stored in super[].
414:   if dir is left then
415:     i+= parent.blocks[superBlock-1].sumdeq-right
416:   else
417:     i+= parent.blocks[superBlock-1].sumdeq + blocks[superBlock].sumdeq-left                                ▷ consider dequeues from the right child
418:   end if
419:   return INDEXDEQ(this.parent, superBlock, i)
420: end INDEX
```

Algorithm Leaf

```
501: void APPEND(block blk)
502:   head+=1                                ▷ Lines 503 to 502 are done by one process at time.
503:   blk.group= head                                ▷ Append is only called by the owener of the leaf.
504:   blocks[head]= blk
505:   parent.PROPAGATE()
506: end APPEND

507: element GETENQ(int b, int i)
508:   return blocks[b].element
509: end GETENQ
```

Algorithm Root

```
601: <int, int> INDEXDEQ(int b, int i)
602:   return <i, b>
603: end INDEXDEQ
```

appendEnd
pendStart

appendEnd

```

► PRBTree[rootBlock]
A persistant red-black tree supporting append(b, key),get(key=i),split(j).
append(b, key) returns true in case successful. Since order, sum_enqueue
both strictly increasing we can use one of them for another.
701: void RBTAPPEND(block b)          ▷ adds block b to the root.blocks
702:   step= root.head
703:   if step%p2==0 then ▷ Help every often p2 operations appended to the
      root. Used in lemma's using the size of the PBRT.
704:     Help()
705:     CollectGarbage()
706:   end if
707:   b.numfinished= 0
708:   return root.blocks.append(b, b.order)
709: end RBTAPPEND

710: void HELP          ▷ Helps pending operations
711:   for leaf l in leaves do ▷ if the tree is implemented with an array we
      can iterate over the second half of the array.
712:     last= l.head-1    ▷ l.blocks[last] can not be null because of
      appendSplitAnd
      lines 503-502.
713:     if l.blocks[last].element==null then ▷ operation is dequeue
714:       goto DeqRest
      215 with these values <> ▷ run Dequeue() for
      l.ops[last] after Propagate(). TODO
715:       l.responses[last]= response

716:   end if
717: end for
718: end HELP

719: void COLLECTGARBAGE          ▷ Collects the root blocks that are done.
720:   s=FindMostRecentDone(Root.Blocks.root) ▷ Lemma: If block b is
      done after helping then all blocks before b are done as well.
721:   t1,t2= RBT.split(order, s)
722:   RBTRoot.CAS(t2.root)
723: end COLLECTGARBAGE

724: Block FINDMOSTRECENTDONE(b)
725:   for leaf l in leaves do
726:     max= Max(l.maxOld, max)
727:   end for
728:   return max          ▷ This snapshot suffies.
729: end FINDYOUNGESTOLD

730: response FALLBACK(op i)          ▷ really necessary?
731:   if root.blocks.get(num_enqueue), i is null then ▷ this enqueue was already
      finished
732:     return this.leaf.response(block.order)
733:   end if
734: end FALLBACK

```
