

1 Pseudocode

Algorithm Tree Fields Description

◇ *Shared*

- A binary tree of Nodes with one leaf for each process. root is the root node.

◇ *Local*

- *Node* leaf: process's leaf in the tree.

◇ *Structures*

► *Node*

- **Node* left, right, parent : initialized when creating the tree.
- *BlockList*
- *int* head= 1: #blocks in blocks. blocks[0] is a block with all integer fields equal to zero.
- *int* num_{propagated}= 0 : # groups of blocks that have been propagated from the node to its parent. Since it is incremented after propagating, it may be behind by 1.

► *Block*

- *int* group : the value read from num_{propagated} when appending this block to the node.

► *LeafBlock* extends *Block*

- *Object* element : Each block in a leaf represents a single operation. If the operation is enqueue(x) then element=x, otherwise element=null.
- *int* sum_{enq}, sum_{deq} : # enqueue, dequeue operations in the prefix for the block

► *InternalBlock* extends *Block*

- *int* end_{left}, end_{right} : indices of the last subblock of the block in the left and right child
- *int* sum_{enq-left} : # enqueue operations in the prefix for left.blocks[end_{left}]
- *int* sum_{deq-left} : # dequeue operations in the prefix for left.blocks[end_{left}]
- *int* sum_{enq-right} : # enqueue operations in the prefix for right.blocks[end_{right}]
- *int* sum_{deq-right} : # dequeue operations in the prefix for right.blocks[end_{right}]

► *RootBlock* extends *InternalBlock*

- *int* size : size of the queue after performing all operations in the prefix for this block
-

Abbreviations:

- blocks[b].sum_x=blocks[b].sum_{x-left}+blocks[b].sum_{x-right} (for b≥0 and x ∈ {enq, deq})
- blocks[b].sum=blocks[b].sum_{enq}+blocks[b].sum_{deq} (for b≥0)
- blocks[b].num_x=blocks[b].sum_x-blocks[b-1].sum_x
(for b>0 and x ∈ {∅, enq, deq, enq-left, enq-right, deq-left, deq-right})

Algorithm Queue

```
201: void ENQUEUE(Object e) ▷ Creates a block with element e and adds it to the tree.
202:   block newBlock= NEW(LeafBlock)
203:   newBlock.element= e
204:   newBlock.sumenq= leaf.blocks[leaf.head].sumenq+1
205:   newBlock.sumdeq= leaf.blocks[leaf.head].sumdeq
206:   leaf.APPEND(newBlock)
207: end ENQUEUE

208: Object DEQUEUE() ▷ Creates a block with null value element, appends it to the tree, computes its order among operations, and returns its response.
209:   block newBlock= NEW(LeafBlock)
210:   newBlock.element= null
211:   newBlock.sumenq= leaf.blocks[leaf.head].sumenq
212:   newBlock.sumdeq= leaf.blocks[leaf.head].sumdeq+1
213:   leaf.APPEND(newBlock)
214:   <b, i>= INDEXDEQ(leaf.head, 1)
215:   output= FINDRESPONSE(b, i)
216:   return output
217: end DEQUEUE

218: <int, int> FINDRESPONSE(int b, int i)
219:   if root.blocks[b-1].size + root.blocks[b].numenq - i < 0 then
220:     return null
221:   else
222:     e= i - root.blocks[b-1].size + root.blocks[b-1].sumenq
223:     return root.GetENQ(root.DSEARCH(e, b))
224:   end if
225: end FINDRESPONSE
```

deqRest

checkEmpty

computeE

findAnswer

Algorithm Node

```
301: void PROPAGATE()
302:   if not REFRESH() then
303:     REFRESH()
304:   end if
305:   if this is not root then
306:     parent.PROPAGATE()
307:   end if
308: end PROPAGATE

309: boolean REFRESH()
310:   h= head
311:   <new, npleft, npright>= CREATEBLOCK(h)  ▷ npleft, npright are the
values read from the children's numpropagated field.
312:   if new.num==0 then return true  ▷ The block contains nothing.
313:   else if blocks.tryAppend(new, h) then
314:     for each dir in {left, right} do
315:       CAS(dir.super[npdir], null, h)  ▷ Write would work too.
316:       CAS(dir.numpropagated, npdir, npdir+1)
317:     end for
318:     CAS(head, h, h+1)
319:     return true
320:   else
321:     CAS(head, h, h+1)  ▷ Even if another process wins, help
to increase the head. The winner might have fallen sleep before increasing
head.
322:     return false
323:   end if
324: end REFRESH

325: int BSEARCH(field f, int i, int start, int end)
▷ Does binary search for the value
i of the given prefix sum field. Returns the index of the leftmost block in
blocks[start..end] whose field f is ≥ i.
326: end BSEARCH
```

↪ Precondition: blocks[start..end] contains a block with field $f \geq i$

```
325: int BSEARCH(field f, int i, int start, int end)
▷ Does binary search for the value
i of the given prefix sum field. Returns the index of the leftmost block in
blocks[start..end] whose field f is ≥ i.
326: end BSEARCH
```

```
327: <Block, int, int> CREATEBLOCK(int i)  ▷ Creates a block
to be inserted as  $i$ th block in blocks. Returns the created block as well as
values read from each child's numpropagated field. These values are used for
incrementing the children's numpropagated field if the block was appended to
blocks successfully.
328:   block newBlock= NEW(block)
329:   newBlock.group= numpropagated
330:   for each dir in {left, right} do
331:     indexlast= dir.head-1
332:     indexprev= blocks[i-1].enddir
333:     newBlock.enddir= indexlast
334:     blocklast= dir.blocks[indexlast]
335:     blockprev= dir.blocks[indexprev]
336:     ▷ newBlock includes dir.blocks[indexprev+1..indexlast].
337:     npdir= dir.numpropagated
338:     newBlock.sumenq-dir= blocks[i-1].sumenq-dir + blocklast.sumenq
- blockprev.sumenq
339:     newBlock.sumdeq-dir= blocks[i-1].sumdeq-dir + blocklast.sumdeq
- blockprev.sumdeq
340:   end for
341:   if this is root then
342:     newBlock.size = max(root.blocks[i-1].size + newBlock.numenq
- newBlock.numdeq, 0)
343:   end if
344:   return <b, npleft, npright>
345: end CREATEBLOCK
```

Algorithm Root

↪ Precondition: root.blocks[end].sum_{enq} ≥ e

```
801: <int, int> DSEARCH(int e, int end)  ▷ Returns <b,i> if  $E_{root,e} = E_{root,b,i}$ .
802:   start= end-1
803:   while root.blocks[start].sumenq ≥ e do
804:     start= max(start-(end-start), 0)
805:   end while
806:   b= root.BSearch(sumenq, e, start, end)
807:   i= e- root.blocks[b-1].sumenq
808:   return <b,i>
809: end DSEARCH
```

Algorithm Node	
	<pre> \rightsquigarrow Precondition: <code>blocks[b].num_{enq} ≥ i ≥ 1</code> </pre>
	<pre> 401: <i>element</i> GETENQ(<i>int</i> b, <i>int</i> i) ▷ Returns the element of $E_{this,b,i}$. 402: if this is leaf then 403: return blocks[b].element 404: else if $i \leq \text{blocks}[b].\text{num}_{\text{enq-left}}$ then ▷ $E_{this,b,i}$ is in the left child of this node. 405: subBlock= left.BSEARCH(sum_{enq}, i+blocks[b-1].sum_{enq-left}, blocks[b-1].end_{left}+1, blocks[b].end_{left}) 406: return left.GETENQ(subBlock, i) 407: else 408: i= i-blocks[b].num_{enq-left} 409: subBlock= right.BSEARCH(sum_{enq}, i+right.blocks[b-1].sum_{enq-right}, blocks[b-1].end_{right}+1, blocks[b].end_{right}) 410: return right.GETENQ(subBlock, i) 411: end if 412: end GETENQ </pre>
tBaseCase	
ftOrRight	
tChildGet	
tChildGet	
	<pre> \rightsquigarrow Precondition: bth block of the node has propagated up to the root and <code>blocks[b].num_{enq} ≥ i</code>. </pre>
	<pre> 413: <<i>int</i>, <i>int</i>> INDEXDEQ(<i>int</i> b, <i>int</i> i) ▷ Returns <x, y> if $D_{this,b,i} = D_{root,x,y}$. 414: if this is root then 415: return <b, i> 416: else 417: dir= (parent.left==n)? left: right ▷ check if this node is a left or a right child 418: superBlock= parent.BSEARCH(sum_{deq-dir}, i+blocks[b-1].sum_{deq}, super[blocks[b].group]-p, super[blocks[b].group]+p) ▷ superblock's group has at most p difference with the value stored in <code>super[]</code>. 419: if dir is right then 420: i+= blocks[superBlock].num_{deq-left} ▷ consider the dequeues from the right child 421: end if 422: return this.parent.INDEXDEQ(superBlock, i) 423: end if 424: end INDEXDEQ </pre>
xBaseCase	
puteSuper	
iderRight	
Algorithm Leaf	
	<pre> 601: <i>void</i> APPEND(<i>block</i> blk) ▷ Append is only called by the owner of the leaf. 602: blk.group= head 603: blocks[head]= blk 604: head+=1 605: parent.PROPAGATE() 606: end APPEND </pre>
pendStart	
appendEnd	
Algorithm BlockList	
	<p>▷ : Supports two operations <code>blocks.tryAppend(Block b)</code>, <code>blocks[i]</code>. Initially empty, when <code>blocks.tryAppend(b, n)</code> returns true b is appended to <code>blocks[n]</code> and <code>blocks[i]</code> returns ith block in the blocks. If some instance of <code>blocks.tryAppend(b, n)</code> returns false there is a concurrent instance of <code>blocks.tryAppend(b', n)</code> which has returned true.blocks[0] contains an empty block with all fields equal to 0 and <code>end_{left}</code>, <code>end_{right}</code> pointers to the first block of the corresponding children.</p> <p><i>block[]</i> blocks: array of blocks</p> <p><i>int[]</i> super: super[i] stores an approximate index of the superblock of the blocks in blocks whose group field have value i.</p> <pre> 701: <i>boolean</i> TRYAPPEND(<i>block</i> blk, <i>int</i> n) 702: return CAS(blocks[n], null, blk) 703: end TRYAPPEND </pre>

2 Proof of Linearizability

TEST Fix the logical order of definitions (cyclic references).

TEST Is it better to show $\text{ops}(\text{EST}_n, t)$ with EST_n, t ?

Question A good notation for *the index of the b*?

Question How to remove the notion of time? To say $\text{pre}(n, i)$ contains $n.\text{blocks}[0..i]$ instead of $\text{EST}(n, t)$ which $\text{head}=i$ at time t . Is it good? Furthermore, can we remove the notion of established blocks?

Definition 1 (Block). A block is an object storing some statistics, as described in Algorithm Queue. A block in a node's blocklist implicitly represents a set of operations. If $n.\text{blocks}[i] == b$ we call i the *index* of block b . Block b is before block b' in node n if and only if the index of the b is smaller than the index of the b' 's. For a block in a `BlockList` we define *the prefix for the block* to be the blocks in the `BlockList` up to and including the block.

Lemma 2 (head Increment). *Let R be an instance of Refresh on node n that reaches Line 313. After R terminates $n.\text{head}$ is greater than h , the value read in line 310 of R .*

Proof. If Line 318 or 321 are successful then the claim holds, otherwise another process has incremented the head from h to $h+1$. \square

Invariant 3 (headPosition). If the value of $n.\text{head}$ is h then, $n.\text{blocks}[i] = \text{null}$ for $i > h$ and $n.\text{blocks}[i] \neq \text{null}$ for $i < h$.

Proof. The invariant is true initially since 1 is assigned to $n.\text{head}$ and $n.\text{blocks}[x]$ is null for every x . The truth of the invariant may be affected by writing into $n.\text{blocks}$ or incrementing $n.\text{head}$. We show the invariant still holds after these two changes.

In the algorithm, some value is appended to $n.\text{blocks}[]$ by writing into $n.\text{blocks}[\text{head}]$ only in Line 313. Writing into $n.\text{blocks}[\text{head}]$ preserves the invariant, since the claim does not talk about $n.\text{blocks}[\text{head}]$. The value of $n.\text{head}$ is modified only in lines 318 and 321. Depending on whether the `TryAppend()` in Line 313 succeeded or not, we show that the claim holds after the increment of $n.\text{head}$ in either case. If $n.\text{head}$ is incremented to h it is sufficient to show $n.\text{blocks}[h] \neq \text{null}$ to prove the invariant still holds. In the first case the process applied a successful `TryAppend(new, h)` in line 314, which means $n.\text{blocks}[h]$ is not null anymore. Note that whether 318 or 318 return true or false, after they finish we know that $n.\text{head}$ has been incremented from the value read in Line 310 (Lemma 2). The failure case is also the same since it means some non-null value has been written into $n.\text{blocks}[\text{head}]$ by some process. \square

Explain More

Lemma 4 (headProgress). $n.\text{head}$ is non-decreasing over time. If $n.\text{blocks}[i] \neq \text{null}$ and $i > 0$ then $n.\text{blocks}[i].\text{end}_{\text{left}} \geq n.\text{blocks}[i-1].\text{end}_{\text{left}}$ and $n.\text{blocks}[i].\text{end}_{\text{right}} \geq n.\text{blocks}[i-1].\text{end}_{\text{right}}$.

Proof. The first claim follows trivially from the pseudocode since $n.\text{head}$ is only incremented in the pseudocode in lines 318 and 321 of `Refresh()`.

Consider the block b written into $n.\text{blocks}[i]$ by `TryAppend()` at Line 313. It is created by the `CreateBlock(i)` called at Line 311. Prior to this call to `CreateBlock(i)`, $n.\text{head}=i$ at Line 310, so $n.\text{blocks}[i-1]$ is already a non-null value b' by Invariant 3. Thus the `CreateBlock(i-1)` that creates b' terminates before `CreateBlock(i)` that creates b is invoked. The value written into $b.\text{end}_{\text{left}}$ at Line 333 of `CreateBlock(i)` was read from $n.\text{left.head}-1$ at Line 331 of `CreateBlock(i)`. Similarly, the value in $n.\text{blocks}[i-1].\text{end}_{\text{left}}$ was read from $n.\text{left.head}-1$ during the call to `CreateBlock(i-1)`. Since $n.\text{left.head}$ is non-decreasing $b'.\text{end}_{\text{left}} \leq b.\text{end}_{\text{left}}$. The proof for $\text{end}_{\text{right}}$ is similar. \square

Definition 5 (Subblock). Block b is a *direct subblock* of $n.\text{blocks}[i]$ if it is in $n.\text{left.blocks}[n.\text{blocks}[i-1].\text{end}_{\text{left}}+1..n.\text{blocks}[i].\text{end}_{\text{left}}] \cup n.\text{right.blocks}[n.\text{blocks}[i-1].\text{end}_{\text{right}}+1..n.\text{blocks}[i].\text{end}_{\text{right}}]$. Block b is a subblock of $n.\text{blocks}[i]$ if b is a direct subblock of $n.\text{blocks}[i]$ or a subblock of a direct subblock of $n.\text{blocks}[i]$.

Definition 6 (Superblock). Block b is *direct superbloc* of block c if c is a direct subblock of b . Block b is *superblock* of block c if c is a subblock of b .

Definition 7 (Operations of a block). A leaf block b in a leaf represents `enqueue(x)` if $b.element \neq \text{null}$. Else if $b.element = \text{null}$ b represents a `dequeue()`. The set of operations of block b are the operations in the subblocks of b . We denote the set of operations of block b by $ops(b)$.

We say block b is *propagated to node n* if b is in $n.blocks$ or is a subblock of a block in $n.blocks$. We also say b contains op if $op \in ops(b)$.

Definition 8. A block b in $n.blocks$ is *established* at time t if $n.head > \text{index of } b \text{ at time } t$. $EST_{n, t}$ is the set of established blocks of node n at time t .

Observation 9. Once a block b is written in $n.blocks[i]$ then $n.blocks[i]$ never changes.

Lemma 10. Every block has at most one direct superbloc.

Proof. To show this we are going to refer to the way $n.blocks[]$ is partitioned while propagating blocks up to $n.parent$. $n.CreateBlock(i)$ merges the blocks in $n.left.blocks[n.blocks[i-1].end_{left}..n.blocks[i].end_{left}]$ and $n.right.blocks[n.blocks[i-1].end_{right}..n.blocks[i].end_{right}]$. (Lines lastLine, pr). Since end_{left}, end_{right} are non-decreasing ($n.blocks[i].end_{left|right} > n.blocks[i-1].end_{left|right}$), so the range of the subblocks of $n.blocks[i]$ which is $(n.blocks[i-1].end_{dir}+1..n.blocks[i].end_{dir})$ does not overlap with the range of the subblocks of $n.blocks[i-1]$. \square

Corollary 11 (No Duplicates). If op is in $n.blocks[i]$ then there is no $j \neq i$ such that $op \in ops(n.blocks[j])$.

Lemma 12 (establishedOrder). If time $t < \text{time } t'$, then $ops(EST_{n, t}) \subseteq ops(EST_{n, t'})$.

Proof. Blocks are only appended (not modified) with CAS to $n.blocks[n.head]$ and $n.head$ is non-decreasing, so the set of operations in established blocks of a node can only grow. \square

useless?

ueRefresh

Lemma 13 (trueRefresh). *Let t_i be the time an instance R of $n.Refresh()$ is invoked and t_t be the time it terminates. If the $TryAppend(new, s)$ of R returns **true**, then $ops(EST_{n.left, t_i}) \cup ops(EST_{n.right, t_i}) \subseteq ops(EST_n, t_t)$.*

Proof. Since $TryAppend$ returns true a block **new** is written into $n.blocks[h]$ in Line ^{cas}313.

We show $ops(EST_{n.left, t_i}) \subseteq ops(EST_n, t_t)$. Let h be the value $n.Refresh()$ reads from $n.head$ at line ^{readHead}310, $h_{left,i}$ be the value of $n.left.head$ at t_i and $h_{left,read}$ be the value read from $n.left.head-1$ at line ^{lastLine}331. end_{left} field of the block returned by $CreateBlock(i)$ is $h_{left,read}$. By lines ^{prevLine}332 and ^{lastLine}331 the new block in $n.blocks[h]$ contains $n.left.blocks[n.blocks[h-1].end_{left}+1..h_{left,read}]$. Since $left.head$ is read after t_i then $h_{left,read} > h_{left,i}$ which means $ops(EST_{n.left, t_i}) \subseteq ops(n.left.blocks[0..h_{left,read}])$. After the successful $TryAppend$ in line ^{cas}313 we know all blocks in $n.left.blocks[0..h_{left,read}-1]$ are subblocks of $n.blocks[0..h]$ by the definition of subblock. At t_t we have $n.head > h$ by Lemma ^{lem:headProgress}4. So $n.blocks[1..h]$ are in EST_{n,t_t} by definition of EST . Note that after line ^{incrementHead2}321 we are sure that the head is incremented by Lemma ^{lem:headInc}2) which means $n.head = h+1$ at t_t so the new block is established at t_t and the new block contains the new operations which is what we wanted to show. The proof for $ops(EST_{n.right, t_i}) \subseteq ops(EST_n, t_t)$ is the same. □

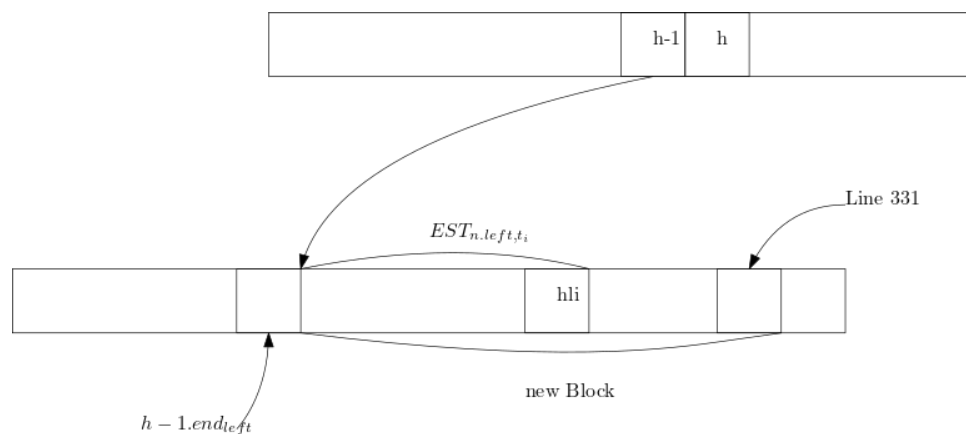


Figure 1: New established operations of the left child are in the new block.

ueRefresh

Lemma 14 (Stronger True Refresh). *Let t_i be the time an instance of $n.Refresh()$ read the head (Line ^{readHead}310) and t_t be the time its $TryAppend(new, s)$ terminates with and returns **true** (Line ^{cas}313). We have $ops(EST_{n.left, t_i}) \cup ops(EST_{n.right, t_i}) \subseteq ops(n.blocks)$.*

Lemma 15 (Double Refresh). *Consider two consecutive failed instances R_1, R_2 of $\mathbf{n.Refresh}()$ by some process. Let t_1 be the time R_1 is invoked and t_2 be the time R_2 terminated. We have $\text{ops}(\text{EST}_{\mathbf{n.left}}, t_1) \cup \text{ops}(\text{EST}_{\mathbf{n.right}}, t_1) \subseteq \text{ops}(\text{EST}_{\mathbf{n}}, t_2)$.*

Proof.

If Line 313 of R_1 or R_2 returns **true**, then the claim is held by Lemma 13. Let R_1 read i and R_2 read $i+1$ from Line 310. If R_2 reads some value greater than $i+1$ in Line 310 it means a successful instance of $\mathbf{Refresh}()$ started after Line 310 of R_1 and finished its Line 318 or 321 before 310 of R_2 , from Lemma 13 by the end of this instance $\text{ops}(\text{EST}_{\mathbf{n.left}}, t_1) \cup \text{ops}(\text{EST}_{\mathbf{n.right}}, t_1)$ has been propagated.

Since R_2 's $\mathbf{TryAppend}()$ returns **false** then there is another successful instance R'_2 of $\mathbf{n.Refresh}()$ that has done $\mathbf{TryAppend}()$ successfully into $\mathbf{n.blocks}[i+1]$ before R_2 tries to append. In Figure 1 we see why the block R'_2 is appending contains established block in the \mathbf{n} 's children at t_1 , since it create a block reading the head after t_1 . By Lemma 14 after R'_2 's CAS we have $\text{ops}(\text{EST}_{\mathbf{n.left}}, t_1) \cup \text{ops}(\text{EST}_{\mathbf{n.right}}, t_1) \subseteq \text{ops}(\mathbf{n.blocks})$. Also by Lemma 2 of R'_2 head is more than $i+1$ after R'_2 's 321 line, so the block appended by R'_2 to \mathbf{n} is established by then. To summarized t_1 is before R'_2 's read head and R'_2 's CAS is before R_2 's termination. So $\text{ops}(\text{EST}_{\mathbf{n.left}}, t_1) \cup \text{ops}(\text{EST}_{\mathbf{n.right}}, t_1) \subseteq \text{ops}(\text{EST}_{\mathbf{n}}, t_2)$. \square

last sentence need more detail and should be earlier. define i and tell why R_2 prime exists

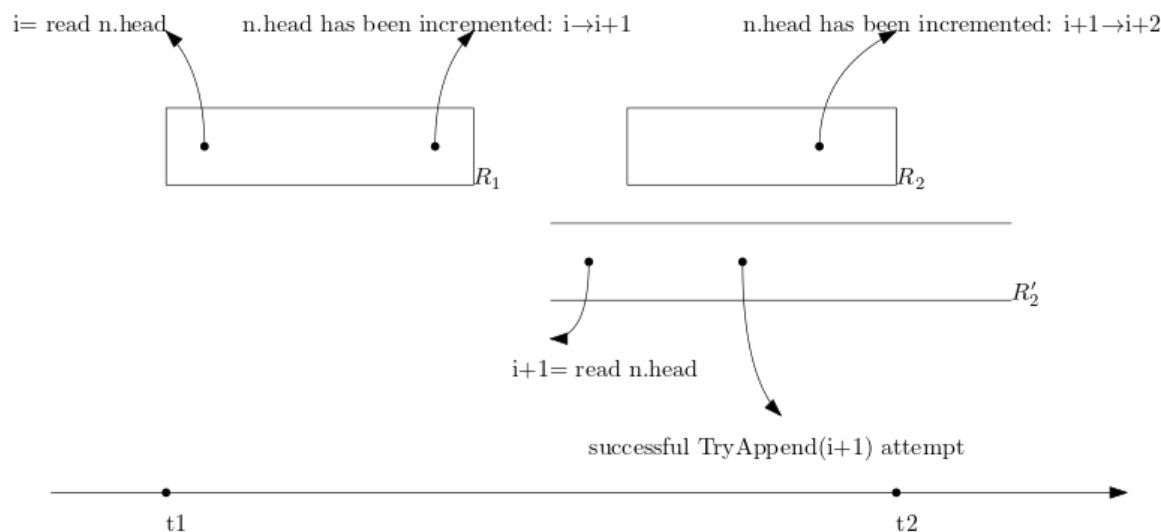


Figure 2: $t_1 < r_1$ reading head $<$ incrementing $\mathbf{n.head}$ from i to $i+1 < R'_2$ reading head $< \mathbf{TryAppend}(i+1) <$ incrementing $\mathbf{n.head}$ from $i+1$ to $i+2 < t_2$

this chain with more depth should be in the proof

lyRefresh

Corollary 16 (Propagate Step). *All operations in \mathbf{n} 's children's established blocks before running line 302 of a **Propagate** routine are guaranteed to be in \mathbf{n} 's established blocks after line 303.*

Proof. If 302 or 303 succeed, the claim is true by Lemma 13. Otherwise Lines 302 and 303 satisfy the preconditions of Lemma 15. \square

actlyOnce

Corollary 17. *After **Append**(blk) finishes $\text{ops}(\text{blk}) \subseteq \text{ops}(\text{root.blocks}[\mathbf{x}])$ for exactly one \mathbf{x} .*

Proof. After **Append**(blk)'s termination, blk is in **root.blocks** since blk is established in the leaf it has been added to. By applying Lemma 16 inductively it is propagated up to the root. Finally Lemma 11 shows only one block in the root contains blk. \square

blockSize **Lemma 18** (Block Size Upper Bound). *Each block contains at most one operation of each process.*

Proof. To derive a contradiction, assume there are 2 operations op_1 and op_2 from process p in block b in node n . WLOG op_1 is invoked earlier than op_2 . A process cannot invoke more than one operations concurrently, so one operation has to be finished before the other starts. So op_1 has terminated before op_2 started. By Corollary lem::appendExactlyOnce 17 before appending op_2 to the tree op_1 exists in every node from the path of p 's leaf to the root. Because op_1 's **Append** is finished before op_2 's **Append** starts. So there is some block b' before b in n containing op_1 . op_1 existing in b and b' contradicts Lemma append 11. □

blocksBound **Lemma 19** (Subblocks Upperbound). *Each block has at most p direct subblocks.*

Proof. It follows directly from Lemma blockSize 18 and the observation that each block appended to the tree contains at least one operation, induced from Line addOP 312. □

Definition 20 (Ordering of operations inside the nodes). ► Note that processes are numbered from 1 to p , left to right in the leaves of the tree and from Lemma 118 we know there is at most one operation from each process in a given block.

- We call operations strictly before op in the sequence of operations S , prefix of the op .
- $E(n, b)$ is the sequence of enqueue operations $\in \text{ops}(\mathbf{n.blocks}[b])$ ordered by their process id.
- $E_{n,b,i}$ is the i th enqueue in $E(n, b)$.
- $D(n, b)$ is the sequence of dequeue operations $\in \text{ops}(\mathbf{n.blocks}[b])$ ordered by their process id.
- $D_{n,b,i}$ is the i th dequeue in $D(n, b)$.
- Order of the enqueue operations in n : $E(n) = E(n, 1).E(n, 2).E(n, 3)\dots$
- $E_{n,i}$ is the i th enqueue in $E(n)$.
- Order of the dequeue operations in n : $D(n) = D(n, 1).D(n, 2).D(n, 3)\dots$
- $D_{n,i}$ is the i th dequeue in $D(n)$.
- Linearization: $L = E(\text{root}, 1).D(\text{root}, 1).E(\text{root}, 2).D(\text{root}, 2).E(\text{root}, 3).D(\text{root}, 3)\dots$

Note that in the non-root nodes we only order enqueues and dequeues among the operations of their own type. Since `GetENQ()` only searches among enqueues and `IndexDEQ()` works on dequeues.

Lemma 21 (Get correctness). *If $n.blocks[b].num_{enq} \geq i$ then $n.GetENQ(b, i)$ returns $E_{n,b,i}$.*

Proof. We are going to prove this lemma by induction on the height of node n . The base case, where n is a leaf, is straightforward (Line getBaseCase 403). Leaf blocks each contain exactly one operation, so only $n.GetENQ(b, 1)$ can be called where n is a leaf. $n.GetENQ(b, 1)$ returns the operation stored in the b th block of leaf n .

For the induction step we prove $n.GetENQ(b, i)$ returns $E_{n,b,i}$, if $n.child.GetENQ(b, i)$ returns $E_{n.child,b,i}$. For non-leaf nodes it is decided that the i th enqueue in block b of internal node n is in the $n.blocks[b]$'s subblocks in the left child of n or in the $n.blocks[b]$'s subblocks in the right child of n (line leftOrRight 404). From Definition ordering 20 we know enqueue operations in a block are ordered by their process id and since in leaves of the tree are ordered by process id from left to right, thus operations from the left subblocks come before operations from the right subblocks in a block (See Figure figGet 3). Furthermore $b.num_{enq-left}$ stores the number of `enqueue()` operations from the $n.blocks[b]$'s subblocks in the left child of n . So if i is greater than $b.num_{enq-left}$ it means i th operation is propagated from the right child, otherwise we should search for the i th enqueue in the left child. By definition def::opdef::subblock 7 and b we need to search in subblocks of $n.blocks[b]$ from the range $n.left.blocks[n.blocks[i-1].end_{left}+1..n.blocks[i].end_{left}] \cup n.right.blocks[n.blocks[i-1].end_{right}+1..n.blocks[i].end_{right}]$.

If the i th enqueue of $n.blocks[b]$ is in the left child it would be i th enqueue in $n.left.blocks[n.blocks[i-1].end_{left}+1..n.blocks[i].end_{left}]$ by Definition def::subblock 5. Also we know there are $eb = n.blocks[b-1].sum_{enq-left}$ enqueues in the blocks before this range, so $E_{n,b,i}$ is $E_{n.left,i+eb}$, which is $E_{n.left,b',i'}$ for some b' and i' . We can compute b' search for $i + eb$ th enqueue in $n.left$ and i' is $i+eb-n.left.blocks[b'-1].sum_{enq}$. The parameters in leftChildGet 405 are for searching $E_{n.left,i+eb}$ in $n.left.block$ in the expected range of blocks, so this BSearch returns the index of the subblock containing $E_{n,b,i}$.

Else if the enqueue we are looking for is in the right child then there are $n.blocks[b].num_{enq-left}$ enqueues ahead of it in $n.blocks[b]$ but not in $n.right.blocks[n.blocks[i-1].end_{right}+1..n.blocks[i].end_{right}]$. So we need to search for $i - n.blocks[b].num_{enq-left} + n.blocks[b-1].sum_{enq-right}$ (Line rightChildGet 409). Other parameters are assigned similar for the left child. So in both cases the direct subblock containing $E_{n,b,i}$ is computed in Lines leftChildGet 405 and rightChildGet 409.

Finally, $n.child.GetENQ()$ is invoked on the subblock containing $E_{n,b,i}$ which returns $E_{n,b,i}$ by the hypothesis of the induction. \square

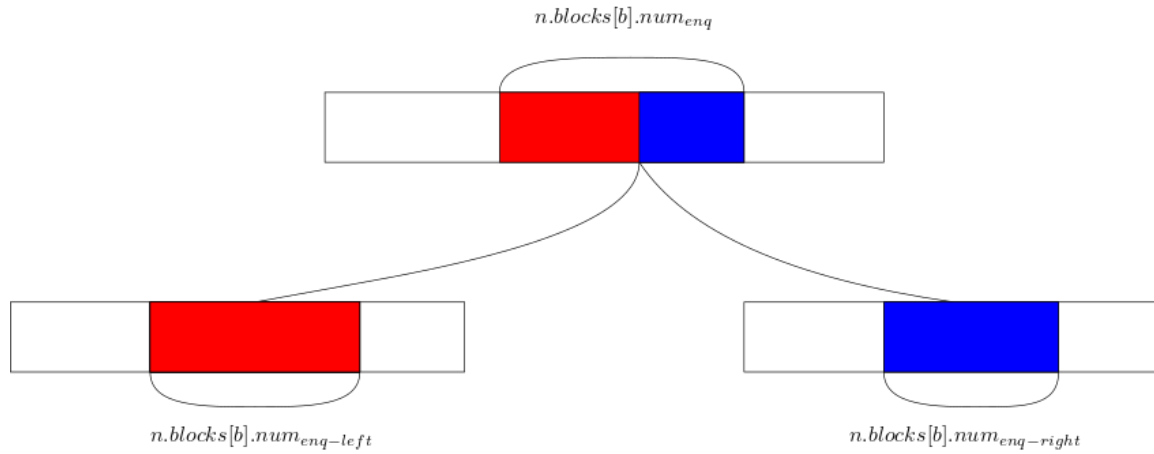


Figure 3: The number and ordering of the enqueues propagated from the left and the right child to $n.blocks[b]$. Enqueue operations from the left subblocks, colored red, are ordered before the enqueue operations from the right child, colored blue.

dsearch

Lemma 22 (DSearch correctness). Assume $\text{root.blocks}[\text{end}].\text{sum}_{\text{enq}} \geq e$ and $E_{\text{root},e}$'s element is the response to some `Dequeue()` operation in $\text{root.blocks}[\text{end}]$. $\text{DSearch}(e, \text{end})$ returns $\langle b, i \rangle$ such that $E_{\text{root},b,i} = E_{\text{root},e}$.

Proof. It is trivial to see that the doubling search from $\text{root.blocks}[\text{end}]$ to $\text{root.blocks}[0]$ will find $E_{\text{root},e}$ eventually. Because $\text{root.blocks}[].\text{sum}_{\text{enq}}$ is an increasing value from 0 to some value greater than e . So there is a b that $\text{root.blocks}[b].\text{sum}_{\text{enq}} > e$ but $\text{root.blocks}[b-1].\text{sum}_{\text{enq}} < e$.

First we show $\text{end}-b \leq 2 \times \text{root.blocks}[b].\text{size} + \text{root.blocks}[\text{end}].\text{size} + 1$. From line ^{addOP} 312, we know that size of the every block in the tree is greater than 0. So each block in $\text{root.blocks}[b..\text{end}]$ contains at least one `Enqueue` or at least one `Dequeue`. Suppose there were more than $\text{root.blocks}[b].\text{size}$ `Dequeues` in $\text{root.blocks}[b+1..\text{end}-1]$. Then the queue would become empty at some point after $\text{blocks}[b]$'s last operations and before $\text{root.blocks}[\text{end}]$'s first operation. Which means the response to a `Dequeue` in $\text{root.blocks}[\text{end}]$ could not be in $E(n, b)$. Furthermore since the size of the queue would become $\text{root.blocks}[\text{end}].\text{size}$ after the $\text{root.blocks}[\text{end}]$, there cannot be more than $\text{root.blocks}[b].\text{size} + \text{root.blocks}[\text{end}].\text{size}$ `Enqueues`. Because there can be at most $\text{root.blocks}[b].\text{size}$ `Dequeues` and the final size is $\text{root.blocks}[\text{end}].\text{size}$. Overall there can be at most $2 \times \text{root.blocks}[b].\text{size} + \text{root.blocks}[\text{end}].\text{size}$ operations in $\text{root.blocks}[b+1..\text{end}-1]$ and since each block size is ≥ 1 thus there are at most $2 \times \text{root.blocks}[b].\text{size} + \text{root.blocks}[\text{end}].\text{size}$ blocks in between $\text{root.blocks}[b]$ and $\text{root.blocks}[\text{end}]$. So $\text{end}-b \leq 2 \times \text{root.blocks}[b].\text{size} + \text{root.blocks}[\text{end}].\text{size} + 1$. See Figure ^{end-b} 5.

Now that we know there are at most $\text{root.blocks}[b].\text{size} + \text{root.blocks}[\text{end}].\text{size}$ blocks in between $\text{root.blocks}[b]$ and $\text{root.blocks}[\text{end}]$ then with doubling search in $\Theta(\log(\text{root.blocks}[b].\text{size} + \text{root.blocks}[\text{end}].\text{size}))$ steps we reach $\text{start}=c$ that the $\text{root.blocks}[c].\text{sum}_{\text{enq}}$ is less than e and $\text{end}-c$ is not more than $2 \times \text{root.blocks}[b].\text{size} + \text{root.blocks}[\text{end}].\text{size}$. Beause otherwise, then $(\text{end}-c)/2$ satisfied the $\text{root.blocks}[(\text{end}-c)/2].\text{sum}_{\text{enq}} < e$. In line ^{doubling} 804 the differenece between end and start is doubled. See Figure ^{fig::doubling} 4.

After computing b , the value i is computed via the definition of sum_{enq} in constant time (Line ^{DSearchCompute i} 807). So the routine non constant part is the binary search which takes $\Theta(\log \text{root.blocks}[b].\text{size} + \text{root.blocks}[\text{end}].\text{size})$ steps from the first paragraph.

□

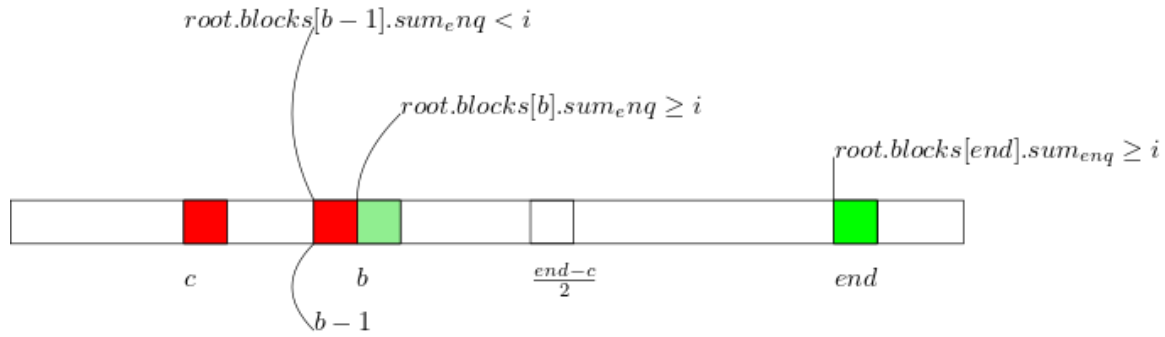


Figure 4: Distance relations between b, c, end

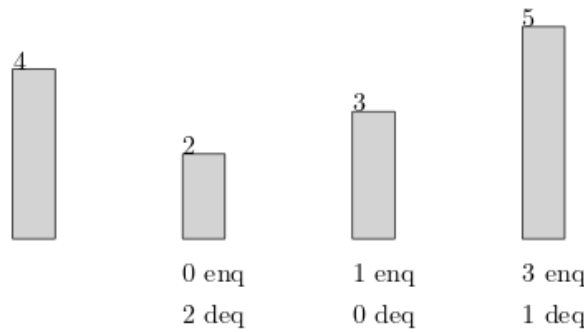


Figure 5: The number written on top of the bars is the queue size. the first block is b and the last block is end .

end-b

Lemma 23 (Index correctness). $n.\text{IndexDEQ}(b, i)$ returns the rank in $D(\text{root})$ of $D_{n,b,i}$.

Proof. We will prove this by induction on the distance of n from the root. We can see the base case $\text{root}.\text{IndexDEQ}(b, i)$ is trivial (Line indexBaseCase 415). In the non-root nodes $n.\text{IndexDEQ}(b, i)$ computes the superblock of the i th Dequeue in the b th block of n in $n.\text{parent}$ by Lemma superBlockComputeSuper 24 (Line 418). After that the order in $D(n.\text{parent}, \text{superblock})$ is computed and $\text{index}()$ is called on $n.\text{parent}$ recursively. Then if the operation was propagated from the right child the number of dequeues from the left child are added to it (Line considerRight 420), because the left child operations come before the right child operations (Definition ordering 20). \square

Do I need to talk about the computation of the order in the parent which is based on the definition of ordering of dequeues in a block?
Make sure to show preconditions of all invocation of BSearch are satisfied.

Lemma 24 (Computing SuperBlock). After computing line computeSuper 418 of $n.\text{IndexDEQ}(b, i)$, $n.\text{parent}.\text{blocks}[\text{superblock}]$ contains $D(n, b, i)$.

Proof. Lemmas 28,29,30,31. \square

Lemma 25. Value read for $\text{super}[b.\text{group}]$ in line 418 is not null.

Proof. Values np_{dir} read in lines setNP 337, super are set before incrementing in lines setSuperNP 315,316. So before incrementing $\text{num}_{\text{propagated}}$, $\text{super}[\text{num}_{\text{propagated}}]$ is set so it cannot be null while reading. \square

Lemma 26. $\text{super}[]$ preserves order from child to parent; i.e. if in node n block b is before c then $b.\text{group} \leq c.\text{group}$

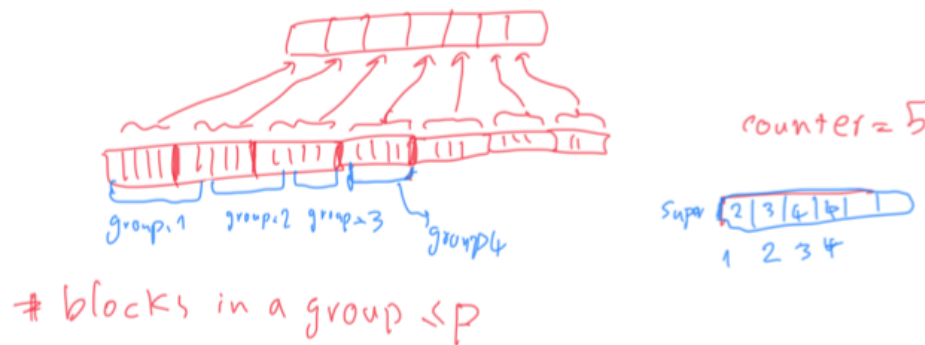
Proof. Line setGroup 329. Since $\text{num}_{\text{propagated}}$ is increasing. \square

Lemma 27. Let b, c be in node n , if $b.\text{group} \leq c.\text{group}$ then $\text{super}[b.\text{group}] \leq \text{super}[c.\text{group}]$

Proof. Line setSuper 315. \square

Lemma 28. The number of the blocks with $\text{group}=i$ in a node is $\leq p$.

Proof. For the sake of simplicity we assumed all the blocks are propagated from the left child. \square

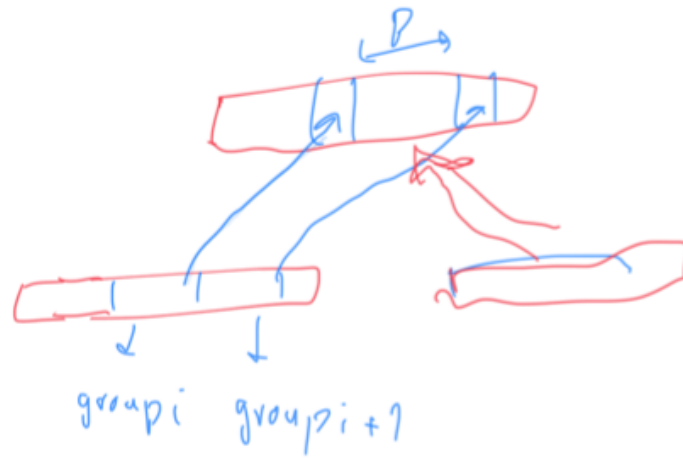


Lemma 29. $\text{super}[i+1] - \text{super}[i] \leq p$

Proof. In a Refresh with successful CAS in line 46, super and counter are set for each child in lines 48,49. Assume the current value of the counter in node n is $i+1$ and still $\text{super}[i+1]$ is not set. If an instance of successful $\text{Refresh}(n)$ finishes $\text{super}[i+1]$ is set a new value and a block is added after $n.\text{parent}[\text{super}[i]]$. There could be at most p successful unfinished concurrent instances of $\text{Refresh}()$ that have not reached line 49. So the distance between $\text{super}[i+1]$ and $\text{super}[i]$ is less than p . \square

Lemma 30 (super property). If $\text{super}[i] \neq \text{null}$ in node n , then $\text{super}[i]$ is the index of the superblock of a block with $\text{time}=i$ in $n.\text{parent}.\text{blocks}$.

Lemma 31. Superblock of b is within range $\pm 2p$ of the $\text{super}[b.\text{time}]$.



Proof. $\text{super}[i]$ is the index of the superblock of a block containing block b , followed by Lemma ^{superCounter} 30. $\text{super}(b)$ is the real superblock of b . $\text{super}(t)$ is the index of the superblock of the last block with time t . If $b.\text{time}$ is t we have:

$$\text{super}[t] - p \leq \text{super}[t - 1] \leq \text{super}(t - 1) \leq \text{super}(b) \leq \text{super}(t + 1) \leq \text{super}(t + 1) \leq \text{super}[t] + p$$

□

Lemma 32. Search in each level of `IndexDeq()` takes $O(\log p)$ steps.

Proof. Show preconditions are satisfied and the range is p .

□

Definition 33. Assume the operations in L are applied on an empty queue. If element of `enqueue e` is the response to `dequeue d` then we say $R(d)=e$. If d 's response is `null` (queue is empty) then $R(d)=\text{null}$.

Definition 34. In an execution on a queue, the dequeue operations that return some value are called *non-null dequeues*.

Observation 35. k th non-null dequeue in an execution returns the element of k th enqueue.

Lemma 36. `root.blocks[b].size` is the size of the queue if the operations in the prefix for the b th block in the root are applied with the order of L .

Proof. need to say? :: If the size of a queue is greater than 0 then a `Dequeue()` would decrease the size of the queue, otherwise the size of the queue remains 0. By definition ordering 20 enqueue operations come before dequeue operations in a block in L .

We prove the claim by induction on b . Base case $b=0$ is trivial since the queue is initially empty and `root.blocks[0].size=0`. For $b=i$ we are going to use the hypothesis for $b=i-1$. If there are more than `root.blocks[i-1].size+ root.blocks[i].sum_enq` dequeue operations in `root.blocks[i]` then the queue would become empty after `root.blocks[i]`. Otherwise we can compute the size of the queue after b th block using with this equality `root.blocks[b].size= root.blocks[b-1].size+ root.blocks[b].sum_enq- root.blocks[b].sum_deq` (Line computeLength 342). See Table qhistory 1 for an example of running some blocks of operations on an empty queue. \square

Lemma 37 (Duality of #non-null dequeues and `block.size`). If the operations are applied with the order of L , the number of non-null dequeues in the prefix for a block b is `b.sum_enq-b.size`

Proof. There are `b.sum_enq` enqueue operations in the prefix for b , then the size of the queue after the prefix for b is `#enqs - #non-null dequeues` in the prefix for b , by Observation 35. So `#non-null dequeues` is `b.sum_enq-b.size`. The correctness of the `block.size` field is shown in Lemma sizeCorrectness 36. \square

Lemma 38. $R(D_{\text{root},b,i})$ is null iff `root.blocks[b-1].size + root.blocks[b].num_enq- i < 0`.

Lemma 39 (Computing Response). `FindResponse(b,i)` returns $R(D_{\text{root},b,i}).\text{element}$.

Proof. First note that by Definition ordering 20 the linearization ordering of operations will not change as new operations come so instead of talking about the linearization of operations before the $E_{\text{root},b,i}$ we talk about what if the whole operation in the linearization are applied on a queue.

$D_{\text{root},b,i}$ is $D_{\text{root},\text{root.blocks}[b-1].\text{sum_deq}+i}$ from the definition ordering 20 and `sum_enq`. $D_{\text{root},b,i}$ returns null if `root.blocks[b-1].size + root.blocks[b].num_enq- i < 0` by Lemma nullReturnCheckEmpty 38 (Line 220). Otherwise if it is d' th non-null dequeue in L it returns d' th enqueue by Observation responseToADeq 35. By Lemma numberOfNND 37 there are `root.blocks[b-1].sum_enq - root.blocks[b-1].size` non-null dequeue operations before prefix for `root.blocks[b-1]`. Note that the dequeues in `root.blocks[b]` before the i th dequeue are non-null dequeues. So the response is $E_{\text{root},i-\text{root.blocks}[b-1].\text{size}+\text{root.blocks}[b-1].\text{sum_deq}}$ (Line computeE 222). See figure computeResponseDetail 6.

After computing e we can find b,i such that $E_{\text{root},b,i} = E_{\text{root},e}$ using `DSearch` and then find its element using `GetEnq` (Line findAnswer 223). \square

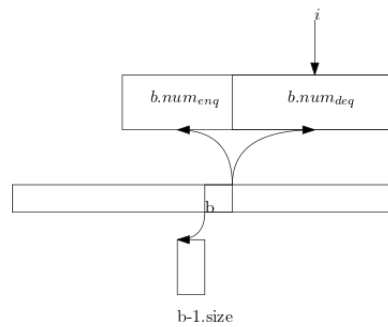


Figure 6: The position of $E_{\text{root},b,i}$.

	DEQ()	ENQ(5), ENQ(2), ENQ(1), DEQ()	ENQ(3), DEQ()	ENQ(4), DEQ(), DEQ(), DEQ(), DEQ()
#enqueues	0	3	1	1
#dequeues	1	1	1	4
#non-null dequeues	0	1	2	5
size	0	2	2	0

Table 1: An example of root blocks fields. Blocks are from left to right and operations in the blocks are also from the left to right.

qhistory

Theorem 40 (Main). *The queue implementation is linearizable.*

Proof. We choose L in Definition [20](#) to be linearization ordering of operations and prove if we linearize operations as L the queue works consistently. \square

Lemma 41 (satisfiability). *L can be a linearization ordering.*

Proof. To show this we need to say if in an execution, op_1 terminates before op_2 starts then op_1 is linearized before op_2 . If op_1 terminates before op_2 starts it means $op_1.\text{Append}()$ is terminated before $op_2.\text{Append}()$ starts. From Lemma [11](#) op_1 is in `root.blocks` before op_2 propagates so op_1 is linearized before op_2 by Definition [20](#). \square

Once some operations are aggregated in one block they will be propagated together up to the root and we can linearize them in any order among themselves. Furthermore in L we arbitrary choose the order to be by process id, since it makes computations in the blocks faster. \square

Lemma 42 (correctness). *If operations are applied as L on a sequential queue, the sequence of the responses would be the same as our algorithm.*

Proof. Old parts to review We show that the ordering L stored in the root, satisfies the properties of a linearizable ordering.

1. If op_1 ends before op_2 begins in E , then op_1 comes before op_2 in T .
 - This is followed by Lemma [11](#). The time op_1 ends it is in root, before op_2 , by Definition [20](#) op_1 is before op_2 .
2. Responses to operations in E are same as they would be if done sequentially in order of L .
 - Enqueue operations do not have any response so it does no matter how they are ordered. It remains to prove Dequeue d returns the correct response according to the linearization order. By Lemma [39](#) it is deduced that the head of the queue at time of the linearization of d is computed properly. If the Queue is not empty by Lemma [21](#) we know that the returning response is the computed index element.

\square

Lemma 43 (Amortized time analysis). **Enqueue()** and **Dequeue** take $O(\log^2 p + q)$ steps (amortized analysis), which p is the number of processes and q is the size of the queue at the time of invocation.

Proof. **Enqueue(x)** consists of creating a **block(x)** and appending it to the tree. The first part takes constant time. To propagate **x** to the root the algorithm tries two **Refreshes** in each node of the path from the leaf to the root (Lines ~~302, 303~~ ^{firstRefresh, Refresh}). Each **Refresh** takes $O(1)$ steps since creating a block is done in constant time and does $O(1)$ CASes. Since the height of the tree is $O(\log p)$, **Enqueue(x)** takes $O(\log p)$ steps.

A **Dequeue()** creates a block with null value element, appends it to the tree, computes its order among operations, and returns the response. The first two part is similar to an **Enqueue** operation. To compute the order there are some constant steps and **IndexDeq** is called. **IndexDeq** does a search with range p in each level (Lemma ~~31~~ ^{superRange}) which takes $O(\log^2 p)$ in the tree. In the **FindResponse()** routine **DSearch()** in the root takes $\Theta(\log(\text{root.blocks}[b].\text{size} + \text{root.blocks}[\text{end}].\text{size}))$ by Lemma ~~22~~ ^{ldsearch}, which is $O(\log \text{size of the queue when enqueue is invoked} + \log \text{size of the queue when dequeue is invoked})$. Each search in **GetEnq()** takes $O(\log p)$ since there are $\leq p$ subblocks in a block (Lemma ~~19~~ ^{subBlocksBound}), so **GetEnq()** takes $O(\log^2 p)$ steps.

If we split **DSearch** time cost between the corresponding **Enqueue**, **Dequeue**, in amortized we have **Enqueue** takes $O(\log p + q)$ and **Dequeue** takes $O(\log^2 p + q)$ steps. □