

◇ *Shared*

- *Node\** `tree[]` : A binary tree of Nodes such that `Tree[0]` is the root and the left child and the right child and the parent of `Tree[i]` are `Tree[2i+1]`, `Tree[2i+2]` and `Tree[i/2]`.

◇ *Local*

- *int* `leaf`: `tree[leaf]` is the process's leaf in the tree.

◇ *Structures*► *Node*

- *BlockList* `blocks`: Supports two operations `blocks.tryAppend(Block b)`, `blocks[i]`. Initially empty, when `blocks.tryAppend(b)` returns true `b` is appended to the end of the list and `blocks[i]` returns *i*th block in the blocks. If some instance of `blocks.tryAppend(b)` returns false there is a concurrent instance of `blocks.tryAppend(b')` which has returned true. `blocks[0]` contains an empty block with all fields equal to 0 and end pointers to the first block of the corresponding children.
- *int* `numpropagated` = 0 : # groups of blocks have been propagated from the node. It may be behind its real value.
- *int[]* `super`: `super[i]` stores the index of the superblock of some block in blocks that its group field is *i*.

► *Root* extends *Node*

- *PBRT* `blocks`  
Implemented with a persistent red-black tree.

► *NonRootNode* extends *Node*

- *Block[]* `blocks`  
Implemented with an array and using CAS for appending to the head.
- *int* `head` = 1: #blocks in the blocks.

► *Leaf* extends *NonRootNode*

- *int* `lastdone`  
Each process stores the index of the most recent block in the root that the process has finished the last operation of the block. `enqueue(e)` is finished if `e` is returned by some `dequeue()` and `dequeue()` is finished when it computes its response.

► *Block* *b*      ▷ If *b* is `blocks[i]` (*i*!=0) then `b[-1]` is `blocks[i-1]`.

- *int* `group` : the value read from the node `n.numpropagated` when appending this block to `n`.

► *LeafBlock* extends *Block*

- *Object* `element` : Each block in a leaf represents an operation. The element shows the operation's argument if it is an enqueue, and if it is a dequeue this value is null.
- *int* `sumenq`, `sumdeq` : number of enqueue, dequeue operations in the leaf's containing this block till this block
- *Object* `response`  
response stores the response of the operation in the *LeafBlock*.

► *InternalBlock* extends *Block*

- *int* `sumenq-left` : #enqueues from the subblocks in the left child + `b[-1].sumenq-left`
- *int* `sumdeq-left` : #dequeues from the subblocks in the left child + `b[-1].sumdeq-left`
- *int* `sumenq-right` : #enqueues from the subblocks in the right child + `b[-1].sumenq-right`
- *int* `sumdeq-right` : #dequeues from the subblocks in the right child + `b[-1].sumdeq-right`
- *int* `endleft`, `endright` : index of the last subblock of the block in the left and right child

► *RootBlock* extends *InternalBlock*

- *int* `size` : size of the queue after this block's operations finish
- *counter* `numfinished` : number of finished operations in the block
- *int* `order` : the index of the block in the node containing the block. Useful in the root since in the PBRT we do not keep indices as key.

---

## Algorithm *Queue*

---

```

201: void ENQUEUE(Object e) ▷ Creates a block with element e and appends
    it to the tree.
202:   block newBlock= NEW(LeafBlock)
203:   newBlock.element= e
204:   newBlock.sumenq= tree[leaf].blocks[tree[leaf].head].sumenq+1
205:   newBlock.sumdeq= tree[leaf].blocks[tree[leaf].head].sumdeq
206:   tree[leaf].APPEND(newBlock)
207: end ENQUEUE

208: Object DEQUEUE()
209:   block newBlock= NEW(LeafBlock) ▷ Creates a block with null value
    element, appends it to the tree, computes its order among operations, then
    computes its response; if it exists returns the response's element.
210:   newBlock.element= null
211:   newBlock.sumenq= tree[leaf].blocks[tree[leaf].head].sumenq
212:   newBlock.sumdeq= tree[leaf].blocks[tree[leaf].head].sumdeq+1
213:   tree[leaf].APPEND(newBlock)
214:   <rdeq, brdeq>= tree[leaf].INDEX(tree[leaf].head, 1) ▷ r is the
    rank among all dequeues of the dequeue and brdeq is the index of the block
    in the root containing the dequeue.
215:   <renq, brenq>= FINDRESPONSE(rdeq, brdeq) ▷ rdeq is the rank of the
    enqueue which is the response to the dequeue or -1 if the response is null.
216:   if renq== -1 then
217:     output= null
218:     tree[0].blocks[brdeq].numfinished.inc() ▷ shared counter
219:     if tree[0].blocks[brdeq].numfinished==tree[0].blocks[brdeq].num
        then ▷ all the operations in the block containing the dequeue are finished.
220:       tree[leaf].lastdone= brdeq
221:     end if
222:   else
223:     output= GET(indexresponse)
224:     root.blocks[brenq].numfinished.inc()
225:     root.blocks[brenq].numfinished.inc()
226:     if root.blocks[br].numfinished==root.blocks[br].num then
227:       tree[leaf].lastdone= brenq
228:     else if root.blocks[brdeq].numfinished==root.blocks[brdeq].num
        then ▷ root.blocks[br] comes after root.blocks[bi].
229:       tree[leaf].lastdone= brdeq
230:     end if
231:   end if
232:   return output
233: end DEQUEUE

234: int, int FINDRESPONSE(int i, int b) ▷ Computes the rank and
    index of the block in the root of the enqueue that is the response of the ith
    dequeue in the root's bth block. Returns j-1,0 if the queue is empty.
235:   if root.blocks[b-1].size + root.blocks[b].numenq - i < 0 then
236:     return -1, 0
237:   else
238:     res= root.blocks[b-1].sumdeq- root.blocks[b-1].size + i
239:     return <res, tree[0].blocks.get(enq, res).order>
240:   end if
241: end FINDRESPONSE

```

---

---

**Algorithm** *Node*


---

<pre> 301: void PROPAGATE() 302:   if not this.REFRESH() then 303:     this.REFRESH()           ▷ Lemma Double Refresh 304:   end if 305:   if this is not root then   ▷ To check a node is the root we can check                                 its index is 0. 306:     this.parent().PROPAGATE() 307:   end if 308: end PROPAGATE  309: boolean REFRESH() 310:   h= head 311:   &lt;new, n<sub>left</sub>, n<sub>right</sub>&gt;= this.CREATEBLOCK(h)   ▷ n<sub>left</sub>, n<sub>right</sub> are the   values read from the children's num<sub>propagated</sub>s. 312:   if new.num==0 then return true   ▷ The block contains nothing. 313:   else if root.blocks.tryAppend(new) then 314:     for each dir in {left, right} do 315:       CAS(dir.super[n<sub>dir</sub>], null, h+1) 316:       CAS(dir.num<sub>propagated</sub>, n<sub>dir</sub>, n<sub>dir</sub>+1) 317:     end for 318:     CAS(head, h, h+1) 319:     return true 320:   else 321:     CAS(head, h, h+1)           ▷ Even if another process wins, help to                                 increase the head. It might fell sleep before increasing. 322:     return false 323:   end if 324: end REFRESH </pre>	<pre> 327: &lt;Block, int, int&gt; CREATEBLOCK(int i)                                 ▷ Creates a block to be inserted into this.blocks[i]. Returns                                 the created block as well as values read from each childnum<sub>propagated</sub> field.                                 The values are used for incrementing children's num<sub>propagated</sub>s if the block                                 was appended to this.blocks successfully. 328:   block newBlock= NEW(block) 329:   newBlock.group= num<sub>propagated</sub> 330:   newBlock.order= i 331:   for each dir in {left, right} do 332:     index<sub>last</sub>= n.dir.head 333:     index<sub>prev</sub>= n.blocks[i-1].end<sub>dir</sub> 334:     block<sub>last</sub>= n.dir.blocks[index<sub>last</sub>] 335:     block<sub>prev</sub>= n.dir.blocks[index<sub>prev</sub>] 336:     ▷ newBlock includes n.dir.blocks[index<sub>prev</sub>+1..index<sub>last</sub>]. 337:     n<sub>dir</sub>= n.dir.num<sub>propagated</sub> 338:     newBlock.end<sub>dir</sub>= index<sub>last</sub> 339:     newBlock.num<sub>enq-dir</sub>= block<sub>last</sub>.sum<sub>enq</sub> - block<sub>prev</sub>.sum<sub>enq</sub> 340:     newBlock.num<sub>deq-dir</sub>= block<sub>last</sub>.sum<sub>deq</sub> - block<sub>prev</sub>.sum<sub>deq</sub> 341:     newBlock.sum<sub>enq-dir</sub>= n.blocks[i-1].sum<sub>enq-dir</sub> + b.num<sub>enq-dir</sub> 342:     newBlock.sum<sub>deq-dir</sub>= n.blocks[i-1].sum<sub>deq-dir</sub> + b.num<sub>deq-dir</sub> 343:   end for 344:   if n is root then 345:     newBlock.size= max(root.blocks[i-1].size + b.num<sub>enq</sub> -                                 b.num<sub>deq</sub>, 0) 346:   end if 347:   return b, c<sub>left</sub>, c<sub>right</sub> 348: end CREATEBLOCK </pre>
--	---

↪ Precondition: blocks[start..end] contains a block with field  $f \geq i$

```

325: int BSEARCH(field f, int i, int start, int end)
                                ▷ Does binary search for the value
                                i of the given prefix sum feild. Returns the index of the leftmost block in
                                blocks[start..end] whose field f is  $\geq i$ .
326: end BSEARCH

```

---

---

**Algorithm Node**

---

↪ Precondition: `n.blocks[b]` contains *i*th enqueue in the node.

```
401: element GET(int b, int i)
402:   if this is leaf then return this.blocks[b].element
403:   else
404:     if  $i \leq$  this.blocks[b].numenq-left then                                ▷ i exists in the left child of this node
405:       subBlock= this.leftBSEARCH(sumenq, i, this.blocks[b-1].endleft+1, this.blocks[b].endleft)
406:       return this.left.GET(i-this.left.blocks[subBlock-1].sumenq, subBlock)
407:     else
408:       i= i-this.blocks[b].numenq-left
409:       subBlock= this.rightBSEARCH(sumenq, i, this.blocks[b-1].endright+1, this.blocks[b].endright)
410:       return this.right.GET(i-this.right.blocks[subBlock-1].sumenq, subBlock)
411:     end if
412:   end if
413: end GET
```

↪ Precondition: *b*th block of the node has propagated up to the root and *i*th dequeue in the node is in `this.blocks[b]`.

```
414: <int, int> INDEX(int b, int i)                                ▷ Returns the order in the root of ith dequeue in the bth block of node n among dequeues.
415:   if this is root then return <i, b>
416:   else
417:     dir= (this.parent().left==n)? left: right                                ▷ check n is a left or a right child
418:     superBlock= BSEARCH(this.parent(), this.sumdeq-dir, i, super[this.blocks[b].group]-p, super[this.blocks[b].group]+p) ▷ superblock's
group has at most p difference with the value stored in super[].
419:     if dir is left then
420:       i+= this.parent().blocks[superBlock-1].sumdeq-right
421:     else
422:       i+= this.parent.blocks[superBlock-1].sumdeq + this.blocks[superBlock].sumdeq-left                                ▷ consider dequeues from the right child
423:     end if
424:     return INDEX(this.parent(), superBlock, i)
425:   end if
426: end INDEX
```

---

---

**Algorithm Leaf**

---

```
501: void APPEND(block b)
502:   head+=1                                ▷ Lines 503 to 502 are done by one process at time.
503:   b.group= head                                ▷ Append is only called by the owner of the leaf.
504:   blocks[head]= b
505:   this.parent().PROPAGATE()
506: end APPEND
```

---

---

**Algorithm Root**

---

```
601: element GET(int i)                                ▷ Returns ith Enqueue.
602:   res= root.blocks.get(enq, i).order
603:   return GET(root, res, i-root.blocks[res-1].sumenq)
604: end GET
```

---

appendEnd  
pendStart

---

```

► PRBTree[rootBlock]
A persistant red-black tree supporting append(b, key),get(key=i),split(j).
append(b, key) returns true in case successful. Since order, sum_enqueue
both strictly increasing we can use one of them for another.
701: void RBTAPPEND(block b)          ▷ adds block b to the root.blocks
702:   step= root.head
703:   if step%p2==0 then ▷ Help every often p2 operations appended to the
      root. Used in lemma's using the size of the PBRT.
704:     Help()
705:     CollectGarbage()
706:   end if
707:   b.num_finished= 0
708:   return root.blocks.append(b, b.order)
709: end RBTAPPEND

710: void HELP          ▷ Helps pending operations
711:   for leaf l in leaves do ▷ if the tree is implemented with an array we
      can iterate over the second half of the array.
712:     last= l.head-1    ▷ l.blocks[last] can not be null because of
      appendSplitAnd
      lines 503-502.
713:     if l.blocks[last].element==null then ▷ operation is dequeue
714:       goto DeqRest
      215 with these values <> ▷ run Dequeue() for
      l.ops[last] after Propagate(). TODO
715:       l.responses[last]= response

716:   end if
717: end for
718: end HELP

719: void COLLECTGARBAGE          ▷ Collects the root blocks that are done.
720:   s=FindMostRecentDone(Root.Blocks.root) ▷ Lemma: If block b is
      done after helping then all blocks before b are done as well.
721:   t1,t2= RBT.split(order, s)
722:   RBTRoot.CAS(t2.root)
723: end COLLECTGARBAGE

724: Block FINDMOSTRECENTDONE(b)
725:   for leaf l in leaves do
726:     max= Max(l.maxOld, max)
727:   end for
728:   return max          ▷ This snapshot suffies.
729: end FINDYOUNGESTOLD

730: response FALLBACK(op i)          ▷ really necessary?
731:   if root.blocks.get(num_enqueue), i is null then ▷ this enqueue was already
      finished
732:     return this.leaf.response(block.order)
733:   end if
734: end FALLBACK

```

---