

---

**Algorithm** Fields description

---

◇ *Shared*

- *Tree tree* : A binary tree of Nodes. *root* is a pointer to the root node.

◇ *Local*

- *\*Node leaf* : a pointer to the process's leaf in the tree.

◇ *Structures*► *Node*

- *\*Node left, right, parent* : initialized when creating the tree.
- *BlockList blocks*
- *int num<sub>propagated</sub> = 0* : # groups of blocks that have been propagated from the node to its parent. Since it is incremented after propagating, it may be behind by 1.
- *int[] super*: *super[i]* stores the index of the superblock of some block in blocks whose group field is *i*.

► *Root* extends *Node*

- *PBRT blocks*  
BlockList is implemented with a persistent red-black tree.

► *NonRootNode* extends *Node*

- *Block[] blocks*  
BlockList is implemented with an array.
- *int size = 1*: #blocks in blocks.

► *Leaf* extends *NonRootNode*

- *int last<sub>done</sub>*  
Stores the index of the block in the root such that the process that owns this leaf has most recently finished the. A block is finished if all of its operations are finished. *enqueue(e)* is finished if *e* is returned by some *dequeue()* and *dequeue()* is finished when it computes its response. *put the definitions before the pseudocode*

► *Block*      ▷ For a block in a blocklist we define *the prefix for the block* to be the blocks in the BlockList up to and including the block. *put the definitions before the pseudocode*

- *int group* : the value read from *num<sub>propagated</sub>* when appending this block to the node.

► *LeafBlock* extends *Block*

- *Object element* : Each block in a leaf represents a single operation. For enqueue operations element is the input of the enqueue and for dequeue operations it is null.
- *Object response* : stores the response of the operation in the LeafBlock.
- *int sum<sub>enq</sub>, sum<sub>deq</sub>* : # enqueue, dequeue operations in the prefix for the block

► *InternalBlock* extends *Block*

- *int end<sub>left</sub>, end<sub>right</sub>* : index of the last subblock of the block in the left and right child
- *int sum<sub>enq-left</sub>* : # enqueue operations in the prefix for *left.blocks[end<sub>left</sub>]*
- *int sum<sub>deq-left</sub>* : # dequeue operations in the prefix for *left.blocks[end<sub>left</sub>]*
- *int sum<sub>enq-right</sub>* : # enqueue operations in the prefix for *right.blocks[end<sub>right</sub>]*
- *int sum<sub>deq-right</sub>* : # dequeue operations in the prefix for *right.blocks[end<sub>right</sub>]*

► *RootBlock* extends *InternalBlock*

- *int length* : length of the queue after performing all operations in the prefix for this block
- *counter num<sub>finished</sub>* : number of finished operations in the block
- *int order* : the index of the block in the BlockList containing the block.

---

*Variable naming:*

- *b<sub>op</sub>*: index of the block containing operation *op*
- *r<sub>op</sub>*: rank of operation *op* i.e. the ordering among the operations of its type according to linearization ordering

*Abbreviations:*

- *blocks[b].sum<sub>x</sub> = blocks[b].sum<sub>x-left</sub> + blocks[b].sum<sub>x-right</sub>* (for  $b \geq 0$  and  $x \in \{\text{enq}, \text{deq}\}$ )
- *blocks[b].sum = blocks[b].sum<sub>enq</sub> + blocks[b].sum<sub>deq</sub>* (for  $b \geq 0$ )
- *blocks[b].num<sub>x</sub> = blocks[b].sum<sub>x</sub> - blocks[b-1].sum<sub>x</sub>*  
(for  $b > 0$  and  $x \in \{\emptyset, \text{enq}, \text{deq}, \text{enq-left}, \text{enq-right}, \text{deq-left}, \text{deq-right}\}$ , *blocks[0].num<sub>x</sub> = 0*)

---

**Algorithm Queue**


---

```

201: void ENQUEUE(Object e)  ▷ Creates a block with element e and appends
    it to the tree.
202:   block newBlock= NEW(LeafBlock)
203:   newBlock.element= e
204:   newBlock.sumenq= leaf.blocks[leaf.size].sumenq+1
205:   newBlock.sumdeq= leaf.blocks[leaf.size].sumdeq
206:   leaf.APPEND(newBlock)
207: end ENQUEUE

208: Object DEQUEUE()
209:   block newBlock= NEW(LeafBlock)          ▷ Creates a block
    with null value element, appends it to the tree, computes its order among
    operations, then computes and returns its response.
210:   newBlock.element= null
211:   newBlock.sumenq= leaf.blocks[leaf.size].sumenq
212:   newBlock.sumdeq= leaf.blocks[leaf.size].sumdeq+1
213:   leaf.APPEND(newBlock)
214:   <rdeq, bdeq>= leaf.INDEXDEQ(leaf.size, 1)
    ▷ r is the rank among all dequeues of the dequeue and bdeq is the index of
    the block in the root containing the dequeue.
215:   <renq, benq>= FINDRESPONSE(rdeq, bdeq)
    ▷ renq is the rank of the enqueue whose element is the response to
    the dequeue and bdeq is the index of the block of it in the blocklist. If the
    response is null then rdeq is -1.
216:   if renq== -1 then
217:     output= null
218:     root.blocks[bdeq].numfinished.inc()          ▷ shared counter
219:     if root.blocks[bdeq].numfinished==root.blocks[bdeq].num then
220:       leaf.lastdone= bdeq
221:     end if
222:   else
223:     output= GETENQ(renq, benq)          ▷ getting the reponse's element.
224:     root.blocks[benq].numfinished.inc()
225:     root.blocks[benq].numfinished.inc()
226:     if root.blocks[bdeq].numfinished==root.blocks[bdeq].num then
227:       leaf.lastdone= bdeq
228:     else if root.blocks[br].numfinished==root.blocks[br].num then
229:       leaf.lastdone= benq
230:     end if
231:   end if
232:   return output
233: end DEQUEUE

234: int, int FINDRESPONSE(int i, int b)          ▷ Computes the rank and
    index of the block in the root of the enqueue that is the response of the ith
    dequeue in the root's bth block. Returns <-1,--> if the queue is empty.
235:   if root.blocks[b-1].length + root.blocks[b].numenq - i < 0 then
236:     return <-1,-->
237:   else
    ▷ We call the dequeues that
    return a value non-null dequeues. rth non-null dequeue returns the element
    of th rth enqueue. We can compute # non-null dequeues in the prefix for
    a block this way: #non-null dequeues= length - #enqueues. Note that the
    ith dequeue in the given block is not a non-null dequeue.
238:     renq= root.blocks[b-1].sumenq- root.blocks[b-1].length + i
239:     return <renq, root.blocks.get(enq, renq).order>
240:   end if
241: end FINDRESPONSE

```

---

deqRest

---

**Algorithm Node**


---

```

301: void PROPAGATE()
302:   if not this.REFRESH() then
303:     this.REFRESH()           ▷ Lemma Double Refresh
304:   end if
305:   if this is not root then   ▷ To check a node is the root we can check
                               its index is 0.
306:     this.parent.PROPAGATE()
307:   end if
308: end PROPAGATE

309: boolean REFRESH()
310:   h= size
311:   <new, npleft, npright>= this.CREATEBLOCK(h)  ▷ npleft, npright are
the values read from the children's numpropagateds.
312:   if new.num==0 then return true  ▷ The block contains nothing.
313:   else if root.blocks.tryAppend(new) then
okcas314:     for each dir in {left, right} do
315:       CAS(dir.super[npdir], null, h+1)
316:       CAS(dir.numpropagated, npdir, npdir+1)
317:     end for
318:     CAS(size, h, h+1)
319:     return true
320:   else
321:     CAS(size, h, h+1)           ▷ Even if another process wins, help to
increase the size. It might fell sleep before increasing.
322:     return false
323:   end if
324: end REFRESH

```

↪ Precondition: blocks[start..end] contains a block with field  $f \geq i$

```

325: int BSEARCH(field f, int i, int start, int end)
                               ▷ Does binary search for the value
i of the given prefix sum feild. Returns the index of the leftmost block in
blocks[start..end] whose field f is  $\geq i$ .
326: end BSEARCH

```

---

```

327: <Block, int, int> CREATEBLOCK(int i)
                               ▷ Creates a block to be inserted into this.blocks[i]. Returns
the created block as well as values read from each childnumpropagated field.
The values are used for incrementing children's numpropagateds if the block
was appended to this.blocks successfully.
328:   block newBlock= NEW(block)
329:   newBlock.group= numpropagated
330:   newBlock.order= i
331:   for each dir in {left, right} do
332:     indexlast= dir.size
333:     indexprev= blocks[i-1].enddir
334:     blocklast= dir.blocks[indexlast]
335:     blockprev= dir.blocks[indexprev]
336:     ▷ newBlock includes dir.blocks[indexprev+1..indexlast].
337:     ndir= dir.numpropagated
338:     newBlock.enddir= indexlast
339:     newBlock.sumenq-dir= blocks[i-1].sumenq-dir + blocklast.sumenq
- blockprev.sumenq
340:     newBlock.sumdeq-dir= blocks[i-1].sumdeq-dir + blocklast.sumdeq
- blockprev.sumdeq
341:   end for
342:   if this is root then
343:     newBlock.length= max(root.blocks[i-1].length + b.numenq -
b.numdeq, 0)
344:   end if
345:   return b, npleft, npright
346: end CREATEBLOCK

```

---

**Algorithm Node**

---

↪ Precondition:  $n.blocks[b]$  contains  $i$ th enqueue in the node.

```
401: element GETENQ(int b, int i)
402:   if  $i \leq blocks[b].num_{enq-left}$  then                                ▷  $i$  exists in the left child of this node
403:     subBlock= leftBSEARCH(sumenq, i, blocks[b-1].endleft+1, blocks[b].endleft)
404:     return left.GET(i-left.blocks[subBlock-1].sumenq, subBlock)
405:   else
406:     i= i-blocks[b].numenq-left
407:     subBlock= rightBSEARCH(sumenq, i, blocks[b-1].endright+1, blocks[b].endright)
408:     return right.GET(i-right.blocks[subBlock-1].sumenq, subBlock)
409:   end if
410: end GETENQ
```

↪ Precondition:  $b$ th block of the node has propagated up to the root and  $i$ th dequeue in the node is in  $blocks[b]$ .

```
411: <int, int> INDEXDEQ(int b, int i)                                     ▷ Returns the order in the root of  $i$ th dequeue in the  $b$ th block of node  $n$  among dequeues.
412:   dir= (parent.left==n)? left: right                                ▷ check if a left or a right child
413:   superBlock= BSEARCH(parent, sumdeq-dir, i, super[blocks[b].group]-p, super[blocks[b].group]+p)
                                                                    ▷ superblock's group has at most  $p$  difference with the value stored in super[].
414:   if dir is left then
415:     i+= parent.blocks[superBlock-1].sumdeq-right
416:   else
417:     i+= parent.blocks[superBlock-1].sumdeq + blocks[superBlock].sumdeq-left                ▷ consider dequeues from the right child
418:   end if
419:   return INDEXDEQ(this.parent, superBlock, i)
420: end INDEX
```

---

---

**Algorithm Leaf**

---

```
501: void APPEND(block blk)
502:   size+=1                                                         ▷ Lines 503 to 502 are done by one process at time.
503:   blk.group= size                                                  ▷ Append is only called by the owner of the leaf.
504:   blocks[size]= blk
505:   parent.PROPAGATE()
506: end APPEND

507: element GETENQ(int b, int i)
508:   return blocks[b].element
509: end GETENQ
```

---

---

**Algorithm Root**

---

```
601: <int, int> INDEXDEQ(int b, int i)
602:   return <i, b>
603: end INDEXDEQ
```

---

appendEnd
pendStart

appendEnd

---

**Algorithm** BlockList

---

▷ : Supports two operations `blocks.tryAppend(Block b)`, `blocks[i]`. Initially empty, when

`blocks.tryAppend(b)` returns true `b` is appended to the end of the list and `blocks[i]` returns  $i$ th block in the blocks. If some instance of `blocks.tryAppend(b)` returns false there is a concurrent instance of `blocks.tryAppend(b')` which has returned true. `blocks[0]` contains an empty block with all fields equal to 0 and `endleft`, `endright` pointers to the first block of the corresponding children.

◇ *PBRT implementation*

```
701: boolean TRYAPPEND(block blk)
702:   TODO
703: end TRYAPPEND
```

◇ *Array implementation*

```
704: boolean TRYAPPEND(block blk)
705:   TODO
706: end TRYAPPEND
```

---

---

<pre>► PRBTree[rootBlock] A persistant red-black tree supporting append(b, key), get(key=i), split(j). append(b, key) returns true in case successful. Since order, sum<sub>eng</sub> are both strictly increasing we can use one of them for another. 801: void RBTAPPEND(block b)          ▷ adds block b to the root.blocks 802:   step= root.size 803:   if step%p<sup>2</sup>==0 then ▷ Help every often <math>p^2</math> operations appended to the       root. Used in lemma's using the size of the PBRT. 804:     Help() 805:     CollectGarbage() 806:   end if 807:   b.num<sub>finished</sub>= 0 808:   return root.blocks.append(b, b.order) 809: end RBTAPPEND 810: void HELP          ▷ Helps pending operations 811: for leaf l in leaves do ▷ if the tree is implemented with an array we       can iterate over the second half of the array. 812:   last= l.size-1    ▷ l.blocks[last] can not be null because of       lines 503-502. 813:   if l.blocks[last].element==null then ▷ operation is dequeue 814:     goto 215 with these values &lt;&gt;    ▷ run Dequeue() for       l.ops[last] after Propagate(). TODO 815:     l.responses[last]= response</pre>	<pre>816:   end if 817:   end for 818: end HELP 819: void COLLECTGARBAGE          ▷ Collects the root blocks that are done. 820:   s=FindMostRecentDone(Root.Blocks.root) ▷ Lemma: If block b is       done after helping then all blocks before b are done as well. 821:   t1,t2= RBT.split(order, s) 822:   RBTRoot.CAS(t2.root) 823: end COLLECTGARBAGE 824: Block FINDMOSTRECENTDONE(b) 825: for leaf l in leaves do 826:   max= Max(l.maxOld, max) 827: end for 828: return max          ▷ This snapshot suffies. 829: end FINDYOUNGESTOLD 830: response FALLBACK(op i)          ▷ really necessary? 831: if root.blocks.get(num<sub>eng</sub>), i is null then ▷ this enqueue was already       finished 832:   return this.leaf.response(block.order) 833: end if 834: end FALLBACK</pre>
---	--

---