

- 1 Feilds
- 2 Queue
- 3 Search+Append
- 4 Propagate
- 5 Index+Get
- 6 MaxByProcess
- 7 Help
- 8 FreeMemory
- 9 PBRT

Algorithm Tree Fields Description

◇ *Shared*

- A binary tree of Nodes with one leaf for each process. root is the root node.
- *MaxbyProcess lastDequeuedFrom* Index of the most recent block in the root that has been dequeued from.

◇ *Local*

- *Node leaf*: process's leaf in the tree.

► *Node*

- **Node left, right, parent* : Initialized when creating the tree.
- *PBRT blocks* : Initially `blocks[0]` contains an empty block with all fields equal to 0.
- *int head= 1*: #blocks in `blocks`. `blocks[0]` is a block with all integer fields equal to zero.

► *Block*

- *int super* : approximate index of the superblock, read from `parent.head` when appending the block to the node

► *InternalBlock* extends *Block*

- *int end_{left}, end_{right}* : indices of the last subblock of the block in the left and right child
- *int sum_{enq-left}*: #enqueues in `left.blocks[1..endleft]`
- *int sum_{deq-left}*: #dequeues in `left.blocks[1..endleft]`
- *int sum_{enq-right}*: #enqueues in `right.blocks[1..endright]`
- *int sum_{deq-right}*: #dequeues in `right.blocks[1..endright]`

► *LeafBlock* extends *Block*

- *Object element* : Each block in a leaf represents a single operation. If the operation is `enqueue(x)` then `element=x`, otherwise `element=null`.
- *int sum_{enq}, sum_{deq}* : # enqueue, dequeue operations in this block and its previous blocks in the leaf
- *object response*

► *RootBlock* extends *InternalBlock*

- *int size* : size of the queue after performing all operations in this block and its previous blocks in the root
-

Algorithm Queue

```
1: void Enqueue(Object e)                                ▷ Creates a block with element e and adds it to the tree.
2:   block newBlock= new(LeafBlock)
3:   newBlock.element= e
4:   newBlock.sumenq= leaf.blocks[leaf.head].sumenq+1
5:   newBlock.sumdeq= leaf.blocks[leaf.head].sumdeq
6:   leaf.Append(newBlock)
7: end Enqueue

▷ Creates a block with null value element, appends it to the tree and returns its response.

8: Object Dequeue()
9:   block newBlock= new(LeafBlock)
10:  newBlock.element= null
11:  newBlock.sumenq= leaf.blocks[leaf.head].sumenq
12:  newBlock.sumdeq= leaf.blocks[leaf.head].sumdeq+1
13:  leaf.Append(newBlock)
14:  <b, i>= IndexDequeue(leaf.head, 1)
15:  output= FindResponse(b, i)
16:  return output
17: end Dequeue

▷ Returns the response to  $D_i(root, b)$ , the  $i$ th Dequeue in  $root.blocks[b]$ .

18: element FindResponse(int b, int i)
19:   if root.blocks[b-1].size + root.blocks[b].numenq - i < 0 then           ▷ Check if the queue is empty.
20:     lastDequeudFrom.update(b)
21:     return null
22:   else                               ▷ The response is  $E_e(root)$ , the  $e$ th Enqueue in the root.
23:     e= i + (root.blocks[b-1].sumenq-root.blocks[b-1].size)
24:     <x, y>= root.DoublingSearch(e, b)
25:     lastDequeudFrom.update(x)
26:     return root.GetEnqueue(x,y)
27:   end if
28: end FindResponse
```

Algorithm *Node*

↪ Precondition: `blocks[start..end]` contains a block with `sumenq` greater than or equal to `x`

▷ Update needed: search on RBT does not need start and end, we can search over whole the red-black tree..

```
26: int BinarySearch(int x)
27:   return min{j: blocks[j].sumenq ≥ x}
28: end BinarySearch
```

Algorithm *Root*

↪ Precondition: `root.blocks[end].sumenq ≥ e`

▷ Returns `<b,i>` such that $E_e(\text{root}) = E_i(\text{root}, b)$, i.e., the `e`th Enqueue in the root is the `i`th Enqueue within ▷ the `b`th block in the root.

```
37: <int, int> DoublingSearch(int e, int end) ▷ I think with garbage collection, doubling search is not needed any more
    and a binary search on root.blocks would be good enough.
```

```
38: end DoublingSearch
```

Algorithm *Leaf*

```
46: void Append(block B) ▷ Only called by the owner of the leaf.
```

```
47:   blocks.TryAppend(B, head)
```

```
48:   head= head+1
```

```
49:   parent.Propagate()
```

```
50: end Append
```

Algorithm *Node*

▷ *n*.Propagate propagates operations in **this**.children up to **this** when it terminates.

51: void Propagate()

52: **if not** Refresh() **then**

53: Refresh()

54: **end if**

55: **if this is not root then**

56: parent.Propagate()

57: **end if**

58: **end** Propagate

▷ Creates a block containing new operations of **this**.children, and then tries to append it to **this**.

59: boolean Refresh()

60: h= head

61: **for each** dir in {left, right} **do**

62: h_{dir}= dir.head

63: **if** dir.blocks[h_{dir}] != null **then**

64: dir.Advance(h_{dir})

65: **end if**

66: **end for**

67: new= CreateBlock(h)

68: **if** new.num==0 **then return** true

69: **end if**

70: result= blocks.TryAppend(new, h)

71: this.Advance(h)

72: **return** result

73: **end** Refresh

Algorithm *Node*

```
74: void Advance(int h)                                ▷ Sets blocks[h].super and increments head from h to h+1.
75:   hp = parent.head
76:   blocks[h].super.CAS(null, hp)
77:   head.CAS(h, h+1)
78: end Advance

79: Block CreateBlock(int i)                            ▷ Creates and returns the block to be installed in blocks[i].
80:   block new = new(InternalBlock)
81:   for each dir in {left, right} do
82:     indexprev = blocks[i-1].enddir
83:     new.enddir = dir.head-1                          ▷ new contains dir.blocks[blocks[i-1].enddir..dir.head-1].
84:     blockprev = dir.blocks[indexprev]
85:     blocklast = dir.blocks[new.enddir]
86:     new.sumenq-dir = blocks[i-1].sumenq-dir + blocklast.sumenq - blockprev.sumenq
87:     new.sumdeq-dir = blocks[i-1].sumdeq-dir + blocklast.sumdeq - blockprev.sumdeq
88:   end for
89:   if this is root then
90:     new.type = InternalBlock-->RootBlock
91:     new.size = max(root.blocks[i-1].size + new.numenq - new.numdeq, 0)
92:   end if
93:   return new
94: end CreateBlock

95: int GetLastDequeuedFrom                            ▷ Returns the index that is safe to remove the blocks before that in the node.
96:   x = lastDequeuedFrom.Get()-1
97:   n = root
98:   while n != this do
99:     dir = left (if this is in left subtree of n): otherwise dir = right
100:    x = n.blocks[x].enddir
101:   end while
102: end GetLastDequeuedFrom
```

Algorithm

Node

↪ Precondition: $\text{blocks}[b].\text{num}_{\text{enq}} \geq i \geq 1$

```
95: element GetEnqueue(int b, int i) ▷ Returns the element of  $E_i(\text{this}, b)$ .
```

```

96:   if this is leaf then

```

```
97:         return blocks[b].element
```

```

98:   else if i <= blocks[b].numenq-left then                                ▷  $E_i(\text{this}, b)$  is in the left child of this node.

```

```
99:      subblockIndex= left.BinarySearch(i+blocks[b-1].sumenq-left, blocks[b-1].endleft+1,
      blocks[b].endleft)
```

▷ start and end values are not needed anymore?

```
100:         return left.GetEnqueue(subblockIndex, i)
```

```
101:     else
```

```
102:         i= i-blocks[b].numenq-left
```

```
103:     subblockIndex= right.BinarySearch(i+blocks[b-1].sumenq-right, blocks[b-1].endright+1,
        blocks[b].endright)
```

▷ start and end values are not needed anymore?

```
104:         return right.GetEnqueue(subblockIndex, i)
```

```
105:   end if
```

```

106: end GetEnqueue

```

↪ Precondition: bth block of the node has propagated up to the root and $\text{blocks}[b].\text{num_deq} \geq i$.

```
107: <int, int> IndexDequeue(int b, int i)    ▷ Update needed: return null when superblock in the root was not found.
```

```
108:    if this is root then
```

```
109:     return <b, i>
```

```
110:     else
```

```
111:         dir= (parent.left==n ? left: right)
```

[illegible]

```
113:         if dir is left then
```

```
114:         i+= blocks[b-1].sum_deq-parent.blocks[superblockIndex-1].sum_deq-left
```

```
115:         else
```

```
116:         i+= blocks[b-1].sum_deg-parent.blocks[superblockIndex-1].sum_deg-right
```

```
117:         i+= parent.blocks[superblockIndex].num_deq-left
```

```
118:         end if
```

```
119:     return this.parent.IndexDeque(superblockIndex, i)
```

```
120:    end if
```

```
121: end IndexDeque
```

Algorithm *MaxByProcess*

```
122: int[p] lastDequeuedbyProcess

123: int Get
124:   return max(lastDequeuedbyProcess)
125: end Get

126: Update(int b)
127:   if lastDequeuedbyProcess[pid]<b then
128:     lastDequeuedbyProcess[pid]=b
129:   end if
130: end Update
```

Algorithm *Tree*

```
131: int Help
132:   for each process P
133:     h=P.leaf.head
134:     if P.leaf.blocks[h].num_deq==1 and P.leaf.IndexDequeue(h,1)!=null then
135:       <b, i>= IndexDequeue(h, 1)
136:       output= FindResponse(b, i)
137:       P.leaf.blocks[h].response= output
138:     end if
139:   end for
140: end Help
```

Algorithm *Node*

```
141: FreeMemory(int b)
142:   if not leaf then
143:     left.FreeMemory(blocks[i].end_left-1)
144:     right.FreeMemory(blocks[i].end_right-1)
145:   end if
146:   blocks= blocks.splitGreater(i)
147: end FreeMemory
```

▷ I think CAS is not needed.

Algorithm *PBRT*

```
PBRT prbt
nodes store <key, sumenq-> block
[i] -> GetByBlock(i)

141: GetByBlock(int i)
142:   return rbt.get(i)
143:   if not found then
144:     return written response
145:   end if
146: end GetByBlock

147: TryAppend(block B, int i)
    Tries to append B with key i to the red-black tree.
148:   if i%p2=0 then
149:     Help()
150:     root.FreeMemory()(lastDequeuedFrom.Get()-1)
151:     garbageCollectRound=floor(root.head/p2)
152:   end if
153: end FreeMemory
```

10 Description

In our original algorithm an **Enqueue** or a **Dequeue** remains in the **blocks** array in the tree nodes even after they terminate. This makes the space used by the algorithm factor of the number of operations of invoked on the queue. In this section here we want to free the memory allocated by the operations that are no longer needed and make the space used polynomial of $p + q$.

One way of handling garbage collection is to reallocate the space taken by each operation right after it is not needed anymore, but it might cost too much to do that. So we garbage collect by batching. If we garbage collect the unnecessary blocks in a node every p^2 block appended to the node, the garbage collection cost is amortized over p^2 blocks which is $O(p^3)$ and $\Omega(p^2)$.

In our design a process garbage collects a node when it attempts to append a block in the kp^2 position in the. **Garbage Collect** corresponded to the kp^2 th block in node n is called the k th round of garbage collect on n . Every p^2 block appended to a node at least one garbage collect happens because the node's **head** cannot advance until the garbage collect is done (see Lines).

Lemma 1. *The number of the blocks in the **blocks** is $O(p^2 + q)$.*

An **Enqueue** operation can be removed from the tree after its **element** has been computed to be the response of some **Dequeue**. It is safe to remove a **Dequeue** operation from the tree after the **Dequeue** is terminated, but a **Dequeue** might sleep for a long time and prevents to be garbage collected. In that situation other processes can help the **Dequeue** to compute its response and write down its response somewhere. If we remove a **Dequeue** from the tree after it is helped the **Dequeue** can read the helped response written when it encountered a problem (computing the **Dequeue** index in the **root** or getting the response **Enqueue**). A **Dequeue** operation can be removed from the tree after its response has been computed. We say a block is *finished* if all of its operations can be removed.

Lemma 2. *There are at most $O(p^2 + q)$ unfinished blocks in a node's **blocks** at a time.*

If the i th **Enqueue** gets dequeued in a FIFO queue, it means 1 to $i - 1$ th **Enqueue** operations are also dequeued. This gives us the idea that if a block is finished then all the blocks before it are also finished. If an operation in a block goes to sleep for a long time then other processes help the operation so the block get finished. Note that when an **Enqueue** operation in a block is not finished the block cannot be finished until some **Dequeue** dequeues that **Enqueue**. Since there are at most p idle operations, we can help them before garbage collection and then remove all the finished blocks safely.

Lemma 3. *If all current operations are helped, then there is a block in the root that all of its previous blocks are finished. That block is the most recent block that has been dequeued from.*

The idea above leads us to a poly-log data structure that supports throwing away all the blocks with keys smaller than an index. Red-black trees do this for us. `Get(i)`, `Append()` and `Split(i)` are logarithmic in block trees. We can create a shared red-black tree just creating a new path for the operation and then using CAS to change the root of the tree. See [this] for more.

Observation 4. *PBRT supports poly-log operations*

Lemma 5. *If we replace the arrays we used to implement `blocks` with red-black trees the amortized complexity of the algorithm would be $PolyLog(p, q)$. And also the algorithm is correct.*

We can help a `Dequeue` by computing its response and writing it down. If the process in future failed to execute, it can read the helped value written down.

Lemma 6. *The `response` written is correct.*

But how can we know which blocks in each node are finished or not?

Lemma 7. *If all current operations are helped, then the blocks before the newest block that some `Enqueue` has been dequeued from is safe to remove. If the most current `Dequeue` returned `null` then all the blocks before the block containing the `Dequeue` can be removed.*

There is a shared array among processes which they write the last block dequeued from in it.

Lemma 8. *`GetLastDequeuedFrom(n - index of the last finished block in the node n)` is $O(p)$.*

Lemma 9. *After `FreeMemory`, the space taken by each node of the tree is $O(p^2 + q)$. The total space in the tree is $PolyLog(p + q)$*