

Introducing the Best Shared Queue Ever

February 14, 2022

1 Description of the algorithm

1.1 Queue

1.1.1 Definition

Queue is an ordered set of elements.

1.1.2 Methods

`Enqueue(element v)` Appends `v` after the last appended element to the queue.

`Dequeue()` Removes the first not dequeued element in the queue and returns it.




Figure 1: Queue

1.1.3 Implementation

We use a `Block tree` and a `Queue helper` implement a `Queue`. To do an enqueue, it is sufficient to append an enqueue operation to the block tree. For dequeues, we do the same but at last, compute the head using `Queue helper`, and if the queue is not empty, we return the corresponding enqueue argument.

1.1.4 Complexity

`Enqueue(v)` and `Dequeue()` each take $O(\log Q + \log^2 p)$ steps.

1.2 Block tree

1.2.1 Definition

`Block tree` is a data structure that linearizes invocations of operations by some processes.

Linearization Operations are linearized when they are added to the root. The operations in a root block are linearized as the block's ordering. ~~TODO~~problem: where should we define the linearization ordering?

1.2.2 Methods

`Append(operation op)` Appends op to the linearization.

`Get(int i)` Returns ith linearized operation.

`Index(operation op)` If `Append(op)` is terminated, `Index(op)` returns the order of op in the linearization.

`Mark(int i)` If operation i is marked `Get(i)` is not going to be invoked in future anymore.

1.2.3 Implementation

`Block tree` is a binary tree of nodes, such that a leaf is assigned to each process. Process p adds op to its leaf and propagates it to the root node, which stores the total ordering.




Figure 2: Block tree

Append(op) Adds a block for op to the invoking process's leaf.

Propagate(node n) When **Propagate(node n)** is terminated, operations in n before calling **Propagate(node n)** are in the root.

1. Refresh(n)
2. If 1 was unsuccessful, do Refresh(n) one more time.
3. Propagate(n.parent)

Get(i)

1. B = find the root block of containing i th op
2. Get(B, i-pre(B), root)

pre(B) is the number of operations before block B in B's node.

Index(op)

1. B = the block in a leaf containing operation op
2. Index(n, B, i+pre(B))

`Mark(i)` Mark i in the root node.

1.2.4 Complexity

`Append(op)`, `Get(i)` and `Index(op)` each take $O(\log^2 p)$, $O(\log(n - \text{\#marked}) + \log^2 p)$, $O(\log^2 p)$ steps.

1.3 Queue helper

1.3.1 Definition

Queue helper tells us the head of the queue at the time some dequeue operation has been linearized.

1.3.2 Methods

`ComputeHead(i)` Computes head of the queue when i th dequeue occurs.

`Augment(i)` Augments root block i .

1.3.3 Implementation

By adding the size and number of non-null dequeues in a block, we can compute the head in constant time.

1.3.4 Complexity

Constant time.

1.4 Node

1.4.1 Definition

A `node` contains the ordering of the blocks appended to it by time.

1.4.2 Methods

`Append(Block b)` Adds block b to the end of the `node`'s ordering.

`GetHead()` Returns the index of the last block added.

`Refresh(Node n)` Creates a block from n 's children's blocks which are not already in n , and tries to append it to n .

`Get(B, i, n)` Returns the i th operation in block B in node n .

`Index(n, B, i)` Returns the superblock in n .parent containing block B in node n .

1.4.3 Implementation




Figure 3: Node

`blocks[i]` stores the i th block. `time` shows the number of groups propagated to the parent. `super[i]` is the index of the superblock of the group i .

`Append(Block b)` CASes b to the first empty cell in A .

`GetHead()` Returns the head field, it is updated every successful `Refresh()`.

`Get(B, i, n)`

1. SB = find the subblock of B containing i th op
2. `Get(SB, i-pre(SB), B's node)`

`Index(n, B, i)`

1. SB = find the block in n 's parent that contains i th op in block B of node n
2. `Index(n, B, i+pre(SB))`

Refresh()

1. Create a block from n 's children's blocks which are not already in n
2. Append the block to n
3. If 2 was successful, update the head field

1.4.4 Requirements

If Append() fails it means another successful instance of Append exists that has some time in common.

1.4.5 Complexity

Each level of Get() takes $\log p$ steps, and there are $\log p$ levels. So Get may take $O(\log^2 p)$ steps. Complexity of finding the superblock of a block is $O(\log p)$.

1.5 Root node

A data structure that contains an ordering of not-marked blocks appended to it.

Append(Block b) Adds block b to the end of the Root node's ordering.

Mark(Block b) Marks block b .

Get(i) Returns the block containing i th operation.

1.6 Block

1.6.1 Definition

A block is an ordering of some operations in a node. We can merge some blocks into a block, the initial blocks are called subblocks, and the latter is called superblock.

TODO problem: where should we put end pointer of blocks that help get become faster?

1.6.2 Methods

CreateBlock(operation op) Returns a block that contains op .

`CreateBlock(node n, int start, int last)` Returns a block that contains the blocks from n.start to n.last. The order of the input blocks remains the same.

`GetEnd()` Returns the last left and right subblock of the block.

1.6.3 Implementation




Figure 4: Block

We can do the told before methods on blocks only by knowing

- | | |
|----------------------------------|---|
| 1. #enqueue operation from left | 5. index of the last block from left |
| 2. #enqueue operation from right | 6. index of the last block from left |
| 3. #dequeue operation from left | 7. group number which tells the blocks which have propagated to the parent together |
| 4. #dequeue operation from right | |

TODO problem: how to explain the separation of the enqueues from the dequeues?
till now, we considered all operations from just one type

1.7 Operation

1.7.1 Implementation

Each operation contains the arguments of the operation, the invoking process, and the position in the invoking's process operations.