

A Wait-free Queue with Polylogarithmic Step Complexity

Hossein Naderibeni¹ and Eric Ruppert^{1*}

¹York University, Toronto, Canada.

*Corresponding author(s). E-mail(s): ruppert@eecs.yorku.ca;

Abstract

We present a novel linearizable wait-free queue implementation using single-word CAS instructions. Previous lock-free queue implementations from CAS all have amortized step complexity of $\Omega(p)$ per operation in worst-case executions, where p is the number of processes that access the queue. Our new wait-free queue takes $O(\log p)$ steps per enqueue and $O(\log^2 p + \log q)$ steps per dequeue, where q is the size of the queue. A bounded-space version of the implementation has $O(\log p \log(p + q))$ amortized step complexity per operation.

1 Introduction

There has been a great deal of research in the past several decades on the design of shared queues. Besides being a fundamental data structure, queues are used in significant concurrent applications, including OS kernels [33], memory management (e.g., [5]), synchronization [30], and sharing resources or tasks. We focus on shared queues that are *linearizable* [18], meaning that operations appear to take place atomically, and *lock-free*, meaning that some operation on the queue is guaranteed to complete regardless of how asynchronous processes are scheduled to take steps. We study the *amortized step complexity* per operation, which is measured by taking

the maximum, over all possible executions, of the number of steps in the execution divided by the total number of enqueue and dequeue operations in the execution.

The lock-free *MS-queue* of Michael and Scott [34] is a classic shared queue implementation. It uses a singly-linked list with pointers to the front and back nodes. To dequeue or enqueue an element, the front or back pointer is updated by a compare-and-swap (CAS) instruction. If this CAS fails, the operation must retry. In the worst case, this means that each successful CAS may cause all other processes to fail and retry, leading to an amortized step complexity of $\Omega(p)$ per operation in a system of p processes. Numerous

papers have suggested modifications to the MS-queue [19, 27, 28, 31, 35, 36, 42], but all still have $\Omega(p)$ amortized step complexity as a result of contention on the front and back of the queue. Morrison and Afek [37] called this the *CAS retry problem*. The same problem occurs in array-based implementations of queues [7, 15, 45, 49]. Solutions that tried to sidestep this problem using fetch&increment [37, 39, 40, 50] rely on slower mechanisms to handle worst-case executions and still have $\Omega(p)$ step complexity.

Many concurrent data structures that keep track of a set of elements also have an $\Omega(p)$ term in their step complexity, as observed by Ruppert [44]. For example, lock-free lists [14, 46], stacks [48] and search trees [10] have an $\Omega(c)$ term in their step complexity, where c represents contention, the number of processes that access the data structure concurrently, which can be p in the worst case. Attiya and Fournier [3] proved that amortized $\Omega(c)$ steps per operation are indeed necessary for any CAS-based implementation of a lock-free bag data structure, which provides operations to insert an element or remove an arbitrary element (chosen non-deterministically). Since a queue trivially implements a bag, this lower bound also applies to queues. Although this might seem to settle the step complexity of lock-free queues, the lower bound holds only if c is $O(\log \log p)$ so it should

be stated more precisely as an amortized bound of $\Omega(\min(c, \log \log p))$ steps per operation.

We exploit this loophole. We show it is, in fact, possible for a linearizable queue to have step complexity sublinear in p . Our queue is the first whose step complexity is polylogarithmic in p and in q , the number of elements in the queue. It is also *wait-free*, meaning that every operation is guaranteed to complete within a finite number of its own steps. For ease of presentation, we first give an unbounded-space construction where enqueues take $O(\log p)$ steps and dequeues take $O(\log^2 p + \log q)$ steps, and then modify it to bound the space while having $O(\log p \log(p + q))$ amortized step complexity per operation. Moreover, each operation does $O(\log p)$ CAS instructions in the worst case, whereas previous lock-free queues use $\Omega(p)$ CAS instructions, even in an amortized sense. Since a queue is also a bag, our queue is the first lock-free bag with polylogarithmic step complexity.

Both versions of our queue use single-word CAS on reasonably-sized words. We assume that a word is large enough to store an item to be enqueued (or at least a pointer to it). We also assume that the number of operations performed on the queue can be stored (in binary) in $O(1)$ words. This is analogous to the assumption for the classical RAM model that the number of bits per word is logarithmic in the problem size. For the space-bounded version, we unlink

unneded objects from our data structure. We do not address the orthogonal problem of reclaiming memory; we assume a safe garbage collector, such as the highly optimized one that Java provides.

Our queue uses a binary tree, called the *ordering tree*, where each process has its own leaf. A process adds its operations to its leaf. As in previous work (e.g., [1, 21]), operations are propagated from the leaves up to the root in a cooperative way that ensures wait-freedom and avoids the CAS retry problem. Operations reach the root in batches and are linearized in the order they reach the root. Since a batch can have up to p operations, explicitly recording the list of operations composing a batch or applying them one-by-one to the queue would be too costly. Instead, we use a novel implicit representation of batches of queue operations that allows us to quickly merge two batches from the children of a node, and quickly access any operation in a batch. A preliminary version of this work appeared in [38].

2 Related Work

List-based Queues.

The MS-queue [34] is a lock-free queue that has stood the test of time. The standard Java Concurrency Package includes a version of it. See [34] for a survey of the early history of concurrent queues. As mentioned above, the MS-queue suffers from the CAS retry problem because of contention at

the front and back of the queue. Thus, it is lock-free but not wait-free and has an amortized step complexity of $\Theta(p)$ per operation.

Many papers have described ways to reduce contention in the MS-queue. Moir et al. [36] added an elimination array that allows an enqueue to pass its enqueued value directly to a concurrent dequeue when the queue is empty. However, when there are p concurrent enqueues (and no dequeues), the CAS retry problem is still present. The baskets queue of Hoffman, Shalev, and Shavit [19] attempts to reduce contention by grouping concurrent enqueues into baskets. An enqueue that fails its CAS is put in the basket with the enqueue that succeeded. Enqueues within a basket order themselves without having to access the back of the queue. However, if p concurrent enqueues are in the same basket the CAS retry problem occurs when they order themselves using CAS instructions. Both modifications still have $\Omega(p)$ amortized step complexity.

Kogan and Herlihy [27] improved the MS-queue's performance using *futures*. Operations return future objects instead of responses. Later, when an operation's response is needed, it is evaluated using the future object. This allows batches of enqueues or dequeues to be done at once on an MS-queue. However, the implementation satisfies a weaker correctness condition than linearizability. Milman-Sela et al. [35] extended this approach to allow batches to mix enqueues and dequeues.

In the worst case, where operations require their response right away, batches have size 1, and both of these implementations behave like a standard MS-queue.

In the MS-queue, an enqueue requires two CAS steps. Ladan-Mozes and Shavit [31] presented an optimistic queue implementation that uses a doubly-linked list to reduce the number of CAS instructions to one in the best case. Pointers in the doubly-linked list can be inconsistent, but are fixed when necessary by traversing the list. This fixing is rare in practice, but it yields an amortized complexity of $\Omega(qp)$ steps per operation in the worst case.

Kogan and Petrank [28] used Herlihy’s helping technique [17] to make the MS-queue wait-free. Then, they introduced the fast-path slow-path methodology [29] for making data structures wait-free: the fast path has good performance and the slow path guarantees termination. They applied their methodology to combine the MS-queue (as the fast path) with their wait-free queue (as the slow path). Ramalhete and Correia [42] added a different helping mechanism to the MS-queue. Although these approaches can perform well in practice, the amortized step complexity remains $\Omega(p)$.

Haas [16] described a queue implementation based on timestamping. To enqueue an item, a process adds the item, together with an associated timestamp, to the process’s own single-enqueuer

multi-dequeuer queue, which is implemented as a linked list. Items added by concurrent enqueues may get the same timestamp. Although enqueues are very efficient, dequeues look at all p lists to find and return an item with the oldest timestamp. This is compounded by the fact that dequeues may compete to claim the same item using a CAS, resulting in a CAS retry problem. As a result, the amortized step complexity for dequeue operations is $\Omega(p^2)$.

Array-Based Queues.

Arrays can be used to implement queues with bounded capacity [7, 41, 45, 49]. Dequeues and enqueues update indices of the front and back elements using CAS instructions. Gidenstam, Sundell, and Tsigas [15] avoid the capacity constraint by using a linked list of arrays. These solutions also use $\Omega(p)$ steps per operation due to the CAS retry problem.

Morrison and Afek [37] also used a linked list of (circular) arrays. To avoid the CAS retry problem, concurrent operations try to claim spots in an array using fetch&increment instructions. (It was shown recently that this implementation can be modified to use single-word CAS instructions rather than double-width CAS [43].) If livelock between enqueues and a racing dequeue prevent enqueues from claiming a spot, the enqueues fall back on using a CAS to add a new array to the linked list, and the CAS retry problem reappears.

This approach is similar to the fast-path slow-path methodology [29]. Other array-based queues [39, 40, 50] also used this methodology. In worst-case executions that use the slow path, they also take $\Omega(p)$ steps per operation, due either to the CAS retry problem or helping mechanisms.

Universal Constructions.

One can also build a queue using a universal construction [17]. Afek, Dauber, and Touitou’s universal construction [1] introduced the technique where a process must climb a binary tree while helping any other processes it sees along the way to reach the root, and this technique is the basis of our queue implementation. Jayanti [20] observed that their construction can be modified to use $O(\log p)$ steps per operation, assuming that words can store $\Omega(p \log p)$ bits. However, if more reasonably-sized $O(\log p)$ -bit words are used, the construction would take $\Omega(p \log p)$ steps per operation. There are two obstacles to making this construction more efficient: it is expensive to manipulate lists of up to p operations that must be propagated along a path up to the root, and when a batch of up to p operations reaches the root, it is expensive to perform all of them on the implemented data structure. In our work, we devise a novel way of doing this implicitly for batches of operations on a queue.

Fatourou and Kallimanis [13] used their own universal construction based on fetch&add and

LL/SC instructions to implement a queue, but its step complexity is also $\Omega(p)$ when words are of reasonable size.

Restricted Queues.

David gave the first sublinear-time queue [8], but it works only for a single enqueueer. It uses fetch&increment and swap instructions and takes $O(1)$ steps per operation, but uses unbounded memory. Bounding the space increases the steps per operation to $\Omega(p)$. Jayanti and Petrovic gave a wait-free polylogarithmic queue [21], but only for a single dequeuer. Like our queue, theirs uses the cooperative tree-climbing technique of Afek, Dauber and Touitou [1]. Concurrently with our work, which first appeared in [38], Johnen, Khatatabi and Milani [23] built on [21] to give a wait-free queue that achieves $O(\log p)$ steps for enqueue operations but fails to achieve polylogarithmic step complexity for dequeues: their dequeue operations take $O(k \log p)$ steps if there are k dequeuers.

Other Primitives.

Li [32] implemented a non-blocking linearizable queue using only the weak primitives test&set, fetch&add and swap. This implementation’s amortized step complexity is not bounded as a function of the number of processes and the size of the queue. (It depends on the total number of enqueues in the execution.)

Khanchandani and Wattenhofer [25] gave a wait-free queue with $O(\sqrt{p})$ step complexity using some strong non-standard synchronization primitives called half-increment and half-max, which can be viewed as double-word read-modify-write operations. They use this as evidence to argue that their primitives can be more efficient than CAS since previous CAS-based queues all required $\Omega(p)$ step complexity. Our new implementation undermines that argument.

Fetch&Increment Objects.

Ellen, Ramachandran and Woelfel [11] gave an implementation of fetch&increment objects that uses a polylogarithmic number of steps per operation. Like our queue, they also use a tree structure similar to the universal construction of [1] to keep track of the operations that have been performed. However, our construction requires more intricate data structures to represent sets of operations, since a queue's state cannot be represented as succinctly as the single-word state of a fetch&increment object. Ellen and Woelfel [12] then gave an improved implementation of fetch&increment with better step complexity.

Lower Bounds.

As mentioned in the introduction, it follows from Attiya and Fourn's lower bound on bag data structures [3] that the amortized step complexity of operations on a queue is $\Omega(\min(c, \log \log p))$,




Fig. 1: Known lower bounds on amortized step complexity of queue operations when contention is c in a system of p processes.

where c is contention. Subsequently, Jayanti, Tarjan and Boix-Adserà [22] proved a $\Omega(\log c)$ lower bound on the amortized step complexity of queue operations. The combination of these lower bounds is shown schematically in Figure 1. If we are interested in the complexity only as a function of p , the second result gives us an $\Omega(\log p)$ lower bound, since c can be as high as p .

3 Queue Implementation

3.1 Overview

Our *ordering tree* data structure is used to agree on a total ordering of the operations performed on the queue. It is a static binary tree of height $\lceil \log_2 p \rceil$ with one leaf for each process. Each tree node stores an array of *blocks*, where each block represents a sequence of enqueues and a sequence of dequeues. See Figure 2 for an example. For ease of presentation, we use an infinite array of blocks in each node in this section. Then, in Section 6,

we describe how to replace the infinite array by a representation that uses bounded space.

To perform an operation on the queue, a process P appends a new block containing that operation to the *blocks* array in P 's leaf. Then, P attempts to propagate the operation to each node along the path from that leaf to the root of the tree. We shall define a total order on all operations that have been propagated to the root, which will serve as the linearization ordering of the operations.

To propagate operations from a node v 's children to v , P performs the following three steps. P first observes the blocks in both of v 's children that are not already in v , creates a new block by combining information from those blocks, and attempts to append this new block to v 's *blocks* array using a CAS. Following [21], we call this three-step sequence a *Refresh* on v . A *Refresh*'s CAS may fail if there is another concurrent *Refresh* on v . However, since a successful *Refresh* propagates multiple pending operations from v 's children to v , we can prove that if two *Refreshes* by P on v fail, then P 's operation has been propagated to v by some other process, so P can continue onwards towards the root.

Now suppose P 's operation has been propagated all the way to the root. If P 's operation is an enqueue, it has obtained a place in the linearization ordering and can terminate. If P 's operation is a dequeue, P must use information in the tree to

compute the value that the dequeue must return. To do this, P first determines which block in the root contains its dequeue (since the dequeue may have been propagated to the root by some other process). P does this by finding the dequeue's location in each node along the path from the leaf to the root. Then, P determines whether the queue is empty when its dequeue is linearized. If so, it returns *null* and we call it a *null dequeue*. If not, P computes the rank r of its dequeue among all non-null dequeues in the linearization ordering. (We say that the r th element in a sequence has *rank* r within that sequence.) P then returns the value of the r th enqueue in the linearization.

We must choose what to store in each block so that the following tasks can be done efficiently.

- (T1) Construct a block for node v that represents the operations in consecutive blocks in v 's children, as required for a *Refresh*.
- (T2) Given a dequeue in a leaf that has been propagated to the root, find that operation's position in the root's *blocks* array.
- (T3) Given a dequeue's position in the root, decide if it is a *null dequeue* (i.e., if the queue is empty when it is linearized) or determine the rank r of the enqueue whose value it returns.
- (T4) Find the r th enqueue in the linearization ordering.

Since these tasks depend on the linearization ordering, we describe that ordering next.




Fig. 2: An example ordering tree with four processes. $E(x)$ denotes an Enqueue(x) operation and D_1 to D_6 denote Dequeue operations. We show explicitly the enqueue sequence and dequeue sequence represented by each block in the *blocks* arrays of the seven nodes. The leftmost element of each *blocks* array is a dummy block. Arrows represent the indices stored in *end_{left}* and *end_{right}* fields of blocks (as described in Section 3.3). The fourth process's D_6 is still propagating towards the root. The linearization order for this tree is $E(a) E(e) D_2 \mid E(b) D_4 D_5 \mid E(d) E(f) E(h) D_1 \mid E(c) D_3 \mid E(g)$, where vertical bars indicate boundaries of blocks in the root.

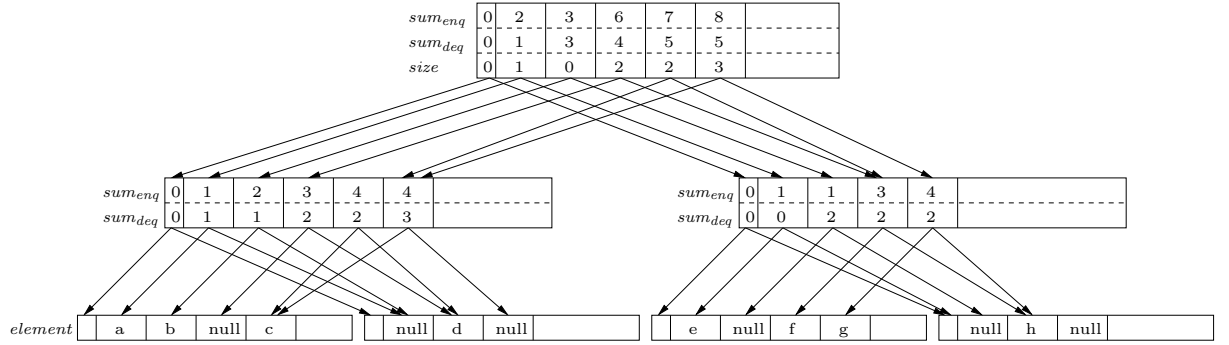


Fig. 3: The actual, implicit representation of the tree shown in Figure 2. The leaf blocks simply show the *element* field. Internal blocks show the *sum_{enq}* and *sum_{deq}* fields, and *end_{left}* and *end_{right}* fields are shown using arrows as in Figure 2. Root blocks also have the additional *size* field. The *super* field is not shown.

3.2 Linearization Ordering

Performing a double Refresh at each node along the path from the leaf to the root ensures a block containing the operation is appended to the root before the operation completes. So, if an operation op_1 terminates before another operation op_2 begins, op_1 will be in an earlier block than op_2 in the root's blocks array. Thus, we linearize operations according to the block they belong to in the

root's array. We can choose how to order operations within the same block, since they must be concurrent.

Each block in a leaf represents one operation. Each block B in an internal node v results from merging several consecutive blocks from each of v 's children. The merged blocks in v 's children are called the *direct subblocks* of B . A block B' is a *subblock* of B if it is a direct subblock of B or a

subblock of a direct subblock of B . A block B represents the set of operations in all of B 's subblocks in leaves of the tree. The operations propagated by a Refresh are all pending when the Refresh occurs, so there is at most one operation per process. Hence, a block represents at most p operations in total. Moreover, we never append empty blocks, so each block represents at least one operation and it follows that a block can have at most p direct subblocks.

As mentioned above, we are free to order operations within a block however we like. We order the enqueues and dequeues separately, and put the operations propagated from the left child before the operations from the right child. More formally, we inductively define sequences $E(B)$ and $D(B)$ of the enqueues and dequeues represented by a block B . If B is a block in a leaf representing an enqueue operation, its enqueue sequence $E(B)$ is that operation and its dequeue sequence $D(B)$ is empty. If B is a block in a leaf representing a dequeue, $D(B)$ is that single operation and $E(B)$ is empty. If B is a block in an internal node v with direct subblocks B_1^L, \dots, B_ℓ^L from v 's left child and B_1^R, \dots, B_r^R from v 's right child, then B 's operation sequences are defined by the concatenations

$$\begin{aligned} E(B) &= E(B_1^L) \cdots E(B_\ell^L) \cdot E(B_1^R) \cdots E(B_r^R) \text{ and} \\ D(B) &= D(B_1^L) \cdots D(B_\ell^L) \cdot D(B_1^R) \cdots D(B_r^R) \end{aligned} \quad (3.1)$$

We say the block B *contains* the operations in $E(B)$ and $D(B)$.

When linearizing operations within one of the root's blocks, we choose to put the block's enqueues before its dequeues. Thus, if the root's blocks array contains blocks B_1, \dots, B_k , the linearization ordering is

$$L = E(B_1) \cdot D(B_1) \cdot E(B_2) \cdot D(B_2) \cdots E(B_k) \cdot D(B_k). \quad (3.2)$$

3.3 Representation of Blocks

In this section, we describe how to represent blocks so that tasks (T1) to (T4) can be done efficiently. Each node of the ordering tree has an infinite array called *blocks*. To simplify the code, *blocks*[0] is initialized with an empty block B_0 , where $E(B_0)$ and $D(B_0)$ are empty sequences. Each node's *head* index stores the position in the *blocks* array to be used for the next attempt to append a block.

If a block contained an explicit representation of its sequences of enqueues and dequeues, it would take $\Omega(p)$ time to construct a block, which would be too slow for task (T1). Instead, the block stores an implicit representation of the sequences. We now explain how we designed the fields for this implicit representation. Refer to Figure 3 for an example showing how the tree in Figure 2 is actually represented, and Figure 4 for the definitions of the fields of blocks and nodes.

460	► Node	
461	• Node <i>left, right, parent</i>	▷ tree pointers initialized when creating the tree
462	• Block[0.. ∞] <i>blocks</i>	▷ blocks that have been propagated to this node;
463		▷ <i>blocks</i> [0] is empty block whose integer fields are 0
464	• int <i>head</i>	▷ position to append next block to <i>blocks</i> , initially 1
465	► Block	
466	• int <i>sum_{enq}, sum_{deq}</i>	▷ # of enqueues, dequeues in <i>blocks</i> array up to this block (inclusive)
467	• int <i>super</i>	▷ approximate index of superblock in <i>parent.blocks</i>
468	▷ Blocks in internal nodes have the following additional fields	
469	• int <i>end_{left}, end_{right}</i>	▷ index of last direct subblock in the left and right child
470	▷ Blocks in leaf nodes have the following additional field	
471	• Object <i>element</i>	▷ <i>x</i> for Enqueue(<i>x</i>) operation; otherwise null
472	▷ Blocks in the root node have the following additional field	
473	• int <i>size</i>	▷ size of queue after performing all operations up to end of this block

Fig. 4: Objects used in the ordering tree data structure.

A block in a leaf represents a single enqueue or dequeue. The block's *element* field stores the value enqueued if the operation is an enqueue, or null if the operation is a dequeue.

Each block in an internal node *v* has fields *end_{left}* and *end_{right}* that store the indices of the block's last direct subblock in *v*'s left and right child. Thus, the direct subblocks of *v.blocks*[*b*] are

$$\begin{aligned}
 &v.\text{left}.\text{blocks}[v.\text{blocks}[b-1].\text{end}_{\text{left}} + 1.. \\
 &\quad v.\text{blocks}[b].\text{end}_{\text{left}}] \text{ and} \\
 &v.\text{right}.\text{blocks}[v.\text{blocks}[b-1].\text{end}_{\text{right}} + 1.. \\
 &\quad v.\text{blocks}[b].\text{end}_{\text{right}}]. \quad (3.3)
 \end{aligned}$$

The *end_{left}* and *end_{right}* fields allow us to navigate to a block's direct subblocks. Blocks also store some prefix sums: *v.blocks*[*b*] has two fields *sum_{enq}* and *sum_{deq}* that store the total numbers of enqueues and dequeues in *v.blocks*[1..*b*]. We use these to search for a particular operation. For

example, consider finding the *r*th enqueue *E_r* in the linearization. A binary search for *r* on the *sum_{enq}* fields of the root's blocks finds the block containing *E_r*. If we know a block *B* in a node *v* contains *E_r*, we can use the *sum_{enq}* field again to determine which child of *v* contains *E_r* and then do a binary search of the *sum_{enq}* fields of the direct subblocks of *B* in that child. Thus, we work our way down the tree until we find the leaf block that stores *E_r* explicitly. We shall show that the binary search in the root can be done in $O(\log p + \log q)$ steps, and the binary search within each other node along the path to a leaf takes $O(\log p)$ steps, for a total of $O(\log^2 p + \log q)$ steps for task (T4).

A block is called the *superblock* of all of its direct subblocks. To facilitate navigating upwards in the ordering tree for task (T2), each block *B* has a field *super* that contains the (approximate) index of its superblock in the parent node's *blocks* array (it may differ from the true index by 1). This

allows a process to determine the true location of the superblock by checking the end_{left} or end_{right} values of just two blocks in the parent node. Thus, starting from an operation in a leaf's block, one can use these indices to track the operation up the path to the root, and determine the operation's location in a root block in $O(\log p)$ time.

Now consider task (T3). To determine whether the queue is empty when a dequeue occurs, each block in the root has a *size* field storing the number of elements in the queue after all operations in the linearization up to that block (inclusive) have been done. We can determine which dequeues in a block B_d in the root are null dequeues using $B_{d-1}.size$, which is the size of the queue just before B_d 's operations, and the number of enqueues and dequeues in B_d . Moreover, the total number of non-null dequeues in blocks B_1, \dots, B_{d-1} is $B_{d-1}.sum_{enq} - B_{d-1}.size$. We can use this information to determine the rank of a non-null dequeue in B_d among all non-null dequeues in the linearization, which is the rank (among all enqueues) of the enqueue whose value the dequeue should return.

Having defined the fields required for tasks (T2), (T3) and (T4), we can easily see how to construct a new block B during a **Refresh** in $O(1)$ time. A **Refresh** on node v reads the values h_ℓ and h_r of the *head* fields of v 's children and stores $h_\ell - 1$ and $h_r - 1$ in $B.end_{left}$ and $B.end_{right}$. Then,

we can compute

$$B.sum_{enq} = v.left.blocks[B.end_{left}].sum_{enq} + v.right.blocks[B.end_{right}].sum_{enq}.$$

For a block B in the root, $B.size$ is computed using the *size* field of the previous block B' and the number of enqueues and dequeues in B :

$$B.size = \max(0, B'.size + (B.sum_{enq} - B'.sum_{enq}) - (B.sum_{deq} - B'.sum_{deq})).$$

The only remaining field is $B.super$. When the block B is created for a node v , we do not yet know where its superblock will eventually be installed in v 's parent. So, we leave $B.super$ blank. Soon after B is installed, some process will set $B.super$ to a value read from the *head* field of v 's parent. We shall show that this happens soon enough that $B.super$ can differ from the true index of B' by at most 1.

3.4 Details of the Implementation

We now discuss the queue implementation in more detail. Pseudocode is provided in Figures 5–7.

An **Enqueue**(e) appends a block to the process's leaf. The block's *element* field is e to indicate it represents an **Enqueue**(e) operation. It suffices to propagate the operation to the root and

```

562
563   ▷ Shared variable
564   • Node root                                ▷ root of ordering tree
565   ▷ Thread-local variable
566   • Node leaf                                ▷ process's leaf in the ordering tree
567
568 1: void Enqueue(Object e)
569 2:   let B be a new Block with fields element := e, sumenq := leaf.blocks[leaf.head - 1].sumenq + 1,
570   sumdeq := leaf.blocks[leaf.head - 1].sumdeq
571 3:   Append(B)
572 4: end Enqueue
573
574 5: Object Dequeue()
575 6:   let B be a new Block with fields element := null, sumenq := leaf.blocks[leaf.head - 1].sumenq,
576   sumdeq := leaf.blocks[leaf.head - 1].sumdeq + 1
577 7:   Append(B)
578 8:    $\langle b, i \rangle := \text{IndexDequeue}(\text{leaf}, \text{leaf.head} - 1, 1)$ 
579 9:   return FindResponse(b, i)
580 10: end Dequeue
581
582 11: void Append(Block B)                                ▷ append block to leaf and propagate to root
583 12:   leaf.blocks[leaf.head] := B
584 13:   leaf.head := leaf.head + 1
585 14:   Propagate(leaf.parent)
586 15: end Append
587
588 16: void Propagate(Node v)                                ▷ propagate blocks from v's children to root
589 17:   if not Refresh(v) then                                ▷ double refresh
590 18:     Refresh(v)
591 19:   end if
592 20:   if v ≠ root then                                ▷ recurse up tree
593 21:     Propagate(v.parent)
594 22:   end if
595 23: end Propagate
596
597 24: boolean Refresh(Node v)                                ▷ try to append a new block B to v.blocks
598 25:   h := v.head
599 26:   for each dir in {left, right} do
600 27:     childHead := v.dir.head
601 28:     if v.dir.blocks[childHead] ≠ null then
602 29:       Advance(v.dir, childHead)
603 30:     end if
604 31:   end for
605 32:   B := CreateBlock(v, h)
606 33:   if B = null then return true
607 34:   else
608 35:     result := CAS(v.blocks[h], null, B)
609 36:     Advance(v, h)
610 37:     return result
611 38:   end if
612 39: end Refresh

```

Fig. 5: Queue implementation's main routines.

```

40: Block CreateBlock(Node v, int i)  $\triangleright$  create new block  $B$  for Refresh to install in  $v.blocks[i]$  613
41:   let  $B$  be a new Block with fields  $end_{left} := v.left.head - 1$ ,  $end_{right} := v.right.head - 1$ , 614
       $sum_{enq} := v.left.blocks[B.end_{left}].sum_{enq} + v.right.blocks[B.end_{right}].sum_{enq}$ , 615
       $sum_{deq} := v.left.blocks[B.end_{left}].sum_{deq} + v.right.blocks[B.end_{right}].sum_{deq}$  616
42:    $num_{enq} := B.sum_{enq} - v.blocks[i-1].sum_{enq}$  617
43:    $num_{deq} := B.sum_{deq} - v.blocks[i-1].sum_{deq}$  618
44:   if  $v = root$  then  $B.size := \max(0, v.blocks[i-1].size + num_{enq} - num_{deq})$  619
45:   end if 620
46:   if  $num_{enq} + num_{deq} = 0$  then return null  $\triangleright$  no blocks need to be propagated to  $v$  621
47:   else return  $B$  622
48:   end if 623
49: end CreateBlock 624

50: void Advance(Node v, int h)  $\triangleright$  set  $v.blocks[h].super$  and increment  $v.head$  from  $h$  to  $h+1$  625
51:   if  $v \neq root$  then 626
52:      $h_p := v.parent.head$  627
53:     CAS( $v.blocks[h].super$ , null,  $h_p$ ) 628
54:   end if 629
55:   CAS( $v.head$ ,  $h$ ,  $h+1$ ) 630
56: end Advance 631

57:  $\langle int, int \rangle$  IndexDequeue(Node v, int b, int i) 632
58:    $\triangleright$  return  $\langle b', i' \rangle$  such that  $i$ th dequeue in  $D(v.blocks[b])$  is  $(i')$ th dequeue of  $D(root.blocks[b'])$  633
59:    $\triangleright$  Precondition:  $v.blocks[b]$  is not null, was propagated to root, and contains at least  $i$  dequeues 634
60:   if  $v = root$  then return  $\langle b, i \rangle$  635
61:   else  $\triangleright$  First, find the superblock of  $v.blocks[b]$  636
62:      $dir := (v.parent.left = v ? left : right)$  637
63:      $sup := v.blocks[b].super$  638
64:     if  $b > v.parent.blocks[sup].end_{dir}$  then  $sup := sup + 1$  639
65:   end if 640
66:    $\triangleright$  compute index  $i$  of dequeue in superblock 641
67:    $i += v.blocks[b-1].sum_{deq} - v.blocks[v.parent.blocks[sup-1].end_{dir}].sum_{deq}$  642
68:   if  $dir = right$  then 643
69:      $i += v.blocks[v.parent.blocks[sup].end_{left}].sum_{deq} -$  644
        $v.blocks[v.parent.blocks[sup-1].end_{left}].sum_{deq}$  645
70:   end if 646
71:   return IndexDequeue( $v.parent$ ,  $sup$ ,  $i$ ) 647
72: end if 648
73: end IndexDequeue 649

74: element FindResponse(int b, int i)  $\triangleright$  find response to  $i$ th dequeue in  $D(root.blocks[b])$  650
75:    $\triangleright$  Precondition:  $1 \leq i \leq |D(root.blocks[b])|$  651
76:    $num_{enq} := root.blocks[b].sum_{enq} - root.blocks[b-1].sum_{enq}$  652
77:   if  $root.blocks[b-1].size + num_{enq} < i$  then 653
78:     return null  $\triangleright$  queue is empty when dequeue occurs 654
79:   else  $\triangleright$  response is the  $e$ th enqueue in the root 655
80:      $e := i + root.blocks[b-1].sum_{enq} - root.blocks[b-1].size$  656
81:     use binary search to find min  $b_e \leq b$  with  $root.blocks[b_e].sum_{enq} \geq e$  657
82:      $i_e := e - root.blocks[b_e-1].sum_{enq}$   $\triangleright$  find rank of enqueue within its block 658
83:     return GetEnqueue( $root$ ,  $b_e$ ,  $i_e$ ) 659
84:   end if 660
85: end FindResponse 661

```

Fig. 6: Queue implementation's helper routines. 662 663

```

664 86: element GetEnqueue(Node  $v$ , int  $b$ , int  $i$ )  $\triangleright$  returns argument of  $i$ th enqueue in  $E(v.blocks[b])$ 
665 87:  $\triangleright$  Preconditions:  $i \geq 1$  and  $v.blocks[b]$  is non-null and contains at least  $i$  enqueues
666 88: if  $v$  is a leaf node then return  $v.blocks[b].element$ 
667 89: else
668 90:    $sum_{left} := v.left.blocks[v.blocks[b].end_{left}].sum_{enq} \triangleright$  # enqueues in  $v.blocks[1..b]$  from  $v.left$ 
669 91:    $prev_{left} := v.left.blocks[v.blocks[b-1].end_{left}].sum_{enq} \triangleright$  # enqueues in  $v.blocks[1..b-1]$  from  $v.left$ 
670 92:    $prev_{right} := v.right.blocks[v.blocks[b-1].end_{right}].sum_{enq} \triangleright$  # enqueues in  $v.blocks[1..b-1]$  from  $v.right$ 
671 93:   if  $i \leq sum_{left} - prev_{left}$  then  $dir := left$   $\triangleright$  required enqueue is in  $v.left$ 
672 94:   else  $\triangleright$  required enqueue is in  $v.right$ 
673 95:      $dir := right$ 
674 96:      $i := i - (sum_{left} - prev_{left})$ 
675 97:   end if
676 98:    $\triangleright$  find enqueue's block in  $v.dir.blocks$  and its rank within block
677 99:   use binary search to find minimum  $b'$  in range  $[v.blocks[b-1].end_{dir}+1..v.blocks[b].end_{dir}]$  such
678 100:   that  $v.dir.blocks[b'].sum_{enq} \geq i + prev_{dir}$ 
679 101:    $i' := i - (v.dir.blocks[b'-1].sum_{enq} - prev_{dir})$ 
680 102:   return GetEnqueue( $v.dir$ ,  $b'$ ,  $i'$ )
681 103: end if
682 103: end GetEnqueue

```

Fig. 7: Queue implementation's GetEnqueue routine.

686 then use its position in the linearization for future
687 Dequeue operations.
688
689 A **Dequeue** also appends a block to the pro-
690 cess's leaf. The block's *element* field is null to
691 indicate that it represents a Dequeue operation.
692 After propagating the operation to the root, the
693 Dequeue calls IndexDequeue to compute its posi-
694 tion in the root and then calls FindResponse to
695 compute its response.

700 **Append**(B) first adds the block B to the invok-
701 ing process's leaf. The leaf's *head* field stores
702 the first empty slot in the leaf's *blocks* array, so
703 the **Append** writes B there and increments *head*.
704 Since **Append** writes only to the process's own
705 leaf, there cannot be concurrent updates to a
706 leaf. **Append** then calls **Propagate** to ensure the
707
708
709
710
711
712
713
714

operation represented by B is propagated to the
root.

Propagate(v) guarantees that any blocks that
are in v 's children when **Propagate** is invoked are
propagated to the root. It uses the double **Refresh**
idea described above and calls **Refresh** on v twice
in Lines 226 and 227. If both calls fail to add
a block to v , it means some other process has
done a successful **Refresh** that propagated blocks
that were in v 's children prior to line 226 to v .
Then, **Propagate** recurses to $v.parent$ to continue
propagating blocks up to the root.

A **Refresh** on node v creates a block repre-
senting the new blocks in v 's children and tries to
append it to $v.blocks$. Line 25 reads $v.head$ into the
local variable h . Line 32 creates the new block B to
install in $v.blocks[h]$. If line 32 returns null instead

of a new block, there were no new blocks in v 's children to propagate to v , so **Refresh** can return true at line 33 and terminate. Otherwise, the CAS at line 35 tries to install B into $v.blocks[h]$. Either this CAS succeeds or some other process has installed a block in this location. Either way, line 36 then calls **Advance** to advance v 's head index from h to $h + 1$ and fill in the *super* field of the most recently appended block. The boolean value returned by **Refresh** indicates whether its CAS succeeded. A **Refresh** may pause after a successful CAS before calling **Advance** at line 36, so other processes help keep *head* up to date by calling **Advance**, either at line 29 during a **Refresh** on v 's parent or line 36 during a **Refresh** on v .

CreateBlock(v, i) is used by **Refresh** to construct a new block B to be installed in $v.blocks[i]$. The end_{left} and end_{right} fields store the indices of the last blocks appended to v 's children, obtained by reading the *head* index in v 's children. Since the sum_{enq} field should store the number of enqueues in $v.blocks[1..i]$ and these enqueues come from $v.left.blocks[1..B.end_{left}]$ and $v.blocks[1..B.end_{right}]$, line 41 sets sum_{enq} to the sum of $v.left.blocks[B.end_{left}].sum_{enq}$ and $v.right.blocks[B.end_{right}].sum_{enq}$. Line 42 sets num_{enq} to the number of enqueues in the new block by subtracting the number of enqueues in $v.blocks[1..i - 1]$ from $B.sum_{enq}$. The values of $B.sum_{deq}$ and num_{deq} are computed similarly. Then, if B is going to be installed in the root,

line 44 computes the *size* field, which represents the number of elements in the queue after the operations in the block are performed. Finally, if the new block contains no operations, **CreateBlock** returns null to indicate there is no need to install it.

Once a dequeue is appended to a block of the process's leaf and propagated to the root, the **IndexDequeue** routine finds the dequeue's location in the root. More precisely, **IndexDequeue**(v, b, i) computes the block in the root and the rank within that block of the i th dequeue of the block B stored in $v.blocks[b]$. Lines 63–65 compute the location of B 's superblock in v 's parent, taking into account the fact that $B.super$ may differ from the superblock's true index by one. The arithmetic in lines 67–70 compute the dequeue's rank within the superblock's sequence of dequeues, using (3.1).

To compute the response of the i th Dequeue in the b th block of the root, **FindResponse**(b, i) determines at line 77 if the queue is empty. If not, line 80 computes the rank e of the Enqueue whose argument is the Dequeue's response. A binary search on the sum_{enq} fields of $root.blocks$ finds the index b_e of the block that contains the e th enqueue. Since the enqueue is linearized before the dequeue, $b_e \leq b$. To find the left end of the range for the binary search for b_e , we can first do a doubling search [6], comparing e to the sum_{enq} fields

715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765

at indices $b-1, b-2, b-4, b-8, \dots$. Then, `GetEnqueue` traces down through the tree to find the required enqueue in a leaf.

GetEnqueue(v, b, i) returns the argument of the i th enqueue in the b th block B of Node v . It recursively finds the location of the enqueue in each node along the path from v to a leaf, which stores the argument explicitly. `GetEnqueue` first determines which child of v contains the enqueue, and then finds the range of blocks within that child that are subblocks of B using information stored in B and the block that precedes B in v . `GetEnqueue` finds the exact subblock containing the enqueue using a binary search on the sum_{enq} field (line 99) and proceeds recursively down the tree.

4 Proof of Correctness

After proving some basic properties in Section 4.1, we show in Section 4.2 that a double refresh at each node suffices to propagate an operation to the root. In Section 4.3 we show `GetEnqueue` and `IndexDequeue` correctly navigate through the tree. Finally, we prove linearizability in Section 4.4.

4.1 Basic Properties

A `Block` object's fields, except for `super`, are immutable: they are written only when the block is created. Moreover, only a CAS at line 53 modifies `super` (from null to a non-null value), so it is

changed only once. Similarly, only a CAS at line 35 modifies an element of a node's `blocks` array (from null to a non-null value), so blocks are permanently added to nodes. Only a CAS at line 55 can update a node's `head` field by incrementing it, which implies the following.

Observation 1. *For each node v , $v.head$ is non-decreasing over time.*

Observation 2. *Let R be an instance of `Refresh`(v) whose call to `CreateBlock` returns a non-null block. When R terminates, $v.head$ is strictly greater than the value R reads from it at line 25.*

Proof. After R 's CAS at line 55, $v.head$ is no longer equal to the value h read at line 25. The claim follows from Observation 1. \square

Now we show $v.blocks[v.head]$ is either the last non-null block or the first null block in node v .

Invariant 3. *For $0 \leq i < v.head$, $v.blocks[i] \neq \text{null}$. For $i > v.head$, $v.blocks[i] = \text{null}$. If $v \neq \text{root}$, $v.blocks[i].super \neq \text{null}$ for $0 < i < v.head$.*

Proof. Initially, $v.head = 1$, $v.blocks[0] \neq \text{null}$ and $v.blocks[i] = \text{null}$ for $i > 0$, so the claims hold.

Assume the claims hold before a change to $v.blocks$, which can be made only by a successful CAS at line 35. The CAS changes $v.blocks[h]$ from null to a non-null value. Since $v.blocks[h]$ is null before the CAS, $v.head \leq h$ by the hypothesis. Since h was read from $v.head$ earlier at line 25, the

current value of $v.head$ is at least h by Observation 1. So, $v.head = h$ when the CAS occurs and a change to $v.blocks[v.head]$ preserves the invariant.

Now, assume the claim holds before a change to $v.head$, which can only be an increment from h to $h + 1$ by a successful CAS at line 55 of **Advance**. For the first two claims, it suffices to show that $v.blocks[head] \neq \text{null}$. **Advance** is called either at line 29 after testing that $v.blocks[h] \neq \text{null}$ at line 28, or at line 36 after the CAS at line 35 ensures $v.blocks[h] \neq \text{null}$. For the third claim, observe that prior to incrementing $v.head$ to $i + 1$ at line 55, the CAS at line 53 ensures that $v.blocks[i].super \neq \text{null}$. \square

It follows that blocks accessed by the **Enqueue**, **Dequeue** and **CreateBlock** routines are non-null.

The following two lemmas show that no operation appears in more than one block of the root.

Lemma 4. *If $b > 0$ and $v.blocks[b] \neq \text{null}$, then*

$$v.blocks[b-1].end_{left} \leq v.blocks[b].end_{left} \text{ and } v.blocks[b-1].end_{right} \leq v.blocks[b].end_{right}.$$

Proof. Let B be the block in $v.blocks[b]$. Before creating B at line 32, the **Refresh** that installed B read b from $v.head$ at line 25. At that time, $v.blocks[b-1]$ contained a block B' , by Invariant 3. Thus, the **CreateBlock**($v, b-1$) that created B' terminated before the **CreateBlock**(v, b) that

created B started. It follows from Observation 1 that the value that line 41 of **CreateBlock**($v, b-1$) stores in $B'.end_{left}$ is less than or equal to the value that line 41 of **CreateBlock**(v, b) stores in $B.end_{left}$. Similarly, the values stored in $B'.end_{right}$ and $B.end_{right}$ at line 41 satisfy the claim. \square

Lemma 5. *If B and B' are two blocks in nodes at the same depth in the ordering tree, their sets of subblocks are disjoint.*

Proof. We prove the lemma by reverse induction on the depth. If B and B' are in leaves, they have no subblocks, so the claim holds. Assume the claim holds for nodes at depth $d + 1$ and let B and B' be two blocks in nodes at depth d . Consider the direct subblocks of B and B' defined by (3.3). If B and B' are in different nodes at depth d , then their direct subblocks are disjoint. If B and B' are in the same node, it follows from Lemma 4 that their direct subblocks are disjoint. Either way, their direct subblocks (at depth $d + 1$) are disjoint, so the claim follows from the induction hypothesis. \square

It follows that each block has at most one superblock. Moreover, we can now prove each operation is contained in at most one block of each node, and hence appears at most once in the linearization L .

Corollary 6. *For $i \neq j$, $v.blocks[i]$ and $v.blocks[j]$ cannot both contain the same operation.*

Proof. A block B contains the operations in B 's subblocks in leaves of the tree. An operation by process P appears in just one block of P 's leaf, so an operation cannot be in two different leaf blocks. By Lemma 5, $v.blocks[i]$ and $v.blocks[j]$ have no common subblocks, so the claim follows. \square

The accuracy of the values stored in the sum_{enq} and sum_{deq} fields on lines 2, 6 and 41 follows easily from the definition of subblocks.

Invariant 7. *If B is a block stored in $v.blocks[i]$, then*

$$B.sum_{enq} = |E(v.blocks[0]) \cdots E(v.blocks[i])| \text{ and}$$

$$B.sum_{deq} = |D(v.blocks[0]) \cdots D(v.blocks[i])|.$$

Proof. Initially, each $blocks$ array contains only an empty block B_0 in location 0. By definition, $E(B_0)$ and $D(B_0)$ are empty sequences. Moreover, $B_0.sum_{enq} = B_0.sum_{deq} = 0$, so the claim is true.

We show that each installation of a block B into some location $v.blocks[i]$ preserves the claim, assuming the claim holds before this installation. We consider two cases.

If v is a leaf, B was created at line 2 or 6. For line 2, B represents a single enqueue, so $|E(B)| = 1$ and $|D(B)| = 0$. Since $B.sum_{enq}$ is set to $v.blocks[i-1].sum_{enq} + 1$ and $B.sum_{deq}$ is set to $v.blocks[i-1].sum_{deq}$, the claim follows from the hypothesis. The proof for line 6, where B represents a single dequeue, is similar.

Now suppose v is an internal node. By the definition of subblocks in (3.3) and Lemma 4, the subblocks of $v.blocks[1..i]$ are $v.left.blocks[1..B.end_{left}]$ and $v.right.blocks[1..B.end_{right}]$. Thus, the enqueues in $E(v.blocks[0]) \cdots E(v.blocks[i])$ are those in $E(v.left.blocks[0]) \cdots E(v.left.blocks[B.end_{left}])$ and those in $E(v.left.blocks[0]) \cdots E(v.left.blocks[B.end_{right}])$. By the hypothesis, the total number of these enqueues is $v.left.blocks[B.end_{left}].sum_{enq} + v.right.blocks[B.end_{right}].sum_{enq}$, which is the value that line 41 stored in $B.sum_{enq}$ when B was created. The proof for sum_{deq} (stored on line 41) is similar. \square

Invariant 7 allows us to prove that every block a Refresh installs contains at least one operation.

Corollary 8. *If a block B is in $v.blocks[i]$ where $i > 0$, then $E(B)$ and $D(B)$ are not both empty.*

Proof. The Refresh that installed B got B as the response to its call to CreateBlock on line 32. Thus, at line 46, $num_{enq} + num_{deq} \neq 0$. By Invariant 7, $num_{enq} = |E(B)|$ and $num_{deq} = |D(B)|$, so these sequences cannot both be empty. \square

4.2 Propagating Operations to the Root

In the next two lemmas, we show two Refreshes suffice to propagate operations from a child to its parent. We say that node v *contains* an operation op if some block in $v.blocks$ contains op . Since

blocks are permanently added to nodes, if v contains op at some time, v contains op at all later times too.

Lemma 9. *Let R be a call to $\text{Refresh}(v)$ that performs a successful CAS on line 35 (or terminates at line 33). After that CAS (or termination, respectively), v contains all operations that v 's children contained when R executed line 25.*

Proof. Suppose v 's child (without loss of generality, $v.\text{left}$) contained an operation op when R executed line 25. Let i be the index such that the block $B_\ell = v.\text{left.blocks}[i]$ contains op . By Observation 1 and Lemma 4, the value of childHead that R reads from $v.\text{left.head}$ in line 27 is at least i . If it is equal to i , R calls Advance at line 29, which ensures that $v.\text{left.head} > i$. Then, R creates a new block B by calling $\text{CreateBlock}(v, h)$ at line 32, where h is the value R reads at line 25. CreateBlock reads a value greater than i from $v.\text{left.head}$ at line 41. Thus, $B.\text{end}_{\text{left}} \geq i$. We consider two cases.

Suppose R 's CAS at line 35 installs B in $v.\text{blocks}$. Then, B_ℓ is a subblock of some block in v , since $B.\text{end}_{\text{left}}$ is greater than or equal to B_ℓ 's index i in $v.\text{left.blocks}$. Hence v contains op , as required.

Now suppose R 's call to CreateBlock returns null, causing R to terminate at line 33. Intuitively, since there are no operations in v 's children to promote, op is already in v . We formalize this

intuition. The value computed at line 41 is

$$\begin{aligned} \text{num}_{\text{enq}} &= v.\text{left.blocks}[B.\text{end}_{\text{left}}].\text{sum}_{\text{enq}} \\ &\quad + v.\text{right.blocks}[B.\text{end}_{\text{right}}].\text{sum}_{\text{enq}} \\ &\quad - v.\text{blocks}[h-1].\text{sum}_{\text{enq}} \\ &= v.\text{left.blocks}[B.\text{end}_{\text{left}}].\text{sum}_{\text{enq}} \\ &\quad + v.\text{right.blocks}[B.\text{end}_{\text{right}}].\text{sum}_{\text{enq}} \\ &\quad - v.\text{left.blocks}[v.\text{blocks}[h-1].\text{end}_{\text{left}}].\text{sum}_{\text{enq}} \\ &\quad - v.\text{right.blocks}[v.\text{blocks}[h-1].\text{end}_{\text{right}}].\text{sum}_{\text{enq}} \end{aligned}$$

It follows from Invariant 7 that num_{enq} is the total number of enqueues in $v.\text{left.blocks}[v.\text{blocks}[h-1].\text{end}_{\text{left}} + 1..B.\text{end}_{\text{left}}]$ and $v.\text{right.blocks}[v.\text{blocks}[h-1].\text{end}_{\text{right}} + 1..B.\text{end}_{\text{right}}]$. Similarly, num_{deq} is the total number of dequeues contained in these blocks. Since $\text{num}_{\text{enq}} + \text{num}_{\text{deq}} = 0$ at line 46, these blocks contain no operations. By Corollary 8, this means the ranges of blocks are empty, so that $v.\text{blocks}[h-1].\text{end}_{\text{left}} \geq B.\text{end}_{\text{left}} \geq i$. Hence, B_ℓ is already a subblock of some block in v , so v contains op . \square

We now show a double Refresh propagates blocks as required.

Lemma 10. *Consider two consecutive terminating calls R_1, R_2 to $\text{Refresh}(v)$ by the same process. All operations contained v 's children when R_1 begins are contained in v when R_2 terminates.*

970 *Proof.* If either R_1 or R_2 performs a successful
 971 CAS at line 35 or terminates at line 33, the claim
 972 follows from Lemma 9. So suppose both R_1 and
 973 R_2 perform a failed CAS at line 35. Let h_1 and
 974 h_2 be the values R_1 and R_2 read from $v.head$ at
 975 line 25. By Observation 2, $h_2 > h_1$. By Lemma
 976 4, $v.blocks[h_2] = \text{null}$ when R_1 executes line 25.
 977 Since R_2 fails its CAS on $v.blocks[h_2]$, some other
 978 Refresh R_3 must have done a successful CAS on
 979 $v.blocks[h_2]$ before R_2 's CAS. R_3 must have exe-
 980 cuted line 25 after R_1 , since R_3 read the value
 981 h_2 from $v.head$ and the value of $v.head$ is non-
 982 decreasing, by Observation 1. Thus, all operations
 983 contained in v 's children when R_1 begins are also
 984 contained in v 's children when R_3 later executes
 985 line 25. By Lemma 9, these operations are con-
 986 tained in v when R_3 performs its successful CAS,
 987 which is before R_2 's failed CAS. \square
 988
 989
 1000
 1001 **Lemma 11.** *When an $\text{Append}(B)$ terminates,*
 1002 *B 's operation is contained in exactly one block in*
 1003 *each node along the path from the process's leaf to*
 1004 *the root.*
 1005
 1006
 1007
 1008
 1009
 1010 *Proof.* Append adds B to the process's leaf and
 1011 calls Propagate , which does a double Refresh on
 1012 each internal node on the path P from the leaf to
 1013 the root. By Lemma 10, this ensures a block in
 1014 each node on P contains B 's operation. There is
 1015 at most one such block in each node, by Corollary
 1016 6. \square

4.3 Correctness of GetEnqueue and IndexDequeue

In this section, we show that the routines used
 to compute responses to Dequeue operations work
 correctly. We first prove the *super* field is accurate,
 since IndexDequeue uses it to trace superblocks up
 the tree. We do this by showing that the *super*
 field of a block B in node v is read from $v.parent$'s
head field close to the time that B 's superblock B_s
 is installed in v 's parent. In particular, $B.super$
 is written after B is installed, but before B_s is
 installed.

Lemma 12. *Let $B = v.blocks[b]$. If $v.parent.blocks[s]$ is the superblock of B then $s - 1 \leq B.super \leq s$.*

Proof. We first show that $B.super \leq s$. Let R_s
 be the instance of $\text{Refresh}(v.parent)$ that installs
 B 's superblock in $v.parent.blocks[s]$. By the defi-
 nition of subblocks (3.3), R_s 's read r of $v.head$ at
 line 41 obtains a value greater than b . By Invari-
 ant 3, $B.super$ is not null when r occurs, which
 means that $B.super$ was set (by line 53) to a
 value read from $v.parent.head$ before r . When r
 occurs, $v.parent.blocks[s] = \text{null}$, since the later
 CAS by R_s at line 35 succeeds. So, by Invariant 3,
 $v.parent.head \leq s$ when r occurs. Since the value
 stored in $B.super$ was read from $v.parent.head$
 before r and the *head* field is non-decreasing by
 Observation 1, it follows that $B.super \leq s$.

Next, we show that $B.super \geq s - 1$. The value stored in $B.super$ at line 53 is read from $v.parent.head$ at line 52 and $head$ is always at least 1, so $B.super \geq 1$. So, if $s \leq 2$, the claim is trivial. Assume $s > 2$ for the rest of the proof. By Lemma 4, $v.parent.blocks[s - 1]$ is not null. Let R_{s-1} be the call to $Refresh(v.parent)$ that installed the block in $v.parent.blocks[s - 1]$. Let r' be the step when R_{s-1} reads $s - 1$ in $v.parent.head$ at line 25. This read r' must be before B is installed in v ; otherwise, Lemma 9 would imply that B is a subblock of one of $v.parent.blocks[1..s - 1]$, contrary to the hypothesis. Now, consider the call to $Advance(v, b)$ that writes $B.super$. It is invoked either at line 29 after seeing $v.blocks[b] \neq \text{null}$ at line 28 or at line 36 after ensuring $v.blocks[b] \neq \text{null}$ at line 35. Either way, the $Advance$ is invoked after B is installed, and therefore after r' . By Observation 1, $v.parent.head$ is non-decreasing, so the value this $Advance$ reads in $v.parent.head$ and writes in $B.super$ is greater than or equal to the value $s - 1$ that r' reads in $v.parent.head$. \square

To show $GetEnqueue$ and $IndexDequeue$ work correctly, we just check that they correctly compute the index of the required block and the operation's rank within the block. For $IndexDequeue$, we use Lemma 12 each time $IndexDequeue$ goes one step up the tree.

Lemma 13. *If $v.blocks[b]$ has been propagated to the root and $1 \leq i \leq |D(v.blocks[b])|$, then*

$IndexDequeue(v, b, i)$ returns $\langle b', i' \rangle$ such that the i th dequeue in $D(v.blocks[b])$ is the (i') th dequeue of $D(\text{root.blocks}[b'])$.

Proof. We prove the claim by induction on the depth of node v . The base case where v is the root is trivial (see Line 60). Assuming the claim holds for v 's parent, we prove it for v . Let $B = v.blocks[b]$ and B' be the superblock of B . $IndexDequeue(v, b, i)$ first computes the index sup of B' in $v.parent$. By Lemma 12, this index is either $B.super$ or $B.super + 1$. The correct index is determined by testing on line 64 whether B is a subblock of $v.parent.blocks[B.super] + 1$.

Next, the position of the required dequeue in $D(B')$ (as defined by Equation (3.1)) is computed in lines 67–70. We first add the number of dequeues in the subblocks of B' in v that precede B on line 67. If v is the right child of its parent, then all of the subblocks of B' from v 's left sibling also precede the required dequeue, so we add the number of dequeues in those subblocks in line 69.

Finally, $IndexDequeue$ is called recursively on v 's parent. Since B has been propagated to the root, so has its superblock B' . Thus, all preconditions of the recursive call are met. By the induction hypothesis, the recursive call returns the location of the required dequeue in the root. \square

Lemma 14. *If $1 \leq i \leq |E(v.blocks[b])|$ then $getEnqueue(v, b, i)$ returns the argument of the i th enqueue in $E(v.blocks[b])$.*

1072 *Proof.* We prove the claim by induction on the
 1073 height of node v . If v is a leaf, the hypothesis
 1074 implies that $i = 1$ and the block $v.blocks[b]$ rep-
 1075 resents an enqueue whose argument is stored in
 1076 $v.blocks[b].element$. `GetEnqueue` returns the argu-
 1077 ment of this enqueue at line 88.

1081 Assuming the claim holds for v 's children,
 1082 we prove it for v . Let B be $v.blocks[b]$. By
 1083 Equation (3.1), $E(B)$ is obtained by concatenat-
 1084 ing the enqueue sequences of the direct subblocks
 1085 of B , which are listed in (3.3). By Invariant 7,
 1086 $sum_{left} - prev_{left}$ is the number of enqueues in
 1087 $E(B)$ that come from B 's subblocks in v 's left
 1088 child. Thus, dir is set to the direction for the
 1089 child of v that contains the required enqueue oper-
 1090 ation. Moreover, when line 97 is reached, i is
 1091 the position of the required enqueue within the
 1092 portion E' of $E(B)$ that comes from that child.
 1093 Thus, line 99 finds the index b' of the subblock B'
 1094 containing the required enqueue. By Invariant 7,
 1095 $v.dir.blocks[b' - 1].sum_{enq} - prev_{dir}$ is the number
 1096 of enqueues in E' before the enqueues of block B' ,
 1097 so the value i' computed on line 100 is the position
 1098 of the required enqueue within $E(B')$. Thus, the
 1099 recursive call on line 101 satisfies its precondition,
 1100 and returns the required result, by the induction
 1101 hypothesis. \square

1116
 1117
 1118
 1119
 1120
 1121
 1122

4.4 Linearizability

We show that the linearization ordering L defined in Equation (3.2) is a legal permutation of a subset of the operations in the execution, i.e., that it includes all operations that terminate and if one operation op_1 terminates before another operation op_2 begins, then op_2 does not precede op_1 in L . We also show each completed dequeue returns the same result as it would in the sequential execution L .

Lemma 15. *L is a legal linearization ordering.*

Proof. By Corollary 6, L is a permutation of a subset of the operations in the execution. By Lemma 11, each terminating operation is propagated to the root before it terminates, so it appears in L . Also, if op_1 terminates before op_2 begins, then op_1 is propagated to the root before op_2 begins, so op_1 appears before op_2 in L . \square

We next show that *size* fields are computed correctly.

Lemma 16. *If the operations of $root.blocks[0..b]$ are applied sequentially in the order of L on an initially empty queue, the resulting queue has $root.blocks[b].size$ elements.*

Proof. We prove the claim by induction on b . The base case when $b = 0$ is trivially true, since the queue is initially empty and $root.blocks[0]$ contains an empty block whose *size* field is 0. Assuming the claim holds for $b - 1$, we prove it

for b . The *size* field of the block B installed in $root.blocks[b]$ is computed at line 44 of a call to $CreateBlock(root, b)$. By the induction hypothesis, $root.blocks[b-1].size$ gives the size of the queue before the operations of block B are performed. By Invariant 7, the values of num_{enq} and num_{deq} are the number of enqueues and dequeues contained in B . Hence, the size of the queue after the operations of B are performed (with enqueues before dequeues as specified by L) is $\max(0, root.blocks[b-1].size + num_{enq} - num_{deq})$. \square

Next, we show each operation returns the same response as it would in the sequential execution L .

Lemma 17. *Each terminating dequeue returns the response it would in the sequential execution L .*

Proof. If a dequeue Deq terminates, it is contained in some block in the root, by Lemma 11. By Lemma 13, Deq 's call to $IndexDequeue$ on line 8 returns a pair $\langle b, i \rangle$ such that Deq is the i th dequeue in the block $B = root.blocks[b]$. Deq then calls $FindResponse(b, i)$ on line 9. By Lemma 16, the queue contains $root.blocks[b-1].size$ elements after the operations in $root.blocks[1..b-1]$ are performed sequentially in the order given by L . By Invariant 7, the value of num_{enq} computed on line 76 is the number of enqueues in B . Since the enqueues in block B precede the dequeues, the queue is empty when the i th dequeue of B occurs if $root.blocks[b-1].size + num_{enq} < i$. So

Deq returns null on line 78 if and only if it would do so in the sequential execution L . Otherwise, the size of the queue after doing the operations in $root.blocks[0..b-1]$ in the sequential execution L is $root.blocks[b-1].sum_{enq}$ minus the number of non-null dequeues in that prefix of L . Hence, line 80 sets e to the rank of Deq among all the non-null dequeues in L . Thus, in the sequential execution L , Deq returns the value enqueued by the e th enqueue in L . By Invariant 7, this enqueue is the i_e th enqueue in $E(root.blocks[b_e])$, where b_e and i_e are the values Deq computes on line 81 and 82. By Lemma 14, the call to $GetEnqueue$ returns the argument of the required enqueue. \square

Combining Lemmas 15 and 17 provides our main result.

Theorem 18. *The queue implementation is linearizable.*

5 Analysis

We now analyze the number of steps and the number of CAS instructions performed by operations.

Proposition 19. *Each Enqueue or Dequeue operation performs $O(\log p)$ CAS instructions.*

Proof. An operation invokes $Refresh$ at most twice at each of the $\lceil \log_2 p \rceil$ levels of the tree. A $Refresh$ does at most 5 CAS steps: one in line 35 and two during each $Advance$ in line 29 or 36. \square

1174 **Lemma 20.** *The search that FindResponse(b, i)*
1175 *does at line 81 to find the index b_e of the block*
1176 *in the root containing the e th enqueue takes*
1177 *$O(\log(\text{root.blocks}[b_e].\text{size} + \text{root.blocks}[b-1].\text{size}))$*
1178 *steps.*
1180
1181 *Proof.* Let the blocks in the root be B_1, \dots, B_ℓ .
1182 The doubling search for b_e takes $O(\log(b - b_e))$
1183 steps, so we prove $b - b_e \leq 2 \cdot B_{b_e}.\text{size} +$
1184 $B_{b-1}.\text{size} + 1$. If $b \leq b_e + 1$, then this is triv-
1185 ial, so assume $b > b_e + 1$. As shown in Lemma
1186 17, the dequeue that calls FindResponse is in B_b
1187 and is supposed to return an enqueue in B_{b_e} .
1188 Thus, there can be at most $B_{b_e}.\text{size}$ dequeues in
1189 $D(B_{b_e+1}) \cdots D(B_{b-1})$; otherwise in the sequen-
1190 tial execution L , all elements enqueued before
1191 the end of $E(B_{b_e})$ would be dequeued before
1192 $D(B_b)$. Furthermore, by Lemma 16, the size
1193 of the queue after the prefix of L correspond-
1194 ing to B_1, \dots, B_{b-1} is $B_{b-1}.\text{size} \geq B_{b_e}.\text{size} +$
1195 $|E(B_{b_e+1}) \cdots E(B_{b-1})| - |D(B_{b_e+1}) \cdots D(B_{b-1})|$.
1196 Thus, $|E(B_{b_e+1}) \cdots E(B_{b-1})| \leq B_{b-1}.\text{size} +$
1197 $|D(B_{b_e+1}) \cdots D(B_{b-1})| \leq B_{b-1}.\text{size} + B_{b_e}.\text{size}$. So,
1198 the total number of operations in $B_{b_e+1}, \dots, B_{b-1}$
1199 is at most $B_{b-1}.\text{size} + 2 \cdot B_{b_e}.\text{size}$. Each of these
1200 $b - 1 - b_e$ blocks contains at least one operation,
1201 by Corollary 8. So, $b - 1 - b_e \leq B_{b-1}.\text{size} + 2 \cdot$
1202 $B_{b_e}.\text{size}$. \square
1203
1204 The following lemma helps bound the time for
1205 GetEnqueue.

Lemma 21. *Each block B in each node contains at most one operation of each process. If c is the execution's maximum point contention, B has at most c direct subblocks.*

Proof. Suppose B contains an operation of process p . Let op be the earliest operation by p contained in B . When op terminates, op is contained in B by Lemma 11. B cannot contain any later operations by p , since B is created before those operations are invoked.

Let t be the earliest termination of any operation contained in B . By Lemma 11, B is created before t , so all operations contained in B are invoked before t . Thus, all are running concurrently at t , so B contains at most c operations. By definition, the direct subblocks of B contain these c operations, and each operation is contained in exactly one of these subblocks, by Lemma 5. By Corollary 8, each direct subblock of B contains at least one operation, so B has at most c direct subblocks. \square

We now bound step complexity in terms of the number of processes p , the maximum contention $c \leq p$, and the size of the queue.

Theorem 22. *Each Enqueue and null Dequeue takes $O(\log p)$ steps and each non-null Dequeue takes $O(\log p \log c + \log q_e + \log q_d)$ steps, where q_d is the size of the queue when the Dequeue is linearized and q_e is the size of the queue when the Enqueue of the value returned is linearized.*

Proof. An **Enqueue** or null **Dequeue** creates a block, appends it to the process's leaf and propagates it to the root. The **Propagate** does $O(1)$ steps at each node on the path from the process's leaf to the root. A null **Dequeue** additionally calls **IndexDequeue**, which also does $O(1)$ steps at each node on this path. So, the total number of steps for either type of operation is $O(\log p)$.

A non-null **Dequeue** must also search at line 81 and call **GetEnqueue** at line 83. By Lemma 20, the doubling search takes $O(\log(q_e + q_d + p))$ steps, since the size of the queue can change by at most p within one block (by Lemma 21). **GetEnqueue** does a binary search within each node on a path from the root to a leaf. Each node v 's search is within the subblocks of one block in v 's parent. By Lemma 21, each such search takes $O(\log c)$ steps, for a total of $O(\log p \log c)$ steps. \square

Corollary 23. *The queue implementation is wait-free.*

6 Bounded-Space Implementation

In the implementation described in Section 3, operations remain in the *blocks* arrays forever. Thus, the space used continues to grow as operations are invoked. Now, we modify the implementation to remove blocks that are no longer needed, so that space usage is polynomial in p and q , while ensuring the (amortized) step complexity is still

polylogarithmic. We replace the *blocks* array in each node by a red-black tree (RBT) that stores the blocks. Each block has an additional *index* field that represents its position within the original *blocks* array, and blocks in a RBT are sorted by *index*. The attempt to install a new block in *blocks*[i] on line 35 is replaced by an attempt to insert a new block with index i into the RBT. Accessing the block in *blocks*[i] is replaced by searching the RBT for the index i . The binary searches for a block in line 81 and 99 can simply search the RBT using the *sum_{enq}* field, since the RBT is also sorted with respect to this field, by Invariant 7.

Known lock-free search trees have step complexity that includes a term linear in p [10, 26]. However, we do not require all the standard search tree operations. Instead of a standard insertion, we allow a **Refresh**'s insertion to fail if another concurrent **Refresh** succeeds in inserting a block, just as the CAS on line 35 can fail if a concurrent **Refresh** does a successful CAS. Moreover, the insertion should succeed only if the newly inserted block has a larger index than any other block in the RBT. Thus, we can use a particularly simple concurrent RBT implementation. A sequential RBT can be made persistent using the classic node-copying technique of Driscoll et al. [9]: all RBT nodes are immutable, and operations on the RBT make a new copy of each RBT node x that must be modified, as well as each RBT node along the path from

the RBT's root to x . The RBT reachable from the new copy of the root is the result of applying the RBT operation. This adds only a constant factor to the running time of any routine designed for a (sequential) RBT. Once a process has performed an update to the RBT representing the blocks of a node v in the ordering tree, it uses a CAS to swing v 's pointer from the previous RBT root to the new RBT root. A search in the RBT can simply read the pointer to the RBT root and perform a standard sequential search on it. Bashari and Woelfel [4] used persistent RBTs in a similar way for a snapshot data structure.

To prevent RBTs from growing without bound, we must discard blocks that are no longer needed. Ensuring the size of the RBT is polynomial in p and q will also keep the running time of our operations polylogarithmic. Blocks should be kept if they contain operations still in progress. Moreover, a block containing an `Enqueue(x)` operation must be kept until x is dequeued.

To maintain good amortized time, we periodically do a garbage collection (GC) phase. If a `Refresh` on a node adds a block whose *index* is a multiple of $G = p^2 \lceil \log p \rceil$, it does GC to remove obsolete blocks from the node's RBT. To determine which blocks can be thrown away, we use a global array `last[1.. p]` where each process writes the index of the last block in the root containing a null dequeue or an enqueue whose element it dequeued. To perform GC, a process reads

`last[1.. p]` and finds the maximum entry m . Then, it helps complete every other process's pending dequeue by computing the dequeue's response and writing it in the block in the leaf that represents the dequeue. Once this helping is complete, it follows from the FIFO property of the queue that elements enqueued in `root.blocks[1.. $m-1$]` have all been dequeued, so GC can discard all subblocks of those. Fortunately, there is an RBT `Split` operation that can remove these obsolete blocks from an RBT in logarithmic time [47, Sec. 4.2].

An operation op 's search of a RBT may fail to find the required block B that has been removed by another process's GC phase. If op is a dequeue, op must have been helped before B was discarded, so op can simply read its response from its own leaf. If op is an enqueue, it can simply terminate.

6.1 Detailed Description

To avoid confusion, we use nodes to refer to the nodes of the ordering tree, and blocks to refer to the nodes of a RBT (since the RBT stores blocks). The space-bounded implementation uses two shared arrays: the *leaf* array allows processes to access one another's leaves to perform helping, and the *last* array is used to determine which blocks are safe to discard.

The *blocks* field of each node in the ordering tree is implemented as a pointer to the root of a RBT of `Blocks` rather than an infinite array. Each RBT is initialized with an empty block with

index 0. Any access to an entry of the *blocks* array is replaced by a search in the RBT. The node's *head* field, which previously gave the next position to insert into the *blocks* array is no longer needed; we can instead simply find the maximum *index* of any block in the RBT. To facilitate this, **MaxBlock** is a query operation on the RBT that returns the block with the maximum *index*. We can store, in the root of the RBT, a pointer to the maximum block so that **MaxBlock** can be done in constant time, without affecting the time of other RBT operations. Similarly, a **MinBlock** query finds the block with the minimum *index* in a RBT.

Blocks no longer require the *super* field. It was used to quickly find a block's superblock in the parent node's *blocks* array, but this can now be done efficiently by searching the parent's *blocks* RBT instead. Each **Block** has an additional field.

- `int index`

that represents the position this block would have in the *blocks* array. To facilitate helping, each **Block** in a leaf has one more additional field

- `Object response`

which is used only for blocks that store a dequeue operation and stores the response of the dequeue in the block.

Pseudocode for the space-bounded implementation appears in Figures 8 to 10. New or modified code appears in blue. The **Propagate** and **GetEnqueue** routines are unchanged. A few lines have been added to **FindResponse** to update the *last*

array to ensure that it stores the value described above. Minor modifications have also been made to **Enqueue**, **Dequeue**, **CreateBlock**, **Refresh** and **Append** to accommodate the switch from an array of blocks to a RBT of blocks (and the corresponding disappearance of the *head* field). In addition, the second half of the **Dequeue** routine is now in a separate routine called **CompleteDeq** so that it can also be used by other processes helping to complete the operation. The **Refresh** routine no longer needs to set the *super* field of blocks since that field has been removed. The **IndexDequeue** routine, which must trace the location of a dequeue along a path from its leaf to the root has a minor modification to search the *blocks* RBT at each level instead of using the *super* field.

The new routines, **AddBlock**, **SplitBlock**, **Help** and **Propagated** are used to implement the garbage collection (GC) phase. When a **Refresh** or **Append** wants to add a new block to a node's *blocks* RBT, it calls the new **AddBlock** routine. Before attempting to add the block to a node's RBT, **AddBlock** triggers a GC phase on the RBT if the new block's *index* is a multiple of the constant G , which we choose to be $p^2 \lceil \log p \rceil$. This ensures that obsolete blocks are removed from the RBT once every G times a new block is added to it. The GC phase uses **SplitBlock** to determine the index s of the oldest block to keep, calls **Help** to help all pending dequeues that have been propagated

1327
1328
1329
1330
1331
1332
1333
1334
1335
1336
1337
1338
1339
1340
1341
1342
1343
1344
1345
1346
1347
1348
1349
1350
1351
1352
1353
1354
1355
1356
1357
1358
1359
1360
1361
1362
1363
1364
1365
1366
1367
1368
1369
1370
1371
1372
1373
1374
1375
1376
1377

```

1378 201: ▷ Shared variables
1379 202: Node root                                ▷ root of ordering tree
1380 203: Node[] leaf[1..p]                       ▷ leaf[k] is the leaf assigned to process k
1381 204: int[] last[1..p]                         ▷ last[k] is max index of a root block that process k saw
1382 205:                                          ▷ contains either a dequeued enqueue or a null dequeue
1383 206: void Enqueue(Object e)
1384 207:   h := MaxBlock(leaf[id].blocks).index+1
1385 208:   let B be a new Block with fields element := e, sumenq := leaf[id].blocks[h-1].sumenq + 1,
1386 209:   sumdeq := leaf[id].blocks[h-1].sumdeq, index := h
1387 210:   Append(B)
1388 211: end Enqueue
1389
1390 212: Object Dequeue()
1391 213:   h := MaxBlock(leaf[id].blocks).index+1
1392 214:   let B be a new Block with fields element := null, sumenq := leaf[id].blocks[h-1].sumenq,
1393 215:   sumdeq := leaf[id].blocks[h-1].sumdeq + 1, index := h
1394 216:   Append(B)
1395 217:   return CompleteDeq(leaf[id], h)
1396 218: end Dequeue
1397 219: Object CompleteDeq(Node leaf, int h)      ▷ finish propagated dequeue in leaf.blocks[h]
1398 220:   ⟨b, i⟩ := IndexDequeue(leaf, h, 1)
1399 221:   return FindResponse(b, i)
1400 222: end CompleteDeq
1401
1402 223: void Append(Block B)                      ▷ append block to leaf and propagate to root
1403 224:   leaf[id].blocks := AddBlock(leaf[id], leaf[id].blocks, B)
1404 225:   Propagate(leaf[id].parent)
1405 226: end Append
1406
1407 227: void Propagate(Node v)                    ▷ propagate blocks from v's children to root
1408 228:   if not Refresh(v) then                  ▷ double refresh
1409 229:     Refresh(v)
1410 230:   end if
1411 231:   if v ≠ root then                        ▷ recurse up tree
1412 232:     Propagate(v.parent)
1413 233:   end if
1414 234: end Propagate
1415
1416 235: boolean Refresh(Node v)                  ▷ try to append a new block B to v.blocks
1417 236:   T := v.blocks
1418 237:   h := MaxBlock(T).index + 1
1419 238:   B := CreateBlock(v, h)
1420 239:   if B = null then return true
1421 240:   else
1422 241:     T' := AddBlock(v, T, B)
1423 242:     return CAS(v.blocks, T, T')
1424 243:   end if
1425 244: end Refresh
1426
1427 245: RBT AddBlock(Node v, RBT T, Block B)    ▷ add block B ≠ null to T; do GC if necessary
1428 246:   if B.index is a multiple of G then    ▷ do garbage collection
1429 247:     s := SplitBlock(v).index
1430 248:     Help
1431 249:     T' := Split(T, s)                  ▷ Split removes blocks with index < s
1432 250:     return Insert(T', B)
1433 251:   else return Insert(T, B)
1434 252:   end if
1435 253: end AddBlock

```

Fig. 8: Bounded-space queue implementation's main routines for process number id .

252:	Block SplitBlock(Node v)	▷ figure out where to split v 's RBT	1429
253:	if $v = \text{root}$ then		1430
254:	$m := 0$		1431
255:	for $k := 1..p$ do $m := \max(m, v.\text{last}[k])$		1432
256:	end for		1433
257:	$B := \text{root.blocks}[m - 1]$		1434
258:	else		1435
259:	$B_p := \text{SplitBlock}(v.\text{parent})$		1436
260:	$\text{dir} := (v = v.\text{parent.left} ? \text{left} : \text{right})$		1437
261:	$B := v.\text{blocks}[B_p.\text{end}_{\text{dir}}]$		1438
262:	end if		1439
263:	return $(B = \text{null} ? \text{MinBlock}(v.\text{blocks}) : B)$ ▷ If B was discarded, use leftmost block instead		1440
264:	end SplitBlock		1441
265:	void Help	▷ help pending operations	1442
266:	for ℓ in $\text{leaf}[1..k]$ do		1443
267:	$B := \text{MaxBlock}(\ell.\text{blocks})$		1444
268:	if $B.\text{element} = \text{null}$ and $B.\text{index} > 0$ and $\text{Propagated}(\ell, B.\text{index})$ then		1445
269:	▷ operation is a propagated dequeue		1446
270:	$B.\text{response} := \text{CompleteDeq}(\ell, B.\text{index})$		1447
271:	end if		1448
272:	end for		1449
273:	end Help		1450
274:	boolean Propagated(Node v , int b)	▷ check if $v.\text{blocks}[b]$ has propagated to root	1451
275:	▷ Precondition: $v.\text{blocks}[b]$ exists		1452
276:	if $v = \text{root}$ then return true		1453
277:	else		1454
278:	$T := v.\text{parent.blocks}$		1455
279:	$\text{dir} := (v.\text{parent.left} = v ? \text{left} : \text{right})$		1456
280:	if $\text{MaxBlock}(T).\text{end}_{\text{dir}} < b$ then return false		1457
281:	else		1458
282:	$B_p :=$ minimum index block in T with $\text{end}_{\text{dir}} \geq b$		1459
283:	return $\text{Propagated}(v.\text{parent}, B_p.\text{index})$		1460
284:	end if		1461
285:	end if		1462
286:	end Propagated		1463
287:	Block CreateBlock(Node v , int i) ▷ create new block B to install in $v.\text{blocks}[i]$		1464
288:	let B be a new Block with fields $\text{end}_{\text{left}} := \text{MaxBlock}(v.\text{left.blocks}).\text{index}$,		1465
	$\text{end}_{\text{right}} := \text{MaxBlock}(v.\text{right.blocks}).\text{index}$, $\text{index} := i$		1466
	$\text{sum}_{\text{enq}} := v.\text{left.blocks}[B.\text{end}_{\text{left}}].\text{sum}_{\text{enq}} + v.\text{right.blocks}[B.\text{end}_{\text{right}}].\text{sum}_{\text{enq}}$		1467
	$\text{sum}_{\text{deq}} := v.\text{left.blocks}[B.\text{end}_{\text{left}}].\text{sum}_{\text{deq}} + v.\text{right.blocks}[B.\text{end}_{\text{right}}].\text{sum}_{\text{deq}}$		1468
289:	$\text{num}_{\text{enq}} := B.\text{sum}_{\text{enq}} - v.\text{blocks}[i - 1].\text{sum}_{\text{enq}}$		1469
290:	$\text{num}_{\text{deq}} := B.\text{sum}_{\text{deq}} - v.\text{blocks}[i - 1].\text{sum}_{\text{deq}}$		1470
291:	if $v = \text{root}$ then $B.\text{size} := \max(0, v.\text{blocks}[i - 1].\text{size} + \text{num}_{\text{enq}} - \text{num}_{\text{deq}})$		1471
292:	end if		1472
293:	if $\text{num}_{\text{enq}} + \text{num}_{\text{deq}} = 0$ then return null ▷ no blocks need to be propagated to v		1473
294:	else return B		1474
295:	end if		1475
296:	end CreateBlock		1476
	Fig. 9: Bounded-space queue implementation's routines for GC and creating new Block.		1477
			1478
			1479


```

1480 297: <int, int> IndexDequeue(Node v, int b, int i)
1481 298:   ▷ return <b', i'> such that i-th dequeue in D(v.blocks[b]) is (i')th dequeue of D(root.blocks[b'])
1482 299:   ▷ Precondition: v.blocks[b] exists and has propagated to root and |D(v.blocks[b])| ≥ i
1483 300:   if v = root then return <b, i>
1484 301:   else                                     ▷ First, find the superblock Bp of v.blocks[b]
1485 302:     dir := (v.parent.left = v ? left : right)
1486 303:     T := v.parent.blocks
1487 304:     Bp := min block in T with enddir ≥ b
1488 305:     B'p := max block in T with enddir < b ▷ predecessor of Bp
1489 306:     ▷ compute index i of dequeue in superblock Bp
1490 307:     i += v.blocks[b-1].sumdeq - v.blocks[B'p.enddir].sumdeq
1491 308:     if dir = right then
1492 309:       i += v.blocks[Bp.endleft].sumdeq - v.blocks[B'p.endleft].sumdeq
1493 310:     end if
1494 311:     return IndexDequeue(v.parent, Bp.index, i)
1495 312:   end if
1496 313: end IndexDequeue
1497 314: element FindResponse(int b, int i) ▷ find response to i-th dequeue in D(root.blocks[b])
1498 315:   ▷ Precondition: 1 ≤ i ≤ |D(root.blocks[b])|
1499 316:   numenq := root.blocks[b].sumenq - root.blocks[b-1].sumenq
1500 317:   if root.blocks[b-1].size + numenq < i then ▷ queue is empty when dequeue occurs
1501 318:     if b > last[id] then last[id] := b
1502 319:   end if
1503 320:   return null
1504 321:   else                                     ▷ response is the e-th enqueue in the root
1505 322:     e := i + root.blocks[b-1].sumenq - root.blocks[b-1].size
1506 323:     use BST search in root.blocks to find minimum index be of a Block with sumenq ≥ e
1507 324:     ie := e - root.blocks[be-1].sumenq ▷ find rank of enqueue within its block
1508 325:     res := GetEnqueue(root, be, ie)
1509 326:     if be > last[id] then last[id] := be
1510 327:   end if
1511 328:   return res
1512 329: end if
1513 330: end FindResponse
1514
1515 331: element GetEnqueue(Node v, int b, int i) ▷ returns argument of i-th enqueue in E(v.blocks[b])
1516 332:   ▷ Preconditions: i ≥ 1 and v.blocks[b] exists and contains at least i enqueues
1517 333:   if v is a leaf node then return v.blocks[b].element
1518 334:   else
1519 335:     sumleft := v.left.blocks[v.blocks[b].endleft].sumenq ▷ # enqueues in v.blocks[1..b] from v.left
1520 336:     prevleft := v.left.blocks[v.blocks[b-1].endleft].sumenq ▷ # enqueues in v.blocks[1..b-1] from v.left
1521 337:     prevright := v.right.blocks[v.blocks[b-1].endright].sumenq ▷ # enqueues in v.blocks[1..b-1] from v.right
1522 338:     if i ≤ sumleft - prevleft then dir := left ▷ required enqueue is in v.left
1523 339:     else                                     ▷ required enqueue is in v.right
1524 340:       dir := right
1525 341:       i := i - (sumleft - prevleft)
1526 342:     end if
1527 343:     ▷ find enqueue's block in v.dir.blocks and its rank within block
1528 344:     use BST search in v.dir.blocks to find minimum index b' of a Block with sumenq ≥ i + prevdir
1529 345:     i' := i - (v.dir.blocks[b'-1].sumenq - prevdir)
1530 346:     return GetEnqueue(v.dir, b', i')
1531 347:   end if
1532 348: end GetEnqueue

```

Fig. 10: Bounded-space queue implementation's routines to compute responses to operations.

to the root (to ensure that all blocks before s can safely be discarded), and then uses the standard RBT **Split** routine [47] to remove all blocks with *index* less than s . Then **AddBlock** inserts the new block at line 248 or 249 and returns the root of the resulting RBT. The **Append** or **Refresh** then stores this root into the node's *blocks* field at line 222 or 240.

To determine the oldest block in a node v to keep, the **SplitBlock** routine first uses the *last* array to find the most recent block B_{root} in the root that contains either an enqueue that has been dequeued or a null dequeue. By the FIFO property of queues, all enqueues in blocks before B_{root} are either dequeued or will be dequeued by a dequeue that is currently in progress. Once those pending dequeues have been helped to complete by line 246, it is safe to discard any blocks in the root older than B_{root} , as well as their subblocks.¹ The **SplitBlock** uses the *end_{left}* and *end_{right}* fields to find the last block in v that is a subblock of B_{root} (or any older block in the root, in case B_{root} has no subblocks in v). While **SplitBlock** is in progress, it is possible that some block that it needs in a node v' along the path from v to the root is discarded by another GC phase. In this case, **SplitBlock** uses the last subblock in v of the oldest block in v' instead

(since a GC phase on v' determined that all blocks older than that are safe to discard anyway).

The **Help** routine is fairly straightforward: it loops through all leaves and helps the dequeue that is in progress there if it has already been propagated to the root. The **Propagated** function is used to determine whether the dequeue has propagated to the root.

In the code, we use $v.blocks[i]$ to refer to the block in the RBT stored in $v.blocks$ with index i . A search for this block may sometimes not find it, if it has already been discarded by another process's GC phase. As mentioned above, if this happens to an enqueue operation, the enqueue can simply terminate because the fact that the block is gone means that another process has helped the enqueue reach the root of the ordering tree. Similarly, if a dequeue operation performs a failed search on a RBT, the dequeue can return the value written in the *response* field of the leaf block that represents the dequeue and terminate, since some other process will have written the *response* there before discarding the needed block. We do not explicitly write this early termination in the pseudocode every time we do a lookup in an RBT. There is one exception to this rule: if an RBT lookup for block B returns null on line 257 or 261 of **SplitBlock** because the required block has been discarded, we continue doing GC, since we do not want a GC phase on one node to be prevented from cleaning up its RBT because a GC phase on

¹If we used the more conservative approach of discarding blocks whose indices are smaller than the *minimum* entry of *last* instead of the maximum, helping would be unnecessary, but then one slow process could prevent GC from discarding any blocks, so the space would not be bounded.

1531
1532
1533
1534
1535
1536
1537
1538
1539
1540
1541
1542
1543
1544
1545
1546
1547
1548
1549
1550
1551
1552
1553
1554
1555
1556
1557
1558
1559
1560
1561
1562
1563
1564
1565
1566
1567
1568
1569
1570
1571
1572
1573
1574
1575
1576
1577
1578
1579
1580
1581

1582 a different node threw away some blocks that were
 1583 needed. Line 263 says what to do in this case.

1584
 1585
 1586

1587 6.2 Correctness

1588
 1589 There are enough changes to the algorithm that a
 1590 new proof of correctness is required. Its structure
 1591 mirrors the proof of the original algorithm, but
 1592 requires additional reasoning to ensure GC does
 1593 not interfere with other routines.

1594
 1595
 1596

1599 6.2.1 Basic Properties

1600
 1601 The following observation describes how the set of
 1602 blocks in a node's RBT can be modified.

1603 **Lemma 24.** *Suppose a step of the algorithm*
 1604 *changes $v.blocks$ from a non-empty tree T to T' .*
 1605 *If the set of index values in T is I , then the set of*
 1606 *index values in T' is $(I \cap [m-1, \infty)) \cup \{\max(I)+1\}$*
 1607 *for some m .*

1608
 1609
 1610
 1611 *Proof.* The RBT of a node is updated only at line
 1612 222 or 240.

1613 If line 222 of an Append operation modifies
 1614 $v.blocks$, then v is a leaf node, and no other process
 1615 ever modifies $v.blocks$. T' was obtained from T by
 1616 calling `AddBlock(v, T, B)`. B was created either by
 1617 the Enqueue or Dequeue that called Append. Either
 1618 way, $B.index = \max(I) + 1$. The `AddBlock` that
 1619 creates T' may optionally Split the RBT at line
 1620 247 and then add B to it. So the claim is satisfied.

1621 If line 240 of a Refresh modifies $v.blocks$, then v
 1622 is an internal node. After reading T from $v.blocks$

at line 234, the Refresh then creates the block B ,
 and calls `AddBlock(v, T, B)` to create T' . Line 288
 sets $B.index = \max(I) + 1$. The `AddBlock` that
 creates T' may optionally Split the RBT and then
 add B to it. So the claim is satisfied. \square

Since each RBT starts with a single block with
 index 0, the following is an easy consequence of
 Lemma 24.

Corollary 25. *The RBT stored in each node v is
 never empty and always stores a set of blocks with
 consecutive indices. Moreover, its maximum index
 can only increase over time.*

Since RBTs are always non-empty, calls to
`MaxBlock` have well-defined answers. Throughout
 the proof, we use $v.blocks[b]$ to refer to the block
 with index b that appeared in v 's tree at some
 time during the execution. It follows from Lemma
 24 and Corollary 25 that each time a new block
 appears in v 's RBT, its index is greater than any
 block that has appeared in v 's RBT earlier. Thus,
 $v.blocks[b]$ is unique, if it exists. We also use this
 notation in the code to indicate that a search of
 the RBT $v.blocks$ should be performed for the
 block with index b .

We now establish that Definition (3.3) of a
 block's subblocks still makes sense by proving the
 analogue of Lemma 4.

Lemma 4'. *If v is an internal node and a block
 with index $h > 0$ has been inserted into $v.blocks$*

then $v.blocks[h-1].end_{left} \leq v.blocks[h].end_{left}$
and $v.blocks[h-1].end_{right} \leq v.blocks[h].end_{right}$.

Proof. The block B with index h was installed into v 's RBT by the CAS at line 240. Suppose that CAS changed the tree from T to T' . Before this CAS, line 234 read the tree T from $v.blocks$, line 235 found a block B' with index $h-1$ in T , and then line 236 created the block B with index $= h$. Since B' was already in T before B was created, the $CreateBlock(v, b-1)$ that created B' terminated before the $CreateBlock(v, b)$ that created B started. By Corollary 25, the value that line 288 of $CreateBlock(v, b-1)$ stores in $B'.end_{left}$ is less than or equal to the value that line 288 of $CreateBlock(v, b)$ stores in $B.end_{left}$. Similarly, the values stored in $B'.end_{right}$ and $B.end_{right}$ at line 288 satisfy the claim. \square

Lemma 4' implies that the nodes of an in-order traversal of any RBT have non-decreasing values of end_{left} (and of end_{right}). Thus, the searches for a block based on end_{left} or end_{right} values at lines 282, 304 and 305, which are used to look for the superblock of a node or its predecessor, can be done using an ordinary BST search.

Lemma 5, Corollary 6, Invariant 7 and Corollary 8 all hold for the modified algorithm. Their proofs are identical to those given in Section 4.1 since they depend only on Lemma 4 (which can be replaced by Lemma 4') and the definition of

subblocks given in (3.3). In particular, Invariant 7 says that nodes in an in-order traversal of a RBT have non-decreasing values of sum_{enq} so the searches for a block based on sum_{enq} values in lines 323 and 344 can be done using an ordinary BST search.

6.2.2 Propagating Operations to the Root

Next, we prove an analogue of Lemma 9. We say a node v contains an operation if some block containing the operation has previously appeared in the RBT $v.block$ (even if the block has been removed from the RBT by a subsequent Split during garbage collection).

Lemma 9'. *Let R be a call to $Refresh(v)$ that performs a successful CAS on line 240 (or terminates at line 237). In the configuration after that CAS (or termination, respectively), v contains all operations that v 's children contained when R executed line 234.*

Proof. Suppose v 's child (without loss of generality, $v.left$) contained an operation op when R executed line 234. Let i be the index of the block B_i containing op that was in $v.left$'s RBT before R executed line 234. We consider two cases.

Suppose R 's call to $CreateBlock$ returns a new block B that is installed in $v.blocks$ by R 's CAS at line 240. The $CreateBlock$ set $B.end_{left}$ to the maximum index in $v.left$'s RBT at line 288. By

1684 Corollary 25, this maximum *index* is bigger than
 1685 *i*. By the definition of subblocks, some block in
 1686 *v* contains B_ℓ as a subblock and therefore *v*
 1688 contains *op*.

1690 Now suppose *R*'s call to CreateBlock returns
 1691 null, causing *R* to terminate at line 237. Let
 1692 *h* be the maximum *index* in *T* plus 1. By
 1693 reasoning identical to the last paragraph of
 1694 Lemma 9's proof, it follows from the fact that
 1695 $num_{enq} + num_{deq} = 0$ at line 293 that the blocks
 1696 $v.left.blocks[v.blocks[h - 1].end_{left} + 1..B.end_{left}]$
 1697 and $v.right.blocks[v.blocks[h - 1].end_{right} +$
 1698 $1..B.end_{right}]$ contain no operations. By Corollary
 1699 8, each block contains at least one opera-
 1700 tion, so these ranges must be empty, and
 1701 $v.blocks[h - 1].end_{left} \geq B.end_{left} \geq i$. This
 1702 implies that the block B_ℓ containing *op* is a
 1703 subblock of some block that has appeared in *v*'s
 1704 RBT, so *op* is contained in *v*. \square

1715 This allows us to show that a double Refresh
 1716 propagates operations up the tree, as in Lemma
 1717 10.

1721 **Lemma 10'.** *Consider two consecutive terminat-*
 1722 *ing calls R_1, R_2 to Refresh(*v*) by the same process.*
 1723 *All operations contained in *v*'s children when R_1*
 1724 *begins are contained in *v* when R_2 terminates.*

1728 *Proof.* If either R_1 or R_2 performs a successful
 1729 CAS at line 240 or terminates at line 237, the claim
 1730 follows from Lemma 9'. So suppose both R_1 and
 1731 R_2 perform a failed CAS at line 240. Then some

other CAS on *v.blocks* succeeds between the time
 each Refresh reads *v.blocks* at line 234 and per-
 forms its CAS at line 240. Consider the Refresh R_3
 that does this successful CAS during R_2 . R_3 must
 have read *v.blocks* after the successful CAS dur-
 ing R_1 . The claim follows from Lemma 9' applied
 to R_3 . \square

Lemma 11 can then be proved in the same way
 as in Section 4.2.

6.2.3 GC Keeps Needed Blocks

The correctness of GetEnqueue and IndexDequeue,
 which are very similar to the original implementa-
 tion, are dependent only on the fact that GC does
 not discard blocks needed by those routines. The
 following results are used to show this.

We say a block is *finished* if

- it has been propagated to the root,
- the value of each enqueue contained in the
 block has either been returned by a dequeue
 or written in the *response* field of a dequeue,
 and
- each dequeue contained in the block has ter-
 minated or some process has written to the
response field in the leaf block that represents
 it.

Intuitively, once a block is finished, operations no
 longer need the block to compute responses to

operations. The following is an immediate consequence of the definition of finished and what it means for an operation to be contained in a block.

Observation 26. *A block is finished if and only all of its subblocks are finished.*

Invariant 27. *If the minimum index of any block in v 's RBT is b_{min} , then each block with index at most b_{min} that was ever added to v is finished.*

Proof. The invariant is true initially, since the minimum *index* block in v 's RBT is the empty block, which is (vacuously) finished.

We show that every step preserves the invariant. We need only consider a step st that modifies a node's RBT. The minimum *index* of v 's RBT can only change when v 's RBT changes, either at line 222 (if v is a leaf) or at line 240 (if v is an internal node). In either case, the step st changes $v.blocks$ from T to T' , where T' is obtained by a call A to $AddBlock(v, T, B)$. (In the case of a leaf v , this is true because only the process that owns the leaf ever writes to $v.blocks$.) If A does not do GC (lines 244–248), then T' is obtained by adding a new block to T , so by Lemma 24 the minimum *index* is unchanged and the invariant is trivially preserved. So consider the case where A performs GC. We must show that any block of v whose index is less than or equal to the minimum *index* in T' is finished when T' is installed in $v.blocks$. Since T' is obtained by discarding all blocks with *index* values less than or equal to s (and adding a

block with a larger index), it suffices to show that all blocks that were ever added to v 's RBT with *index* at most s are finished.

We must examine how A 's call at line 245 to the recursive algorithm `SplitBlock` computes the value of s .

Claim 27.1. *If one of the recursive calls to `SplitBlock(x)` within A 's call to `SplitBlock` returns a block B , then B and all earlier blocks in x are finished when st occurs.*

Proof of Claim. We prove this claim by induction on the depth of x .

For the base case, suppose x is the root. `SplitBlock` finds the maximum value m in *last*, which is the index of some block that contains an operation that is either a null dequeue or an enqueue whose value is the response for a dequeue that has been propagated to the root (since these are the only ways that an entry of *last* can be set to m). By the FIFO property of queues, the values enqueued by enqueues in $root.blocks[1..m-1]$ are all dequeued by operations that have already been installed in *root.blocks* before the end of the `SplitBlock`. Between the termination of A 's call to `SplitBlock` at line 245 and the CAS step st after A terminates, A helps all pending dequeues at line 246. Thus, after this helping (and before step st), $root.blocks[1..m-1]$ are all finished blocks. Since `SplitBlock` returns $root.blocks[m-1]$, the claim is true.

1786 For the induction step, we assume the claim
 1787 holds for x 's parent, and prove it for x . We
 1788 consider two cases.

1790 If `SplitBlock(x)` returns the minimum block of
 1791 x 's RBT at line 263, then the claim follows from
 1792 the assumption that Invariant 27 holds at all times
 1793 before st .

1797 Otherwise, `SplitBlock(x)` returns the block B
 1798 at line 263. By the induction hypothesis, the block
 1799 B_p computed at line 259 (and all earlier blocks
 1800 of x) are finished when st occurs. By Observa-
 1801 tion 26, the block B in x indexed by $B_p.\text{end}_{\text{left}}$ or
 1802 $B_p.\text{end}_{\text{right}}$ is also finished when st occurs. This
 1803 completes the proof of Claim 27.1. \diamond

1809
 1810 If the `Split` at line 247 of A modifies the RBT,
 1811 then it discards all the blocks older than the one
 1812 returned by `SplitBlock` at line 245. By Claim 27.1,
 1813 the minimum block in the new tree will satisfy the
 1814 invariant when st installs the new tree in $v.\text{blocks}$.
 1815 \square

1820
 1821 We remark that Invariant 27 guarantees that
 1822 GC keeps one block that is finished and discards
 1823 all blocks with smaller indices. This is because
 1824 the first unfinished block may still be traversed
 1825 by an operation in the future, and when examin-
 1826 ing that block the operation may need information
 1827 from the preceding block. For example, when `Find-`
 1828 `Response` is called on a block with index b , line 316
 1829 looks up the block with index $b - 1$.

Lemma 28. *If a Dequeue operation fails to find a block in an RBT, then it has been propagated to the root and its result has been written in the response field of the leaf block it created. If an Enqueue operation fails to find a block in an RBT, it has been propagated to the root.*

Proof. Any block that GC removes from an RBT is finished, by Invariant 27, so if an `Enqueue` or `Dequeue` fails to find a block while it is propagating itself up to the root (for example, during the `CreateBlock` routine), then a block containing the operation itself has been removed from a RBT, so the operation has propagated to the root, by the definition of finished. Moreover, by Invariant 27, if the operation is a dequeue, then its result is in its *response* field.

After propagation to the root, a `Dequeue` must access blocks that contain the dequeue, as well as the enqueue whose value it will return (if it is not a null dequeue). If any of those blocks have been removed, it follows from Invariant 27 that the `Dequeue`'s result is written in its *response* field. \square

By Lemma 28, an operation that fails to find a block in an RBT can terminate. If it is a `Dequeue`, it can return the result written in its *response* field. Since no other process updates the RBT in a process's leaf, the block containing the *response* will be the last block in the leaf's RBT, and the last block of an RBT is never removed by GC.

Thus, the *response* field will still be there when a Dequeue needs it.

6.2.4 Linearizability

The correctness of the `IndexDequeue` and `GetEnqueue` operations can be proved in the same way as in Section 4.3, since they are largely unchanged (except for the simplification that `IndexDequeue` can simply search for a block's superblock instead of using the block's *super* field to calculate the superblock's position). They will give the correct response, provided none of the blocks they need to access have been removed by GC. But as we have seen above, if that happens, the `Enqueue` or `Dequeue` can simply terminate.

Similarly, the results of Section 4.4 can be reproved in exactly the same way as for the original algorithm to establish that the space-bounded algorithm is linearizable.

6.3 Analysis

Lemma 21 shows that each block contains at most one operation of each process. Its proof depends only on Lemma 5, Corollary 8 and Lemma 11, which are all still true for the space-bounded implementation. So the same proof of Lemma 21 still applies.

We first bound the size of RBTs. Let q_{max} be the maximum size of the queue at any time during the sequential execution given by the linearization L . Recall that GC is done on a node every G

times its RBT is updated, and we chose G to be $p^2 \lceil \log p \rceil$. Part of the proof of the following lemma is similar to the proof of Lemma 20.

Lemma 29. *If the maximum index in a node's RBT is a multiple of G , then it contains at most $3q_{max} + 5p + 1$ blocks.*

Proof. Consider the invocation A of `AddBlock` that updates a node v 's RBT with the insertion of a block whose *index* is a multiple of G . Then, A performs a GC phase. Let C be the configuration before A invokes `SplitBlock` on line 245. That call to `SplitBlock` recurses up to the root, where it computes m by reading the *last* array. Let L_1 be the prefix of the linearization L corresponding to blocks $1..m$ of the root. Let L_2 be the next segment of the linearization corresponding to blocks $m+1..\ell$ of the root, where ℓ is the last block added to the root's RBT before C .

We first bound the number of operations in L_2 .

The number of enqueues in L_2 whose values are still in the queue at the end of L_2 is at most q_{max} . If the value enqueued by any enqueue in L_2 is not still in the queue at the end of L_2 , then the dequeue in L_2 that dequeued that value must still be in progress at C ; otherwise the process that performed that dequeue would have set its *last* entry to the index of the root block that contains the enqueue, which is greater than m , before C , contradicting the fact that all values in the *last* array at C are less than or equal to m . So, there

are at most p enqueues in L_2 whose values are still in the queue at the end of L_2 . Thus, there are at most $q_{max} + p$ enqueues in L_2 .

If a dequeue in L_2 returns a non-null value in the sequential execution L , then the value it returns was either in the queue at the end of L_1 or it was enqueued during L_2 . Thus, there are at most $q_{max} + (q_{max} + p)$ dequeues in L_2 that return non-null values. Any dequeue in L_2 that returns a null value in the sequential execution L must still be in progress at C ; otherwise the process that performed the dequeue would have set its *last* entry to a value greater than the index of the root block that contains the dequeue prior to C , contradicting the definition of m . So, there are at most p null dequeues in L_2 . Thus, there are at most $2q_{max} + 2p$ dequeues in L_2 .

A 's call to `SplitBlock` determines the index s used to split v 's RBT by following end_{left} and end_{right} pointers from the root down to v . So, the block returned is a subblock of the root block with index m , unless at some point along the path of subblocks the subblock has already been removed by a split, in which case `SplitBlock` returns a subblock of a root block with index $m' > m$.

Next, we bound the number of operations in v 's blocks that are retained when A sets $T' := \text{Split}(T, s)$. Since T was read before C , any operation in T is either in progress at C or has been propagated to the root before C , by Lemma 11. Thus, there are at most p operations in T that

do not appear in $L_1 \cdot L_2$. All the rest of the operations in blocks of T' have been propagated to blocks $m..\ell$ of the root. By Lemma 21, There are at most p operations in block m of the root and we showed above that there are at most $3q_{max} + 3p$ in blocks $m+1..\ell$ of the root. Thus, there are at most $3q_{max} + 5p$ operations in blocks of T' . Since each block is non-empty by Corollary 8, T' contains at most $3q_{max} + 5p$ blocks, and one more block is inserted before A sets v 's *blocks* to the resulting RBT. \square

Corollary 30. *At all times, the size of a node's RBT is $O(q_{max} + p + G)$.*

Proof. Each update to a node's RBT adds at most one block to it, increasing its maximum *index* by 1. Thus, there are at most G updates since the last time its maximum *index* was a multiple of G . The claim follows from Lemma 29. \square

The following theorem bounds the space that is reachable (and therefore cannot be freed by the environment's garbage collector) at any time.

Theorem 31. *The queue data structure uses a maximum of $O(pq_{max} + p^3 \log p)$ words of memory at any time.*

Proof. There are $2p - 1$ nodes in the ordering tree. Aside from the RBT, each node uses $O(1)$ memory words. Each process may hold pointers to $O(1)$ RBTs that are no longer current in local variables.

So the space bound follows from Corollary 30 and the fact that G is chosen to be $p^2 \lceil \log p \rceil$. \square

Performing GC on a node takes $\Theta(p \log p \log(q_{max} + p))$ steps in the worst case (as explained in the following proof), so an individual operation can take up to $\Theta(p \log^2 p \log(q_{max} + p))$ steps if it helps at each node along the path from a leaf to the root. However, we show that operations still have polylogarithmic amortized step complexity.

Theorem 32. *The amortized step complexity of each operation is $O(\log p \log(p + q_{max}))$.*

Proof. It follows from Corollary 30 and our choice of $G = p^2 \lceil \log p \rceil$ that all the RBT routines we use to perform Split, Insert and searches for blocks with a particular index or for a sum_{enq} value (in line 323 or 344) can be done in $O(\log(p + q_{max}))$ steps.

First, we bound the number of steps taken *excluding* the GC phase in line 244–248. An Enqueue or null Dequeue does $O(1)$ RBT operations and other work at each level of the tree during Propagate, for a total of $O(\log p \log(p + q_{max}))$ steps. A non-null Dequeue must also search for a block in the root at line 323 and call GetEnqueue. At each level of the tree, GetEnqueue does $O(1)$ RBT operations (including a search at line 344) and $O(1)$ other steps. Thus, a Dequeue also takes $O(\log p \log(p + q_{max}))$ steps.

Now we consider the additional steps a process takes while doing GC in line 244–248 and show that the amortized number of GC steps each operation performs is also $O(\log p \log(p + q_{max}))$. If a process does GC in a call to AddBlock(v, T, B) where B has *index* $r \cdot G$ for some integer r , we call this the process's r th GC phase on v .

We argue that each process P can do an r th GC phase on v at most once. Consider P 's first call A to AddBlock that does an r th GC phase on v . Let v, T, B be the arguments of A . Any call to AddBlock on internal node v is from line 239 of Refresh, so $B.index$ is the maximum *index* in T plus 1. The Refresh that called A performed a CAS at line 240. Either the CAS succeeds or it fails because some other CAS changes $v.blocks$ from T to another tree. Either way, by Lemma 24, v 's RBT's maximum *index* will be at least $r \cdot G$ at all times after this CAS. So if a subsequent Refresh by process P ever calls AddBlock on v again, the block it passes as the third argument will have *index* $> r \cdot G$, so P will not perform an r th GC phase on v again.

Each GC phase takes $O(p)$ steps for SplitBlock to read the *last* array and figure out where to split the *blocks* RBT, $O(p \log p \log(p + q_{max}))$ steps in Help, and $O(\log(p + q_{max}))$ steps to split and insert a new node into the RBT. Thus, for each integer r and each node v , a total of $O(p^2 \log p \log(p + q_{max}))$ steps are performed by all processes during their r th helping phase on v . We can amortize

1939
1940
1941
1942
1943
1944
1945
1946
1947
1948
1949
1950
1951
1952
1953
1954
1955
1956
1957
1958
1959
1960
1961
1962
1963
1964
1965
1966
1967
1968
1969
1970
1971
1972
1973
1974
1975
1976
1977
1978
1979
1980
1981
1982
1983
1984
1985
1986
1987
1988
1989

these steps over the operations that appear in $v.blocks[(r-1)G+1..rG]$. By Corollary 8, there are at least G such operations, so each operation's amortized number of steps for GC at each node along the path from its leaf to the root is $O(p^2 \log p \log(p+q_{max})/G) = O(\log(p+q_{max}))$. Hence each operation's amortized number of GC steps is $O(\log p \log(p+q_{max}))$. \square

The implementation remains wait-free: the depth of recursion in each routine is bounded by the height of the tree and the only loop is the counted loop in the `Help` routine. Moreover, since each operation still does only two CAS instructions at each level of the tree (at line 240), the following proposition still holds for the space-bounded version of the queue.

Proposition 19'. *Each operation performs $O(\log p)$ CAS instructions in the worst case.*

7 Future Directions

Our focus was on optimizing step complexity for worst-case executions. However, our queue has a higher cost than the MS-queue in the best case (when an operation runs by itself). Perhaps our queue could be made adaptive by having an operation capture a starting node in the ordering tree (as in [1]) rather than starting at a statically assigned leaf. A possible application of our queue might be to use it as the slow path in the fast-path slow-path methodology [29] to get a queue

that has good performance in practice while also having good worst-case step complexity.

It would be interesting to close the gap that remains between our queue, which takes $O(\log^2 p + \log q)$ steps per operation, and Jayanti, Tarjan and Boix-Adserà's $\Omega(\log p)$ lower bound [22]. For the more relaxed bag data structure the gap is larger between the $\Omega(\min(c, \log \log p))$ lower bound [3] and our upper bound of $O(\log^2 p + \log q)$. Could the complexity for either queues or bags be made polylogarithmic in p while being independent of the size q of the data structure?

The approach used here to implement a lock-free queue could be applied to obtain other lock-free data structures with a polylogarithmic step complexity. For example, we can easily adapt the routines of the implementation in Section 3 to implement a restricted kind of vector data structure that stores a sequence and provides three operations: `Append(e)` to add an element e to the end of the sequence, `Get(i)` to read the i th element in the sequence, and `Index(e)` to compute the position of element e in the sequence. Only the `Append` operations need to be propagated to the root of the ordering tree since the other two operations do not affect the state of the object. An `Append(e)` is implemented like `Enqueue(e)` in $O(\log p)$ steps. A `Get(i)` is similar to `GetEnqueue`, taking $O(\log n + \log^2 p)$ steps when the vector has n elements. An `Index` is similar to `IndexDequeue` (except operating

on enqueues instead of dequeues) and would take $O(\log p)$ steps if the argument is a pointer to the leaf block that contains the element e .

Building on the work described in this paper, Asbell and Ruppert [2] have designed a doubly-ended queue with polylogarithmic amortized step complexity, which also yields a stack as a special case. This required a substantially different representation of the data stored in the ordering tree. Whether the ordering tree could also be used to obtain a priority queue with polylogarithmic step complexity remains an open question.

Acknowledgements

We thank Franck van Breugel, Siddhartha Jayanti, Philipp Woelfel and the anonymous referees for their helpful comments. Funding was provided by the Natural Sciences and Engineering Research Council of Canada.

References

- [1] Yehuda Afek, Dalia Dauber, and Dan Touitou. Wait-free made fast. In *Proc. 27th ACM Symposium on Theory of Computing*, pages 538–547, 1995.
- [2] Shalom Asbell and Eric Ruppert. Wait-free stacks and dequeues with polylogarithmic step complexity. Manuscript, 2023.
- [3] Hagit Attiya and Arie Fouren. Lower bounds on the amortized time complexity of shared objects. In *21st International Conference on Principles of Distributed Systems*, volume 95 of *LIPICs*, pages 16:1–16:18, 2017.
- [4] Benyamin Bashari and Philipp Woelfel. An efficient adaptive partial snapshot implementation. In *Proc. ACM Symposium on Principles of Distributed Computing*, pages 545–555, 2021.
- [5] Naama Ben-David, Guy E. Blelloch, Panagiota Fatourou, Eric Ruppert, Yihan Sun, and Yuanhao Wei. Space and time bounded multiversion garbage collection. In *Proc. 35th International Symposium on Distributed Computing*, volume 209 of *LIPICs*, pages 12:1–12:20, 2021.
- [6] Jon Louis Bentley and Andrew Chi-Chih Yao. An almost optimal algorithm for unbounded searching. *Information Processing Letters*, 5(3):82–87, 1976.
- [7] Robert Colvin and Lindsay Groves. Formal verification of an array-based nonblocking queue. In *10th International Conference on Engineering of Complex Computer Systems*, pages 507–516. IEEE, 2005.
- [8] Matei David. A single-enqueuer wait-free queue implementation. In *Proc. 18th International Conference on Distributed Computing*, volume 3274 of *LNCS*, pages 132–143. Springer, 2004.
- [9] James R. Driscoll, Neil Sarnak, Daniel D. Sleator, and Robert E. Tarjan. Making data structures persistent. *Journal of Computer*

- 2092 *and System Sciences*, 38(1):86–124, February
2093 1989.
2094
- 2095 [10] Faith Ellen, Panagiota Fatourou, Joanna
2096 Helga, and Eric Ruppert. The amortized
2097 complexity of non-blocking binary search
2098 trees. In *Proc. 33rd ACM Symposium on*
2099 *Principles of Distributed Computing*, pages
2100 332–340, 2014.
2101
- 2105 [11] Faith Ellen, Vijaya Ramachandran, and
2106 Philipp Woelfel. Efficient fetch-and-
2107 increment. In *Proc. International Symposium*
2108 *on Distributed Computing*, volume 7611 of
2109 *LNCS*, pages 16–30. Springer, 2012.
2110
- 2113 [12] Faith Ellen and Philipp Woelfel. An optimal
2114 implementation of fetch-and-increment. In
2115 *Proc. 27th International Symposium on Dis-*
2116 *tributed Computing*, volume 8205 of *LNCS*,
2117 pages 284–298. Springer, 2013.
2118
- 2122 [13] Panagiota Fatourou and Nikolaos D. Kalli-
2123 manis. Highly-efficient wait-free synchro-
2124 nization. *Theory of Computing Systems*,
2125 55(3):475–520, 2014.
2126
- 2128 [14] Mikhail Fomitchev and Eric Ruppert. Lock-
2129 free linked lists and skip lists. In *Proc. 23rd*
2130 *ACM Symposium on Principles of Distributed*
2131 *Computing*, pages 50–59, 2004.
2132
- 2135 [15] Anders Gidenstam, Håkan Sundell, and
2136 Philippas Tsigas. Cache-aware lock-free
2137 queues for multiple producers/consumers and
2138 weak memory consistency. In *Proc. 14th*
2139 *International Conference on Principles of*
2140 *Distributed Systems*, volume 6490 of *LNCS*,
2141 pages 302–317. Springer, 2010.
- 2142 [16] Andreas Haas. *Fast Concurrent Data Struc-*
tures Through Timestamping. PhD thesis,
University of Salzburg, 2015.
- [17] Maurice Herlihy. Wait-free synchroniza-
tion. *ACM Trans. Program. Lang. Syst.*,
13(1):124–149, 1991.
- [18] Maurice P. Herlihy and Jeannette M. Wing.
Linearizability: A correctness condition for
concurrent objects. *ACM Trans. Program.*
Lang. Syst., 12(3):463–492, 1990.
- [19] Moshe Hoffman, Ori Shalev, and Nir Shavit.
The baskets queue. In *Proc. 11th Inter-*
national Conference on Principles of Dis-
tributed Systems, volume 4878 of *LNCS*,
pages 401–414. Springer, 2007.
- [20] Prasad Jayanti. A time complexity lower
bound for randomized implementations of
some shared objects. In *Proc. 17th ACM*
Symposium on Principles of Distributed
Computing, pages 201–210, 1998.
- [21] Prasad Jayanti and Srdjan Petrovic.
Logarithmic-time single deleter, multiple
inserter wait-free queues and stacks. In
Foundations of Software Technology and
Theoretical Computer Science, volume 3821
of *LNCS*, pages 408–419. Springer, 2005.
- [22] Siddhartha V. Jayanti, Robert E. Tarjan, and
Enric Boix-Adserà. Randomized concurrent
set union and generalized wake-up. In *Proc.*

- ACM Symposium on Principles of Distributed Computing*, pages 187–196, 2019.
- [23] Colette Johnen, Adnane Khattabi, and Alessia Milani. Efficient wait-free queue algorithms with multiple enqueueers and multiple dequeuers. In *Proc. 26th International Conference on Principles of Distributed Systems*, volume 253 of *LIPICs*, pages 4:1–4:19, February 2023.
- [24] Michael Kenzel, Stefan Lemme, Richard Membarth, Matthias Kurtenacker, Hugo Dvillers, Markus Steinberger, and Philipp Slusallek. AnyQ: An evaluation framework for massively-parallel queue algorithms. In *Proc. 37th IEEE International Parallel and Distributed Processing Symposium*, 2023. To appear.
- [25] Pankaj Khanchandani and Roger Wattenhofer. On the importance of synchronization primitives with low consensus numbers. In *Proc. 19th International Conference on Distributed Computing and Networking*, pages 18:1–18:10, 2018.
- [26] Jeremy Ko. The amortized analysis of a non-blocking chromatic tree. *Theoretical Computer Science*, 840:59–121, November 2020.
- [27] Alex Kogan and Maurice Herlihy. The future(s) of shared data structures. In *Proc. ACM Symposium on Principles of Distributed Computing*, pages 30–39, 2014.
- [28] Alex Kogan and Erez Petrank. Wait-free queues with multiple enqueueers and dequeuers. In *Proc. 16th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*, pages 223–234, 2011.
- [29] Alex Kogan and Erez Petrank. A methodology for creating fast wait-free data structures. *ACM SIGPLAN Not.*, 47(8):141–150, 2012.
- [30] Nikita Koval, Dan Alistarh, and Roman Elizarov. Fast and scalable channels in Kotlin coroutines. In *Proc. ACM Symposium on Principles and Practice of Parallel Programming*, pages 107–118, 2023.
- [31] Edya Ladan-Mozes and Nir Shavit. An optimistic approach to lock-free FIFO queues. *Distributed Computing*, 20(5):323–341, 2008.
- [32] Zongpeng Li. Non-blocking implementations of queues in asynchronous distributed shared-memory systems. Master’s thesis, University of Toronto, 2001. Available from <https://tspace.library.utoronto.ca/bitstream/1807/16583/1/MQ62967.pdf>.
- [33] Henry Massalin and Carlton Pu. A lock-free multiprocessor OS kernel. Technical Report CUCS-005-91, Department of Computer Science, Columbia University, 1991.
- [34] Maged M. Michael and Michael L. Scott. Nonblocking algorithms and preemption-safe locking on multiprogrammed shared memory multiprocessors. *Journal of Parallel and Distributed Computing*, 21:181–191, 1991.

- 2194 *Distributed Computing*, 51(1):1–26, 1998.
- 2195 [35] Gal Milman-Sela, Alex Kogan, Yossi Lev,
2196 Victor Luchangco, and Erez Petrank. BQ: A
2197 lock-free queue with batching. *ACM Trans.*
2198 *Parallel Comput.*, 9(1):5:1–5:49, March 2022.
- 2200 [36] Mark Moir, Daniel Nussbaum, Ori Shalev,
2201 and Nir Shavit. Using elimination to imple-
2202 ment scalable and lock-free FIFO queues.
2203 In *Proc. 17th ACM Symposium on Paral-*
2204 *lelism in Algorithms and Architectures*, pages
2205 253–262, 2005.
- 2206 [37] Adam Morrison and Yehuda Afek. Fast con-
2207 current queues for x86 processors. In *Proc.*
2208 *ACM SIGPLAN Symposium on Principles*
2209 *and Practice of Parallel Programming*, pages
2210 103–112, 2013.
- 2211 [38] Hossein Naderibeni. A wait-free queue
2212 with poly-logarithmic worst-case step com-
2213 plexity. Master’s thesis, York University,
2214 Toronto, Canada, November 2022. Avail-
2215 able from [https://yorkspace.library.yorku.](https://yorkspace.library.yorku.ca/xmlui/handle/10315/40975)
2216 [ca/xmlui/handle/10315/40975](https://yorkspace.library.yorku.ca/xmlui/handle/10315/40975).
- 2217 [39] Ruslan Nikolaev. A scalable, portable, and
2218 memory-efficient lock-free FIFO queue. In
2219 *Proc. 33rd International Symposium on Dis-*
2220 *tributed Computing*, volume 146 of *LIPIcs*,
2221 pages 28:1–28:16, 2019.
- 2222 [40] Ruslan Nikolaev and Binoy Ravindran. wCQ:
2223 A fast wait-free queue with bounded mem-
2224 ory usage. In *Proc. 34th ACM Symposium on*
2225 *Parallelism in Algorithms and Architectures*,
2226 pages 307–319, 2022.
- 2227 [41] Peter Pirkelbauer, Reed Milewicz, and
2228 Juan Felipe Gonzalez. A portable lock-free
2229 bounded queue. In *Proc. 16th Interna-*
2230 *tional Conference on Algorithms and Archi-*
2231 *tectures for Parallel Processing*, volume 10048
2232 of *LNCS*, pages 55–73, 2016.
- 2233 [42] Pedro Ramalhete and Andreia Correia.
2234 Poster: A wait-free queue with wait-free
2235 memory reclamation. *ACM SIGPLAN Not.*,
2236 52(8):453–454, January 2017.
- 2237 [43] Raed Romanov and Nikita Koval. The state-
2238 of-the-art LCRQ concurrent queue algorithm
2239 does NOT require CAS2. In *Proc. ACM Sym-*
2240 *posium on Principles and Practice of Parallel*
2241 *Programming*, pages 14–26, 2023.
- 2242 [44] Eric Ruppert. Analysing the average time
2243 complexity of lock-free data structures. Pre-
2244 sented at BIRS Workshop on Complexity and
2245 Analysis of Distributed Algorithms, 2016.
Available from [http://www.birs.ca/videos/](http://www.birs.ca/videos/2016)
[2016](http://www.birs.ca/videos/2016).
- 2246 [45] Niloufar Shafiei. Non-blocking array-based
2247 algorithms for stacks and queues. In *Proc.*
2248 *10th International Conference on Distributed*
2249 *Computing and Networking*, volume 5408 of
2250 *LNCS*, pages 55–66. Springer, 2009.
- 2251 [46] Niloufar Shafiei. Non-blocking doubly-linked
2252 lists with good amortized complexity. In
2253 *Proc. 19th International Conference on Prin-*
2254 *ciples of Distributed Systems*, volume 46 of

<i>LIPICs</i> , pages 35:1–35:17, 2015.	2245
[47] Robert Endre Tarjan. <i>Data Structures and Network Algorithms</i> . SIAM, Philadelphia, USA, 1983.	2246
	2247
	2248
	2249
	2250
[48] R.K. Treiber. Systems programming: Coping with parallelism. Technical Report RJ 5118, IBM Almaden Research Center, 1986.	2251
	2252
	2253
	2254
	2255
	2256
[49] Philippas Tsigas and Yi Zhang. A simple, fast and scalable non-blocking concurrent FIFO queue for shared memory multiprocessor systems. In <i>Proc. 13th ACM Symposium on Parallel Algorithms and Architectures</i> , pages 134–143, 2001.	2257
	2258
	2259
	2260
	2261
	2262
	2263
	2264
	2265
	2266
[50] Chaoran Yang and John M. Mellor-Crummey. A wait-free queue as fast as fetch-and-add. In <i>Proc. 21st ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming</i> , pages 16:1–16:13, 2016.	2267
	2268
	2269
	2270
	2271
	2272
	2273
	2274
	2275
	2276
	2277
	2278
	2279
	2280
	2281
	2282
	2283
	2284
	2285
	2286
	2287
	2288
	2289
	2290
	2291
	2292
	2293
	2294
	2295